Model Recognition as Planning

Diego Aineto and **Sergio Jiménez** and **Eva Onaindia**

Departamento de Sistemas Informáticos y Computación
Universitat Politècnica de València.
Camino de Vera s/n. 46022 Valencia, Spain
{dieaigar,serjice,onaindia}@dsic.upv.es

Abstract

Given a set of possible action models and a partially observed plan execution, *model recognition* is the task of identifying which of these action models has the highest probability of producing the given plan execution. *Model recognition* is relevant because it enables identifying algorithms by their execution and because, once an action model is recognized, the planning model-based machinery becomes applicable. This paper formalizes the *model recognition* task and proposes a method to assess the probability of a given STRIPS action model to produce a partially observed plan execution. This method, that we called *model recognition as planning*, is robust to missing data in the intermediate states and actions of the observed plan execution besides, it is computable with an off-the-shelf classical planner. The effectiveness of *model recognition as planning* is shown with sets of STRIPS models that represent different *Turing Machines*. We show that *model recognition as planning* effectively identifies the executed *Turing Machine* despite intermediate machine state, applied transitions or tape values, are unobserved.

Introduction

Plan recognition is the task of predicting the future actions of an agent provided the observation of its current behaviour. *Goal recognition* is a subtask of plan recognition with the particular aim of discovering the final objectives of the observed agent. Diverse methods has been proposed for plan and goal recognition (Carberry 2001) such as *rule-based systems*, *parsing* (both conventional and stochastic), *graph-covering*, *Bayesian nets*, . . .

Plan recognition as planning recently defined the model-based approach for plan/goal recognition (Ramírez 2012; Ramírez and Geffner 2009). This approach assumes that the action model of the observed agent is known and leverages it to compute the most likely goal according to the observed plan execution.

In this paper we introduce the *model recognition* task as the task of identifying the action model with the highest probability of producing an observed plan execution. This task is of interest because :

- Once the action model is recognized the planning model-based machinery (Ghallab, Nau, and Traverso 2004;

Geffner and Bonet 2013) becomes applicable.

- It enables identifying algorithms by observing their execution. Many computational models are compilable into classical planning (Baier, Fritz, and McIlraith 2007; Bonet, Palacios, and Geffner 2010; Segovia-Aguas, Jiménez, and Jonsson 2016).

The paper introduces *model recognition as planning*, a novel method to assess the probability of a given STRIPS action model to produce an observed plan execution. This method is robust to missing data in the intermediate states and actions of the observed plan execution besides, it is computable with an off-the-shelf classical planner.

Last but not least, the paper evaluates the effectiveness of *model recognition as planning* with sets of STRIPS models that represent different *Turing Machines*, all of them defined within the same tape alphabet and same machine states. We show that *model recognition as planning* effectively identifies the executed *Turing Machine* despite intermediate machine state, applied transitions or tape values, are unobserved.

Background

This section defines the classical planning model and the STRIPS action model that we use for formalizing the *model recognition as planning* method.

Classical planning

We use F to denote the set of *fluents* (propositional variables) describing a state. A *literal* l is a valuation of a fluent $f \in F$; i.e. either $l = f$ or $l = \neg f$. A set of literals L represents a partial assignment of values to fluents (without loss of generality, we will assume that L does not contain conflicting values). We use $\mathcal{L}(F)$ to denote the set of all literal sets on F; i.e. all partial assignments of values to fluents.

A *state* s is a full assignment of values to fluents; $|s| = |F|$, so the size of the state space is $2^{|F|}$. Explicitly including negative literals $\neg f$ in states simplifies subsequent definitions but often we will abuse of notation by defining a state s only in terms of the fluents that are true in s, as it is common in STRIPS planning.

A *classical planning frame* is a tuple $\Phi = \langle F, A \rangle$, where F is a set of fluents and A is a set of actions. An action $a \in A$ is defined with *preconditions*, $\mathsf{pre}(a) \subseteq \mathcal{L}(F)$, *positive*

effects, $\text{eff}^+(a) \subseteq \mathcal{L}(F)$, and *negative effects* $\text{eff}^-(a) \subseteq \mathcal{L}(F)$. We say that an action $a \in A$ is *applicable* in a state s iff $\text{pre}(a) \subseteq s$. The result of applying a in s is the *successor state* denoted by $\theta(s, a) = \{s \setminus \text{eff}^-(a)) \cup \text{eff}^+(a)\}$.

The result of applying action a in state s is the *successor* state $\theta(s, a) = \{s \setminus \text{eff}_c^-(s, a)) \cup \text{eff}_c^+(s, a)\}$ where $\text{eff}_c^-(s, a) \subseteq triggered(s, a)$ and $\text{eff}_c^+(s, a) \subseteq triggered(s, a)$ are, respectively, the triggered *negative* and *positive* effects.

A *classical planning problem* is a tuple $P = \langle F, A, I, G \rangle$, where I is an initial state and $G \subseteq \mathcal{L}(F)$ is a goal condition. A *plan* for P is an action sequence $\pi = \langle a_1, \ldots, a_n \rangle$ that induces the *state trajectory* $\langle s_0, s_1, \ldots, s_n \rangle$ such that $s_0 = I$ and a_i ($1 \leq i \leq n$) is applicable in s_{i-1} and generates the successor state $s_i = \theta(s_{i-1}, a_i)$. The *plan length* is denoted with $|\pi| = n$. A plan π *solves* P iff $G \subseteq s_n$; i.e. if the goal condition is satisfied in the last state resulting from the application of the plan π in the initial state I.

STRIPS action schemas

This work addresses the learning and evaluation of PDDL action schemas that follow the STRIPS requirement (McDermott et al. 1998; Fox and Long 2003). Figure 1 shows the *stack* action schema, coded in PDDL, from a four-operator *blocksworld* (Slaney and Thiébaux 2001).

```
(:action stack
 :parameters (?v1 ?v2 - object)
 :precondition (and (holding ?v1) (clear ?v2))
 :effect (and (not (holding ?v1)) (not (clear ?v2))
              (handempty) (clear ?v1) (on ?v1 ?v2)))
```

Figure 1: STRIPS operator schema coding, in PDDL, the *stack* action from a four-operator *blocksworld*.

To formalize the target of the learning and evaluation tasks, we assume that fluents F are instantiated from a set of *predicates* Ψ, as in PDDL. Each predicate $p \in \Psi$ has an argument list of arity $ar(p)$. Given a set of *objects* Ω, the set of fluents F is induced by assigning objects in Ω to the arguments of predicates in Ψ, i.e. $F = \{p(\omega) : p \in \Psi, \omega \in \Omega^{ar(p)}\}$ s.t. Ω^k is the k-th Cartesian power of Ω.

Let $\Omega_v = \{v_i\}_{i=1}^{\max_{a \in A} ar(a)}$ be a new set of objects ($\Omega \cap \Omega_v = \emptyset$), denoted as *variable names*, and that is bound by the maximum arity of an action in a given planning frame. For instance, in a three-block *blocksworld* $\Omega = \{block_1, block_2, block_3\}$ while $\Omega_v = \{v_1, v_2\}$ because the operators with the maximum arity, stack and unstack, have arity two.

We define F_v, a new set of fluents s.t. $F \cap F_v = \emptyset$, produced instantiating Ψ using only *variable names*, and that defines the elements that can appear in the action schemes. In *blocksworld* this set contains 11 elements, $F_v = \{$handempty, holding(v_1), holding(v_2), clear(v_1), clear(v_2), ontable(v_1), ontable(v_2), on(v_1, v_1), on(v_1, v_2), on(v_2, v_1), on$(v_2, v_2)\}$.

In more detail, for a given operator schema ξ, we define $F_v(\xi) \subseteq F_v$ as the subset of elements that can appear in that action schema. For instance, for the *stack* action schema $F_v(\text{stack}) = F_v$ while $F_v(\text{pickup}) = \{$handempty, holding(v_1), clear(v_1), ontable(v_1), on$(v_1, v_1)\}$ excludes the elements from F_v that involve v_2 because pickup(v_1) has arity one.

We assume also that actions $a \in A$ are instantiated from STRIPS operator schemas $\xi = \langle head(\xi), pre(\xi), add(\xi), del(\xi) \rangle$ where:

- $head(\xi) = \langle name(\xi), pars(\xi) \rangle$, is the operator *header* defined by its name and the corresponding *variable names*, $pars(\xi) = \{v_i\}_{i=1}^{ar(\xi)}$. The headers of a four-operator *blocksworld* are pickup(v_1), putdown(v_1), stack(v_1, v_2) and unstack(v_1, v_2).

- The preconditions $pre(\xi) \subseteq F_v$, the negative effects $del(\xi) \subseteq F_v$, and the positive effects $add(\xi) \subseteq F_v$ such that, $del(\xi) \subseteq pre(\xi)$, $del(\xi) \cap add(\xi) = \emptyset$ and $pre(\xi) \cap add(\xi) = \emptyset$.

Given the set of predicates Ψ and the header of the operator schema ξ, $2^{2|F_v(\xi)|}$ defines the size of the space of possible STRIPS models for that operator. Note that the previous constraints require that negative effects appear as preconditions and that they cannot be positive effects and also, that a positive effect cannot appear as a precondition. For the *blocksworld*, $2^{2|F_v(stack)|} = 4194304$ while for the pickup operator this number is only 1024.

We say that two STRIPS operator schemes ξ and ξ' are *comparable* if both schemas have the same parameters so they share the same space of possible STRIPS models (formally, iff $pars(\xi) = pars(\xi')$. For instance, we can claim that blocksworld operators stack and unstack are *comparable* while stack and pickup are not. Last but not least, two STRIPS action models $\mathcal{M}$ and $\mathcal{M}'$ are *comparable* iff there exists a bijective function $\mathcal{M} \mapsto \mathcal{M}^*$ that maps every $\xi \in \mathcal{M}$ to a comparable action schema $\xi' \in \mathcal{M}'$ and vice versa.

Model Recognition

Given a finite and non-empty set of *comparable* action models $M = \{\mathcal{M}_1, \ldots, \mathcal{M}_m\}$, and a partially observed plan execution $\mathcal{T} = \langle s_0, a_1, s_1, \ldots, a_n, s_n \rangle$ that is obtained watching the execution of a plan $\pi = \langle a_1, \ldots, a_n \rangle$ s.t., for each $1 \leq i \leq n$, a_i is applied in s_{i-1} and generates s_i. *Model recognition* is the task of identifying which model $\mathcal{M} \in M$ has the highest probability of producing $\mathcal{T}$.

The STRIPS edit distance

Our approach for *model recognition* is to assess how well an action model $\mathcal{M} \in M$ explains $\mathcal{T}$ according to the amount of *edition* required by the model $\mathcal{M}$ to produce $\mathcal{T}$.

In this work we assume that action models $\mathcal{M} \in M$ are represented with the STRIPS formalism. We define here the two allowed *operations* to edit an STRIPS action model:

- *Deletion*. A fluent $pre_f(\xi)/del_f(\xi)/add_f(\xi)$ is removed from the operator schema $\xi \in \mathcal{M}$, such that $f \in F_v(\xi)$.

- *Insertion*. A fluent $pre_f(\xi)/del_f(\xi)/add_f(\xi)$ is added to the operator schema $\xi \in \mathcal{M}$, s.t. $f \in F_v(\xi)$.

We can now formalize an *edit distance* that quantifies how similar two given STRIPS action models are. The distance is symmetric and meets the *metric axioms* provided that the two *edit operations*, deletion and insertion, have the same positive cost.

Definition 1. *Let $\mathcal{M}$ and $\mathcal{M}'$ be two* comparable STRIPS *action models. The* **edit distance**, *denoted as $\delta(\mathcal{M}, \mathcal{M}')$, is the minimum number of* edit operations *that is required to transform $\mathcal{M}$ into $\mathcal{M}'$.*

Since F_v is a bound set, the maximum number of edits that can be introduced to a given action model defined within F_v is bound as well. In more detail, for an operator schema $\xi \in \mathcal{M}$ the maximum number of edits that can be introduced to their precondition set is $|F_v(\xi)|$ while the max number of edits that can be introduced to the effects is twice $|F_v(\xi)|$.

Definition 2. *The **maximum edit distance** of an* STRIPS *action model $\mathcal{M}$ built from the set of possible elements F_v is $\delta(\mathcal{M}, *) = \sum_{\xi \in \mathcal{M}} 3 \times |F_v(\xi)|.$*

We define now an edit distance to asses the matching of a learned action model with respect to a plan trace $\mathcal{T}$.

Definition 3. *Given $\mathcal{M}$, a* STRIPS *action model built from F_v and $\mathcal{T}$, a plan trace built with fluents in F and actions in A. The **observation edit distance**, denoted by $\delta(\mathcal{M}, \mathcal{T})$, is the minimal edit distance from $\mathcal{M}$ to any* comparable *model $\mathcal{M}'$ s.t. $\mathcal{M}'$ produces a valid plan trace $\mathcal{T}$;*

$$\delta(\mathcal{M}, \mathcal{T}) = \min_{\forall \mathcal{M}' \to \mathcal{T}} \delta(\mathcal{M}, \mathcal{M}')$$

The $\delta(\mathcal{M}, \mathcal{T})$ distance could also be defined assessing the edition required by the observed plan execution to match the given model. This implies defining *edit operations* that modify $\mathcal{T}$ instead of modifying $\mathcal{M}$ (Sohrabi, Riabov, and Udrea 2016). Our definition of the *observation edit distance* is more practical since normally F_v is smaller than F (the number of *variable objects* is normally smaller than the number of objects in the observed states).

The STRIPS probability distribution

According to the *Bayes* rule, the probability of an hypothesis $\mathcal{H}$ provided the observation $\mathcal{O}$ is given by the expression $P(\mathcal{H}|\mathcal{O}) = \frac{P(\mathcal{O}|\mathcal{H})P(\mathcal{H})}{P(\mathcal{O})}$. In our scenario, the hypotheses are about the set of possible STRIPS action models $\mathcal{M} \in M$ while the given observation is a partially observed plan execution $\mathcal{T}$.

We call the STRIPS *probability distribution*, denoted by $P(\mathcal{M}|\mathcal{T})$, to the probability distribution of the comparable STRIPS models (within the $F_v(\xi)$ sets) provided that the plan execution $\mathcal{T}$ is observed. This probability distribution can be estimated in three-steps by:

1. Estimating the a priori probabilities $P(\mathcal{T})$ and $P(\mathcal{H})$. For instance, given the set of predicates Ψ and the given a set of operator headers (in other words, given the $F_v(\xi)$ sets) the size of the set of possible STRIPS models set is $\prod_\xi 2^{2|F_v(\xi)|}$. If we assume that a priori all models are equiprobable this means that $P(\mathcal{M}) = \frac{1}{\prod_\xi 2^{2|F_v(\xi)|}}$. With respect to the observations, given Ψ and a set of objects

Ω, the number of possible state observations of length n in a trace $\mathcal{T}$ is $2^{n \times |F|}$ while the number of possible action observations of length n in a trace $\mathcal{T}$ is $2^{n \times |A|}$. If we assume that a priori all traces are equiprobable, this means $P(\mathcal{T}) = \frac{1}{2^{n \times |F|} \times 2^{n \times |A|}}$.

2. Estimating the conditional probability $P(\mathcal{T}|\mathcal{M})$. For instance, we can compute the *observation edit distance* $\delta(\mathcal{M}, \mathcal{T})$ for every possible model $\mathcal{M} \in M$ and map this distance into a likelihood with the following expression $1 - \frac{\delta(\mathcal{M}, \mathcal{T})}{\delta(\mathcal{M}, *)}$.

3. Applying the Bayes rule to obtain the normalized posterior probabilities since these probabilities must sum 1.

Model Recognition as Planning

Now we show a classical planning compilation for computing the *observation edit distance* $\delta(\mathcal{M}, \mathcal{T})$ for each $\mathcal{M} \in M$. The compilation follows these assumptions

1. The initial state $s_0 \in \mathcal{T}$ is *fully observable*.

2. Intermediate actions $a_i \in \mathcal{T}$ and states $s_i \in \mathcal{T}$ s.t. $1 \leq i \leq n$, can be *partially observed*. This means that some fluents in s_i may be missing because its value is unknown. In the extreme, entire states s_i can be missing.

3. $\mathcal{T}$ is *noiseless*, meaning that if a fluent or an action is observed it is the correct value.

Conditional effects

Conditional effects allow us to compactly define the actions output by our compilation. An action $a \in A$ with conditional effects is defined as a set of *preconditions* $\mathsf{pre}(a) \in \mathcal{L}(F)$ and a set of *conditional effects* $\mathsf{cond}(a)$. Each conditional effect $C \rhd E \in \mathsf{cond}(a)$ is composed of two sets of literals $C \in \mathcal{L}(F)$, the *condition*, and $E \in \mathcal{L}(F)$, the *effect*. An action $a \in A$ is *applicable* in a state s if and only if $\mathsf{pre}(a) \subseteq s$, and the *triggered effects* resulting from the action application are the effects whose conditions hold in s:

$$triggered(s, a) = \bigcup_{C \rhd E \in \mathsf{cond}(a), C \subseteq s} E,$$

The result of applying action a in state s is the *successor* state $\theta(s, a) = \{s \setminus \mathsf{eff}_c^-(s, a)) \cup \mathsf{eff}_c^+(s, a)\}$ where $\mathsf{eff}_c^-(s, a) \subseteq triggered(s, a)$ and $\mathsf{eff}_c^+(s, a) \subseteq triggered(s, a)$ are, respectively, the triggered *negative* and *positive* effects.

The compilation

The intuition behind the compilation is that a solution to the resulting classical planning task is a sequence of actions that:

1. **Edits the action model $\mathcal{M}$ to build $\mathcal{M}'$.** A solution plan starts with a *prefix* that modifies the preconditions and effects of the action schemes in $\mathcal{M}$ using to the two *edit operations* defined above, *deletion* and *insertion*.

2. **Validates the edited model $\mathcal{M}'$ in the observed plan trace.** The solution plan continues with a postfix that validates the edited model $\mathcal{M}'$ on the given observations $\mathcal{T}$.

To illustrate this, Figure 2 shows the plan for editing a given *blockswold* action model where the positive effects `(handempty)` and `(clear ?v1)` of the `stack` schema are missing. The edited action model is validated at a four action plan for inverting a two-block tower. Our interest is not in the resulting action model $\mathcal{M}'$ but in the number of required *edit operations* (insertions and deleitions) for that $\mathcal{M}'$ is validated in the given observations, e.g. $\delta(\mathcal{M}, \mathcal{T}) = 2$ for the example in Figure 2. In this case $\delta(\mathcal{M}, *) = 3 \times 2 \times (11 + 5)$ since there are 4 action schemes (`pickup`, `putdown`, `stack` and `unstack`) and $|F_v| = |F_v(stack)| = |F_v(unstack)| = 11$ while $|F_v(pickup)| = |F_v(putdown)| = 5$. The *observation edit distance* is exactly computed if the classical planning task resulting from our compilation is optimally solved (according to the number of edit actions); is approximated if it is solved with a satisfying planner; and is a less accurate estimate (but faster to be computed) if the solved task is a relaxation of the classical planning task that results from our compilation (Bonet and Geffner 2001).

```
00 : (insert_add_handempty_stack)
01 : (insert_add_clear_stack_var1)
02 : (apply_unstack blockB blockA i1 i2)
03 : (apply_putdown blockB i2 i3)
04 : (apply_pickup blockA i3 i4)
05 : (apply_stack blockA blockB i4 i5)
06 : (validate_1)
```

Figure 2: Plan for editing and validating a given *blockswold* model.

Given $\langle \mathcal{M}, \Psi, \mathcal{T} \rangle$ the compilation outputs a classical planning task $P = \langle F, A, I, G \rangle$:

- F contains:
  - The set of fluents F built instantiating the predicates Ψ with the objects Ω that appear in the plan trace given as input, i.e. the blocks A and B in the example of Figure 2. Formally, $\Omega = \bigcup_{s \in \mathcal{T}} obj(s)$, where obj is a function that returns the objects that appear in a given state.
  - Fluents $pre_f(\xi)$, $del_f(\xi)$ and $add_f(\xi)$, for every $f \in F_v(\xi)$, that represent the programmed action model. If a fluent $pre_f(\xi)/del_f(\xi)/add_f(\xi)$ holds, it means that f is a precondition/negative/positive effect in the schema $\xi \in \mathcal{M}'$. For instance, the preconditions of the *stack* schema (Figure 1) are represented by the pair of fluents `pre_holding_stack_v1` and `pre_clear_stack_v2` set to True.
  - The fluents $F_\pi = \{plan(name(a_i), \Omega^{ar(a_i)}, i)\}_{1 \leq i \leq n}$ to code the i^{th} action in $\mathcal{T}$. The static facts $next_{i,i+1}$ and the fluents at_i, $1 \leq i < n$, are also added to iterate through the n steps of $\mathcal{T}$.
  - The fluents $mode_{prog}$ and $mode_{val}$ to indicate whether the operator schemas are programmed or validated, and the fluents $\{test_i\}_{1 \leq i \leq n}$, indicating the state observation $s_i \in \mathcal{T}$ where the action model is validated.
- I encodes the first state observation, $s_0 \subseteq F$ and sets to true $mode_{prog}$ as well as the fluents F_π plus fluents at_1 and $\{next_{i,i+1}\}$, $1 \leq i < n$, for tracking the plan

step where the action model is validated. Our compilation assumes that initially, operator schemas are programmed with every possible precondition (the most specific learning hypothesis), no negative effect and no positive effect. Therefore fluents $pre_f(\xi)$, for every $f \in F_v(\xi)$, hold also at the initial state.

- $G = \bigcup_{1 \leq i \leq n} \{test_i\}$, requires that the programmed action model is validated in the state observations $s_i \in \mathcal{T}$.
- A comprises three kinds of actions:
  1. Actions for *programming* operator schema $\xi \in \mathcal{M}$:
     - Actions for **removing** a *precondition* $f \in F_v(\xi)$ from the action schema $\xi \in \mathcal{M}$.

     $$\mathsf{pre}(\mathsf{programPre}_{f,\xi}) = \{\neg del_f(\xi), \neg add_f(\xi), mode_{prog}, pre_f(\xi)\},$$
     $$\mathsf{cond}(\mathsf{programPre}_{f,\xi}) = \{\emptyset\} \triangleright \{\neg pre_f(\xi)\}.$$

     - Actions for **adding** a *negative* or *positive* effect $f \in F_v(\xi)$ to the action schema $\xi \in \mathcal{M}$.

     $$\mathsf{pre}(\mathsf{programEff}_{f,\xi}) = \{\neg del_f(\xi), \neg add_f(\xi), mode_{prog}\},$$
     $$\mathsf{cond}(\mathsf{programEff}_{f,\xi}) = \{pre_f(\xi)\} \triangleright \{del_f(\xi)\},$$
     $$\{\neg pre_f(\xi)\} \triangleright \{add_f(\xi)\}.$$

     Besides these actions A also contains the actions for *inserting* a precondition and for *deleting* a negative/positive effect.
  2. Actions for *applying* a programmed operator schema $\xi \in \mathcal{M}$ bound with objects $\omega \subseteq \Omega^{ar(\xi)}$. Since operators headers are given as input, the variables $pars(\xi)$ are bound to the objects in ω that appear at the same position. Figure 3 shows the PDDL encoding of the action for applying a programmed operator $stack$ from *blocksworld*.

     $$\mathsf{pre}(\mathsf{apply}_{\xi,\omega}) = \{pre_f(\xi) \implies p(\omega)\}_{\forall p \in \Psi, f = p(pars(\xi))}$$
     $$\cup \{\neg mode_{val}\},$$
     $$\mathsf{cond}(\mathsf{apply}_{\xi,\omega}) = \{del_f(\xi)\} \triangleright \{\neg p(\omega)\}_{\forall p \in \Psi, f = p(pars(\xi))},$$
     $$\{add_f(\xi)\} \triangleright \{p(\omega)\}_{\forall p \in \Psi, f = p(pars(\xi))},$$
     $$\{mode_{prog}\} \triangleright \{\neg mode_{prog}\},$$
     $$\{\emptyset\} \triangleright \{mode_{val}\}.$$

     The extra conditional effects $\{at_i, plan(name(a_i), \Omega^{ar(a_i)}, i)\} \triangleright \{\neg at_i, at_{i+1}\}_{\forall i \in [1,n]}$ are included in the $\mathsf{apply}_{\xi,\omega}$ actions to validate that actions are applied, exclusively, in the same order as in $\mathcal{T}$.

  3. Actions for *validating* the partially observed state $s_i \in \mathcal{T}, 1 \leq i < n$.
     $$\mathsf{pre}(\mathsf{validate}_i) = s_i \cup \{test_j\}_{j \in 1 \leq j < i}$$
     $$\cup \{\neg test_j\}_{j \in i \leq j \leq n} \cup \{mode_{val}\},$$
     $$\mathsf{cond}(\mathsf{validate}_i) = \{\emptyset\} \triangleright \{test_i, \neg mode_{val}\}.$$

Evaluation

To evaluate the empirical performance of *model recognition as planning* we defined a set of possible STRIPS models, each representing a different *Turing Machine*, but all sharing the same set of machine states and same tape alphabet.

```
(:action apply_stack
 :parameters (?o1 - object ?o2 - object)
 :precondition
  (and (or (not (pre_on_stack_v1_v1)) (on ?o1 ?o1))
       (or (not (pre_on_stack_v1_v2)) (on ?o1 ?o2))
       (or (not (pre_on_stack_v2_v1)) (on ?o2 ?o1))
       (or (not (pre_on_stack_v2_v2)) (on ?o2 ?o2))
       (or (not (pre_ontable_stack_v1)) (ontable ?o1))
       (or (not (pre_ontable_stack_v2)) (ontable ?o2))
       (or (not (pre_clear_stack_v1)) (clear ?o1))
       (or (not (pre_clear_stack_v2)) (clear ?o2))
       (or (not (pre_holding_stack_v1)) (holding ?o1))
       (or (not (pre_holding_stack_v2)) (holding ?o2))
       (or (not (pre_handempty_stack)) (handempty)))
 :effect
  (and (when (del_on_stack_v1_v1) (not (on ?o1 ?o1)))
       (when (del_on_stack_v1_v2) (not (on ?o1 ?o2)))
       (when (del_on_stack_v2_v1) (not (on ?o2 ?o1)))
       (when (del_on_stack_v2_v2) (not (on ?o2 ?o2)))
       (when (del_ontable_stack_v1) (not (ontable ?o1)))
       (when (del_ontable_stack_v2) (not (ontable ?o2)))
       (when (del_clear_stack_v1) (not (clear ?o1)))
       (when (del_clear_stack_v2) (not (clear ?o2)))
       (when (del_holding_stack_v1) (not (holding ?o1)))
       (when (del_holding_stack_v2) (not (holding ?o2)))
       (when (del_handempty_stack) (not (handempty)))
       (when (add_on_stack_v1_v1) (on ?o1 ?o1))
       (when (add_on_stack_v1_v2) (on ?o1 ?o2))
       (when (add_on_stack_v2_v1) (on ?o2 ?o1))
       (when (add_on_stack_v2_v2) (on ?o2 ?o2))
       (when (add_ontable_stack_v1) (ontable ?o1))
       (when (add_ontable_stack_v2) (ontable ?o2))
       (when (add_clear_stack_v1) (clear ?o1))
       (when (add_clear_stack_v2) (clear ?o2))
       (when (add_holding_stack_v1) (holding ?o1))
       (when (add_holding_stack_v2) (holding ?o2))
       (when (add_handempty_stack) (handempty))
       (when (modeProg) (not (modeProg))))))
```

Figure 3: PDDL action for applying an already programmed schema $stack$ (implications are coded as disjunctions).

Modeling *Turing Machines* with STRIPS

A *Turing machine* is a tuple $\mathcal{M} = \langle Q, q_o, Q_\perp, \Sigma, \Upsilon, \square, \delta \rangle$:

- Q, is a finite and non-empty set of machine states such that $q_0 \in Q$ is the initial state of the machine and $Q_\perp \subseteq Q$ is the subset of acceptor states.

- Σ is the *tape alphabet*, that is a finite non-empty set of symbols that contains the *input alphabet* $\Upsilon \subseteq \Sigma$ (the subset of symbols allowed to initially appear in the tape) and the *blank symbol* $\square \in \Upsilon$ (the only symbol allowed to occur on the tape infinitely often).

- $\delta : \Sigma \times (Q \setminus Q_\perp) \to \Sigma \times Q \times \{left, right\}$ is the *transition function*. For each possible pair of tape symbol and machine state, δ defines (1), the tape symbol to print at the current position of the header (2), the new state of the machine and (3), whether the header is shifted *left* or *right* after the print operation. If δ is not defined for the current pair of tape symbol and machine state, the machine halts.

Figure 4 shows the δ function of a *Turing Machine* for recognizing the $\{a^n b^n c^n : n \geq 1\}$ language. The *tape*

	q_0	q_1	q_2	q_3	q_4	$\underline{q_5}$
a	x,r,q_1	a,r,q_1	-	a,l,q_3	-	-
b	-	y,r,q_2	b,r,q_2	b,l,q_3	-	-
c	-	-	z,l,q_3	-	-	-
x	-	-	-	x,r,q_0	-	-
y	y,r,q_4	y,r,q_1	-	y,l,q_3	y,r,q_4	-
z	-	-	z,r,q_2	z,l,q_3	z,r,q_4	-
$\square$	-	-	-	-	$\square$,r,q_5	-

Figure 4: Seven-symbol six-state *Turing Machine* for recognizing the $\{a^n b^n c^n : n \geq 1\}$ language ($\underline{q_5}$ is the only acceptor state).

alphabet is $\Sigma = \{a, b, c, x, y, z, \square\}$, the *input alphabet* $\Upsilon = \{a, b, c, \square\}$ and the possible machine states are $Q = \{q_0, q_1, q_2, q_3, q_4, \underline{q_5}\}$ where $\underline{q_5}$ is the only acceptor state.

A classical planning frame $\Phi = \langle F, A \rangle$ can encode the *transition function* δ of a *Turing Machine* $\mathcal{M}$ as follows:

- Fluents F are instantiated from a set of four *predicates* Ψ: (head ?x) that encodes the current position of the header in the tape. (next ?x1 ?x2) encoding that the cell ?x2 follows cell ?x1 in the tape. (symbol-σ ?x) encoding that the tape cell ?x contains the symbol $\sigma \in \mathcal{T}$. (state-q) encoding that $q \in Q$ is the current machine state. Given a set of *objects* Ω that represent the cells in the tape of the given Turing Machine, the set of fluents F is induced by assigning objects in Ω to the arguments of the predicates in Ψ.

- Actions A are instantiated from STRIPS operator schema. For each transition in δ, a STRIPS action schema is defined:

  - The **header** is transition-id(?xl ?x ?xr) where id uniquely identifies the transition in δ. Parameters ?xl, ?x and ?xr are tape cells.

  - The **preconditions** are (head ?x) and (next ?x1 ?x) (next ?x ?xr) to make ?x the tape cell pointed by the header and ?xl and ?xr its left and right neighbours. Preconditions (symbol-σ ?x) and (state-q) are also included to capture the symbol pointed by the header and the currrent machine state.

  - The **delete effects** remove the symbol pointed by the header and the currrent machine state while the **positive effects** set the new symbol pointed by the header and the new machine state.

The STRIPS action schema of Figure 5 models the rule $a, q_0 \to x, r, q_1$ of the *Turing Machine* defined in Figure 4 (the full encoding of the *Turing Machine* defined in Figure 4 produces a total of sixteen STRIPS action schema).

Since the *transition function* of a *Turing Machine* can be encoded as STRIPS action schemas, executions of that machine are definable as *plan traces* $\mathcal{T} = \langle s_0, a_1, s_1, \ldots, a_n, s_n \rangle$. In this case s_0 encodes the initial state of the tape plus, the position of the header and the initial machine state.

```
(:action transition-1        ;;; a,q_0 → x,r,q_1
 :parameters (?xl ?x ?xr)
 :precondition (and (head ?x) (symbol-a ?x) (state-q_0)
                    (next ?xl ?x) (next ?x ?xr))
 :effect (and (not (head ?x))
              (not (symbol-a ?x)) (not (state-q_0))
              (head ?xr) (symbol-x ?x) (state-q_1)))
```

Figure 5: STRIPS action schema that models the transition $a, q_0 \rightarrow x, r, q_1$ of the Turing Machine defined in Figure 4.

Experimental setup

We randomly generated a $M = \{\mathcal{M}_1, \ldots, \mathcal{M}_{100}\}$ set of one-hundred different *Turing Machines* where each $\mathcal{M} \in M$ is a two-symbol three-state *Turing Machine*. We randomly choose a machine $\mathcal{M} \in M$ and produce an fifty-step execution plan trace $\mathcal{T} = \langle s_0, a_1, s_1, \ldots, a_{50}, s_{50} \rangle$. The experiment is following our *model recognition as planning* method to identify the *Turing Machine* that produced $\mathcal{T}$. The exeperiment is repeated for different amounts of missing information in the input trace $\mathcal{T}$.

Reproducibility MADAGASCAR is the classical planner we used to solve the instances that result from our compilations for its ability to deal with dead-ends (Rintanen 2014). Due to its SAT-based nature, MADAGASCAR can apply the actions for programming preconditions in a single planning step (in parallel) because there is no interaction between these actions. Actions for programming effects can also be applied in a single planning step, thus significantly reducing the planning horizon.

The compilation source code, evaluation scripts and benchmarks (including the used training and test sets) are fully available at this anonymous repository *https://github.com/anonsub/observations-learning* so any experimental data reported in the paper can be reproduced.

Results

Conclussions

References

Baier, J. A.; Fritz, C.; and McIlraith, S. A. 2007. Exploiting procedural domain control knowledge in state-of-the-art planners. In *ICAPS*, 26–33.

Bonet, B., and Geffner, H. 2001. Planning as heuristic search. *Artificial Intelligence* 129(1-2):5–33.

Bonet, B.; Palacios, H.; and Geffner, H. 2010. Automatic derivation of finite-state machines for behavior control. In *AAAI Conference on Artificial Intelligence*.

Carberry, S. 2001. Techniques for plan recognition. *User Modeling and User-Adapted Interaction* 11(1-2):31–48.

Fox, M., and Long, D. 2003. PDDL2.1: An extension to PDDL for expressing temporal planning domains. *Journal of Artificial Intelligence Research* 20:61–124.

Geffner, H., and Bonet, B. 2013. A concise introduction to models and methods for automated planning.

Ghallab, M.; Nau, D.; and Traverso, P. 2004. *Automated Planning: theory and practice*. Elsevier.

McDermott, D.; Ghallab, M.; Howe, A.; Knoblock, C.; Ram, A.; Veloso, M.; Weld, D.; and Wilkins, D. 1998. PDDL – The Planning Domain Definition Language.

Ramírez, M., and Geffner, H. 2009. Plan recognition as planning. In *International Joint conference on Artifical Intelligence*, 1778–1783.

Ramírez, M. 2012. *Plan recognition as planning*. Ph.D. Dissertation, Universitat Pompeu Fabra.

Rintanen, J. 2014. Madagascar: Scalable planning with sat. *Proceedings of the 8th International Planning Competition (IPC-2014)*.

Segovia-Aguas, J.; Jiménez, S.; and Jonsson, A. 2016. Hierarchical finite state controllers for generalized planning. In *International Joint Conference on Artificial Intelligence, IJCAI-16*, 3235–3241. AAAI Press.

Slaney, J., and Thiébaux, S. 2001. Blocks world revisited. *Artificial Intelligence* 125(1-2):119–153.

Sohrabi, S.; Riabov, A. V.; and Udrea, O. 2016. Plan recognition as planning revisited. In *IJCAI*, 3258–3264.