

Learning and Recognition of STRIPS Action Models from State Observations

Diego Aineto and Sergio Jiménez and Eva Onaindia

Departamento de Sistemas Informáticos y Computación

Universitat Politècnica de València.

Camino de Vera s/n. 46022 Valencia, Spain

{dieaigar,serjice,onaindia}@dsic.upv.es

Abstract

This paper presents a classical planning compilation for learning STRIPS action models from state observations. The compilation approach does not require observing the precise actions that produced the state observations because such actions are determined by a classical planner. In addition, the compilation is extensible to (1) semantically evaluate the quality of learned STRIPS models and (2), estimate the probability distribution of the possible STRIPS models given a sequence of state observations.

1 Introduction

Besides *plan synthesis* [Ghallab *et al.*, 2004], planning action models are also useful for *plan/goal recognition* [Ramírez, 2012]. At these planning tasks, automated planners are required to reason about an action model that correctly and completely captures the possible world transitions [Geffner and Bonet, 2013]. Unfortunately, building planning action models is complex, even for planning experts, and this knowledge acquisition task is a bottleneck that limits the potential of *AI planning* [Kambhampati, 2007].

The Machine Learning of planning action models is a promising alternative to hand-coding them and nowadays, there exist sophisticated algorithms like AMAN [Zhuo and Kambhampati, 2013], ARMS [Yang *et al.*, 2007], LOCM [Cresswell *et al.*, 2013] or SLAF [Amir and Chang, 2008]. Motivated by recent advances on the synthesis of different kinds of generative models with classical planning [Bonet *et al.*, 2009; Segovia-Aguas *et al.*, 2016; 2017], this paper presents a novel approach for learning STRIPS action models that introduces the following contributions:

1. Is defined as a classical planning compilation. This fact opens the door to the *bootstrapping* of planning action models and allows us to report results over a wide range of planning domains.
2. Does not require observing the particular executed actions. An off-the-shelf classical planner determines these actions given the state observations.
3. Is extensible to semantically evaluate the quality of learned STRIPS models as well as to estimate the prob-

ability distribution of the possible STRIPS models given a sequence of state observations.

2 Background

This section defines the planning models used in this work as well as the output (an STRIPS action model) of the addressed learning task.

2.1 Classical planning

We use F to denote the set of *fluents* (propositional variables) describing a state. A *literal* l is a valuation of a fluent $f \in F$, i.e. either $l = f$ or $l = \neg f$. A set of literals L represents a partial assignment of values to fluents (WLOG we assume that L does not assign conflicting values to any fluent). We use $\mathcal{L}(F)$ to denote the set of all literal sets on F , i.e. all partial assignments of values to fluents.

A *state* s is a full assignment of values to fluents, i.e. $|s| = |F|$, so the size of the state space is $2^{|F|}$. Explicitly including negative literals $\neg f$ in states simplifies subsequent definitions but often, we will abuse notation by defining a state s only in terms of the fluents that are true in s , as is common in STRIPS planning.

A *classical planning frame* is a tuple $\Phi = \langle F, A \rangle$, where F is a set of fluents and A is a set of actions. Each action $a \in A$ comprises three sets of literals:

- $\text{pre}(a) \subseteq \mathcal{L}(F)$, called *preconditions*, the literals that must hold for the action $a \in A$ to be applicable.
- $\text{eff}^+(a) \subseteq \mathcal{L}(F)$, called *positive effects*, that defines the fluents set to true by the application of the action $a \in A$.
- $\text{eff}^-(a) \subseteq \mathcal{L}(F)$, called *negative effects*, that defines the fluents set to false by the action application.

We say that an action $a \in A$ is *applicable* in a state s iff $\text{pre}(a) \subseteq s$. The result of applying a in s is the *successor state* denoted by $\theta(s, a) = \{s \setminus \text{eff}^-(a)\} \cup \text{eff}^+(a)$.

A *classical planning problem* is a tuple $P = \langle F, A, I, G \rangle$, where I is an initial state and $G \subseteq \mathcal{L}(F)$ is a goal condition. A *plan* for P is an action sequence $\pi = \langle a_1, \dots, a_n \rangle$ that induces the *state trajectory* $\langle s_0, s_1, \dots, s_n \rangle$ such that $s_0 = I$ and, for each $1 \leq i \leq n$, a_i is applicable in s_{i-1} and generates the successor state $s_i = \theta(s_{i-1}, a_i)$. The *plan length* is denoted with $|\pi| = n$. A plan π *solves* P iff $G \subseteq s_n$, i.e. if the goal condition is satisfied at the last state reached after following the application of the plan π in the initial state I .

Classical planning with conditional effects

Our approach for learning STRIPS action models is compiling the learning task into a classical planning task with conditional effects.

An action $a \in A$ with conditional effects is defined as a set of *preconditions* $\text{pre}(a) \in \mathcal{L}(F)$ and a set of *conditional effects* $\text{cond}(a)$. Each conditional effect $C \triangleright E \in \text{cond}(a)$ is composed of two sets of literals $C \in \mathcal{L}(F)$, the *condition*, and $E \in \mathcal{L}(F)$, the *effect*. An action $a \in A$ is *applicable* in a state s if and only if $\text{pre}(a) \subseteq s$, and the *triggered effects* resulting from the action application are the effects whose conditions hold in s :

$$\text{triggered}(s, a) = \bigcup_{C \triangleright E \in \text{cond}(a), C \subseteq s} E,$$

The result of applying action a in state s is the *successor* state $\theta(s, a) = \{s \setminus \text{eff}_c^-(s, a)\} \cup \text{eff}_c^+(s, a)$ where $\text{eff}_c^-(s, a) \subseteq \text{triggered}(s, a)$ and $\text{eff}_c^+(s, a) \subseteq \text{triggered}(s, a)$ are, respectively, the triggered *negative* and *positive* effects.

2.2 STRIPS action schemes

This work addresses the learning of PDDL action schemes that follow the STRIPS requirement [McDermott *et al.*, 1998; Fox and Long, 2003]. Figure 1 shows the *stack* action schema, coded in PDDL, from a four-operator *blocksworld* [Slaney and Thiébaux, 2001].

To formalize the output of the learning task, we assume that fluents F are instantiated from a set of *predicates* Ψ , as in PDDL. Each predicate $p \in \Psi$ has an argument list of arity $\text{ar}(p)$. Given a set of *objects* Ω , the set of fluents F is induced by assigning objects in Ω to the arguments of predicates in Ψ , i.e. $F = \{p(\omega) : p \in \Psi, \omega \in \Omega^{\text{ar}(p)}\}$ s.t. Ω^k is the k -th Cartesian power of Ω .

Let $\Omega_v = \{v_i\}_{i=1}^{\max_{a \in A} \text{ar}(a)}$ be a new set of objects ($\Omega \cap \Omega_v = \emptyset$), denoted as *variable names*, and that is bound by the maximum arity of an action in a given planning frame. For instance, in a three-block *blocksworld* $\Omega = \{\text{block}_1, \text{block}_2, \text{block}_3\}$ while $\Omega_v = \{v_1, v_2\}$ because the operators with the maximum arity, *stack* and *unstack*, have arity two.

Let us also define F_v , a new set of fluents $F \cap F_v = \emptyset$, that results from instantiating Ψ using only the objects in Ω_v and that defines the elements that can appear in an action schema. For the *blocksworld*, $F_v = \{\text{handempty}, \text{holding}(v_1), \text{holding}(v_2), \text{clear}(v_1), \text{clear}(v_2), \text{ontable}(v_1), \text{ontable}(v_2), \text{on}(v_1, v_1), \text{on}(v_1, v_2), \text{on}(v_2, v_1), \text{on}(v_2, v_2)\}$.

We assume also that actions $a \in A$ are instantiated from STRIPS operator schemes $\xi = \langle \text{head}(\xi), \text{pre}(\xi), \text{add}(\xi), \text{del}(\xi) \rangle$ where:

- $\text{head}(\xi) = \langle \text{name}(\xi), \text{pars}(\xi) \rangle$, is the operator *header* defined by its name and the corresponding *variable names*, $\text{pars}(\xi) = \{v_i\}_{i=1}^{\text{ar}(\xi)}$. The headers of a four-operator *blocksworld* are $\text{pickup}(v_1)$, $\text{putdown}(v_1)$, $\text{stack}(v_1, v_2)$ and $\text{unstack}(v_1, v_2)$.
- The preconditions $\text{pre}(\xi) \subseteq F_v$, the negative effects $\text{del}(\xi) \subseteq F_v$, and the positive effects $\text{add}(\xi) \subseteq F_v$

```
(:action stack
:parameters (?v1 ?v2 - object)
:precondition (and (holding ?v1) (clear ?v2))
:effect (and (not (holding ?v1))
(not (clear ?v2))
(handempty) (clear ?v1)
(on ?v1 ?v2)))
```

Figure 1: STRIPS operator schema coding, in PDDL, the *stack* action from the *blocksworld*.

such that, $\text{del}(\xi) \subseteq \text{pre}(\xi)$, $\text{del}(\xi) \cap \text{add}(\xi) = \emptyset$ and $\text{pre}(\xi) \cap \text{add}(\xi) = \emptyset$.

Finally we also define $F_v(\xi) \subseteq F_v$ as the subset of elements that can appear in a given action schema ξ . For instance, for the *stack* action schema $F_v(\text{stack}) = F_v$ while $F_v(\text{pickup}) = \{\text{handempty}, \text{holding}(v_1), \text{clear}(v_1), \text{ontable}(v_1), \text{on}(v_1, v_1)\}$ only contains the fluents from F_v that do not involve v_2 because the action header contained the single parameter v_1 .

3 Learning STRIPS action models

Learning STRIPS action models from fully available input knowledge, i.e. from plans where every action in the plan is available as well as its corresponding *pre-* and *post-states*, is straightforward. In this case, STRIPS operator schemes are derived lifting the literals that change between the pre and post-state of the corresponding action executions. Preconditions are derived lifting the minimal set of literals that appears in all the pre-states of the corresponding actions.

In this paper we address a more challenging learning task, where less input knowledge is available. The addressed learning task corresponds to observing an agent acting in the world but watching only the results of its plan executions, the actual executed actions are unobserved. This learning task is defined as $\Lambda = \langle \Psi, \Xi, \mathcal{O} \rangle$, where:

- Ψ is the set of predicates that define the abstract state space of a given classical planning frame.
- Ξ is the set of empty operator schemes containing only the headers with the name and parameter of each operator schema $\xi \in \Xi$.
- $\mathcal{O} = \langle s_0, s_1, \dots, s_n \rangle$ is a sequence of *state observations* obtained observing the execution of an *unobserved* plan $\pi = \langle a_1, \dots, a_n \rangle$.

A solution to Λ is a set of operator schema Ξ' compliant with the predicates in Ψ , the headers in Ξ and the given sequence of state observations \mathcal{O} . A planning compilation is a suitable approach for addressing a Λ learning task because a solution must not only determine the STRIPS action model Ξ' but also, the *unobserved* plan $\pi = \langle a_1, \dots, a_n \rangle$, that explains \mathcal{O} . Figure 2 shows an example of a Λ task for learning a STRIPS action model in the *blocksworld* from the sequence of five state observations that corresponds to inverting a 2-blocks tower.

3.1 Learning with classical planning

Our approach for addressing the learning task, is compiling it into a classical planning task with conditional effects. The

```

;;; Predicates in  $\Psi$ 
(handempty) (holding ?o - object)
(clear ?o - object) (ontable ?o - object)
(on ?o1 - object ?o2 - object)

;;; Headers in  $\Xi$ 
(pickup v1) (putdown v1)
(stack v1 v2) (unstack v1 v2)

;;; Observations in  $\mathcal{O}$ 
;;; observation #0
(clear block2) (on block2 block1)
(ontable block1) (handempty)

;;; observation #1
(holding block2) (clear block1)
(ontable block1)

;;; observation #2
(clear block1) (ontable block1)
(clear block2) (ontable block2)
(handempty)

;;; observation #3
(holding block1) (clear block2)
(ontable block2)

;;; observation #4
(clear block1) (on block1 block2)
(ontable block2) (handempty)

```

Figure 2: Example of a Λ task for learning a STRIPS action model in the *blocksworld* from a sequence of five state observations.

intuition behind the compilation is that a solution to the resulting classical planning task is a sequence of actions that:

1. Programs the STRIPS action model Ξ' . A solution plan starts with a *prefix* that, for each $\xi \in \Xi$, determines which fluents $f \in F_v(\xi)$ belong to its $pre(\xi)$, $del(\xi)$ and $add(\xi)$ sets.
2. Validates the STRIPS action model Ξ' in \mathcal{O} . The solution plan continues with a *postfix* that produces the given sequence of states $\langle s_0, s_1, \dots, s_n \rangle$ using the programmed action model Ξ' .

Given a learning task $\Lambda = \langle \Psi, \Xi, \mathcal{O} \rangle$ the compilation outputs a classical planning task $P_\Lambda = \langle F_\Lambda, A_\Lambda, I_\Lambda, G_\Lambda \rangle$:

- F_Λ contains:
 - The set of fluents F that is built instantiating the predicates Ψ with the objects appearing in the input observations \mathcal{O} .
 - Fluents representing the programmed action model $pre_f(\xi)$, $del_f(\xi)$ and $add_f(\xi)$, for every $f \in F_v(\xi)$. If a fluent $pre_f(\xi)/del_f(\xi)/add_f(\xi)$ holds, it means that f is a precondition/negative effect/positive effect in the STRIPS operator schema $\xi \in \Xi$. For instance, the preconditions of the *stack* schema (Figure 1) are represented by the fluents `pre.holding_stack.v1` and `pre.clear_stack.v2` set to *True*.

- Fluent $mode_{prog}$ indicating whether the operator schemes are programmed or validated (already programmed) and fluents $\{test_i\}_{1 \leq i \leq n}$, indicating the observation where the action model is validated.

- I_Λ contains the fluents from F that encode s_0 (the first observation) and $mode_{prog}$ set to true. In addition, our compilation assumes that initially operator schemes are programmed with every possible precondition, no negative effect and no positive effect. With this regard, the fluents $pre_f(\xi)$ hold at the initial state for every $f \in F_v(\xi)$.
- $G_\Lambda = \bigcup_{1 \leq i \leq n} \{test_i\}$, indicates that the programmed action model is validated in all the input observations.
- A_Λ comprises three kinds of actions:

1. Actions for *programming* operator schema $\xi \in \Xi$:
 - Actions for **removing** a *precondition* $f \in F_v(\xi)$ from the action schema $\xi \in \Xi$.

$$\begin{aligned}
pre(\text{programPre}_{f,\xi}) &= \{\neg del_f(\xi), \neg add_f(\xi), \\
&\quad mode_{prog}, pre_f(\xi)\}, \\
cond(\text{programPre}_{f,\xi}) &= \{\emptyset\} \triangleright \{\neg pre_f(\xi)\}.
\end{aligned}$$

- Actions for **adding** a *negative* or *positive* effect $f \in F_v(\xi)$ to the action schema $\xi \in \Xi$.

$$\begin{aligned}
pre(\text{programEff}_{f,\xi}) &= \{\neg del_f(\xi), \neg add_f(\xi), \\
&\quad mode_{prog}\}, \\
cond(\text{programEff}_{f,\xi}) &= \{pre_f(\xi)\} \triangleright \{del_f(\xi)\}, \\
&\quad \{\neg pre_f(\xi)\} \triangleright \{add_f(\xi)\}.
\end{aligned}$$

2. Actions for *applying* an already programmed operator schema $\xi \in \Xi$ bound with the objects $\omega \subseteq \Omega^{ar}(\xi)$. We assume that the operators headers are known so the binding of the operator schema is done implicitly by order of appearance of the action parameters, i.e. variables $pars(\xi)$ are bound to the objects in ω appearing at the same position. Figure 3 shows the PDDL encoding of the action for applying a programmed operator *stack*.

$$\begin{aligned}
pre(\text{apply}_{\xi,\omega}) &= \{pre_f(\xi) \implies p(\omega)\}_{\forall p \in \Psi, f = p(pars(\xi))}, \\
cond(\text{apply}_{\xi,\omega}) &= \{del_f(\xi)\} \triangleright \{\neg p(\omega)\}_{\forall p \in \Psi, f = p(pars(\xi))}, \\
&\quad \{add_f(\xi)\} \triangleright \{p(\omega)\}_{\forall p \in \Psi, f = p(pars(\xi))}, \\
&\quad \{mode_{prog}\} \triangleright \{\neg mode_{prog}\}.
\end{aligned}$$

3. Actions for *validating* an observation $1 \leq i \leq n$.

$$\begin{aligned}
pre(\text{validate}_i) &= s_i \cup \{test_j\}_{j \in 1 \leq j < i} \\
&\quad \cup \{\neg test_j\}_{j \in i \leq j \leq n} \cup \{\neg mode_{prog}\}, \\
cond(\text{validate}_i) &= \{\emptyset\} \triangleright \{test_i\}.
\end{aligned}$$

3.2 Compilation properties

Lemma 1. *Soundness.* Any classical plan π that solves P_Λ induces an action model Ξ' that solves the learning task Λ .

```

(:action apply_stack
:parameters (?o1 - object ?o2 - object)
:precondition
  (and (or (not (pre_on_stack_v1_v1)) (on ?o1 ?o1))
        (or (not (pre_on_stack_v1_v2)) (on ?o1 ?o2))
        (or (not (pre_on_stack_v2_v1)) (on ?o2 ?o1))
        (or (not (pre_on_stack_v2_v2)) (on ?o2 ?o2))
        (or (not (pre_ontable_stack_v1)) (ontable ?o1))
        (or (not (pre_ontable_stack_v2)) (ontable ?o2))
        (or (not (pre_clear_stack_v1)) (clear ?o1))
        (or (not (pre_clear_stack_v2)) (clear ?o2))
        (or (not (pre_holding_stack_v1)) (holding ?o1))
        (or (not (pre_holding_stack_v2)) (holding ?o2))
        (or (not (pre_handempty_stack)) (handempty)))
:effect
  (and (when (del_on_stack_v1_v1) (not (on ?o1 ?o1)))
        (when (del_on_stack_v1_v2) (not (on ?o1 ?o2)))
        (when (del_on_stack_v2_v1) (not (on ?o2 ?o1)))
        (when (del_on_stack_v2_v2) (not (on ?o2 ?o2)))
        (when (del_ontable_stack_v1) (not (ontable ?o1)))
        (when (del_ontable_stack_v2) (not (ontable ?o2)))
        (when (del_clear_stack_v1) (not (clear ?o1)))
        (when (del_clear_stack_v2) (not (clear ?o2)))
        (when (del_holding_stack_v1) (not (holding ?o1)))
        (when (del_holding_stack_v2) (not (holding ?o2)))
        (when (del_handempty_stack) (not (handempty)))
        (when (add_on_stack_v1_v1) (on ?o1 ?o1))
        (when (add_on_stack_v1_v2) (on ?o1 ?o2))
        (when (add_on_stack_v2_v1) (on ?o2 ?o1))
        (when (add_on_stack_v2_v2) (on ?o2 ?o2))
        (when (add_ontable_stack_v1) (ontable ?o1))
        (when (add_ontable_stack_v2) (ontable ?o2))
        (when (add_clear_stack_v1) (clear ?o1))
        (when (add_clear_stack_v2) (clear ?o2))
        (when (add_holding_stack_v1) (holding ?o1))
        (when (add_holding_stack_v2) (holding ?o2))
        (when (add_handempty_stack) (handempty))
        (when (modeProg) (not (modeProg)))))

```

Figure 3: Action for applying an already programmed schema *stack* as encoded in PDDL (implications coded as disjunctions).

Proof sketch. The compilation forces that once the preconditions of an operator schema $\xi \in \Xi'$ are programmed, they cannot be altered. The same happens with the positive and negative effects. Furthermore because of the preconditions of the program $\text{Pre}_{f,\xi}$ actions, effects are only programmable after preconditions are programmed. Once the operator schemes Ξ' are programmed, they can only be applied because of the *modeProg* fluent. To solve P_Λ , the goals $\{test_i\}$, $1 \leq i \leq n$ can only be achieved: executing an applicable sequence of programmed operator schemes that reaches every state $s_i \in \mathcal{O}$, starting from s_0 and following the sequence $1 \leq i \leq n$. Therefore if $test_n$ is achieved, it means that the programmed action model Ξ' is compliant with the provided input knowledge and hence, solves Λ . \square

Lemma 2. Completeness. Any STRIPS action model Ξ' that solves a $\Lambda = \langle \Psi, \Xi, \mathcal{O} \rangle$ learning task, is computable solving the corresponding classical planning task P_Λ .

Proof sketch. By definition, $F_v(\xi) \subseteq F_\Lambda$ fully captures the set of elements that can appear in a STRIPS action schema $\xi \in \Xi$. In addition, any possible STRIPS action schema Ξ' that can be built with the fluents in F_v can be computed with the P_Λ compilation. The only STRIPS action models that the compilation cannot compute are the ones that violate the state constraints defined by the \mathcal{O} sequence, which constrain as well the solutions to the Λ learning task. \square

4 Recognizing STRIPS action models

Inspired by *plan recognition as planning* [Ramirez and Geffner, 2009], we show that the previous compilation can be extended to (1) semantically evaluate the quality of learned STRIPS models, without comparing with a reference model and (2), address STRIPS action model recognition tasks.

4.1 The STRIPS edit distance

We assume that how well a given STRIPS action model Ξ explains a given sequence of observations \mathcal{O} , depends on the amount of *edition* that one has to introduce to Ξ to produce the sequence of observations \mathcal{O} . In the extreme, if the given model Ξ perfectly explains the observations \mathcal{O} , no *edit* to the model must be introduced.

Let us define the operations allowed to edit a given STRIPS action model. We call these operations the STRIPS *edit operations*. With the aim of keeping tractable the branching factor of the classical planning task that results from our compilation, we only define two STRIPS *edit operations*:

- *Deletion.* A fluent $pre_f(\xi)/del_f(\xi)/add_f(\xi)$, such that $f \in F_v(\xi)$, is removed from the operator schema $\xi \in \Xi$.
- *Insertion.* A fluent $pre_f(\xi)/del_f(\xi)/add_f(\xi)$, such that $f \in F_v(\xi)$, is added to the operator schema $\xi \in \Xi$.

With this defined, we can now formalize an edit distance metric that assesses how dissimilar two given STRIPS action models are. This distance satisfies the *metric axioms* provided that the two STRIPS *edit operations* have the same positive cost.

Definition 3. Given two STRIPS action model Ξ and Ξ' , such that both are built from the same set of possible elements F_v . The **STRIPS edit distance**, denoted as $\delta(\Xi, \Xi')$, is the minimum number of STRIPS edit operations to transform Ξ into Ξ' .

Since the size of the F_v set is bound, the maximum number of edits that can be introduced to a given STRIPS action model is bound as well. In more detail, for a given operator schema $\xi \in \Xi$ the maximum number of edits that can be introduced to their precondition set is $|F_v(\xi)|$. With regard to effects of ξ , the maximum number of edits that can be introduced are two times $|F_v(\xi)|$.

Definition 4. Given a STRIPS action model Ξ built from the set of possible elements F_v . The **maximum STRIPS edit distance** is the upperbound $\delta(\Xi, *) = \sum_{\xi \in \Xi} 3|F_v(\xi)|$, that indicates the maximum value of the STRIPS edit distance from Ξ to any STRIPS action model definable within F_v .

Likewise the STRIPS *edit distance* can be defined with regard to a sequence of observations \mathcal{O} .

Definition 5. Given a STRIPS action model Ξ built from the set of possible elements F_v and a sequence of observations \mathcal{O} . The **observations STRIPS edit distance**, denoted by $\delta(\Xi, \mathcal{O})$, is the minimum number of STRIPS edit operations to transform Ξ into an action model Ξ' defined within F_v that can produce a plan $\pi = \langle a_1, \dots, a_n \rangle$ such that π induces the state observations $\mathcal{O} = \langle s_0, s_1, \dots, s_n \rangle$.

The *observations STRIPS edit distance* allow us to semantically evaluate the quality of a learned strips action model with respect to a sequence of observations that acts as a test set. The lower this distance is, the better the quality of the learned model. The semantic nature of this evaluation is robust to learning episodes where actions are reformulated and still compliant with the inputs (e.g. the *blocksworld* operator *stack* could be *learned* with the preconditions and effects of the *unstack* operator and vice versa or the roles of actions parameters with the same type could be interchanged).

4.2 Recognition of STRIPS models

Given a set of possible STRIPS models and a sequence of state observations the *recognition of STRIPS models* is the task of computing the probability distribution of these models according to the given observations. Here we show that the *observations STRIPS edit distance* is helpful to compute the probability distribution of the possible STRIPS models given a sequence of state observations. The main idea, taken from *plan recognition as planning* [Ramirez and Geffner, 2009], is to map these distances into likelihoods using the Bayes rule.

According to the Bayes rule, the probability of an hypothesis \mathcal{H} given the observations \mathcal{O} can be computed with $P(\mathcal{H}|\mathcal{O}) = \frac{P(\mathcal{O}|\mathcal{H})P(\mathcal{H})}{P(\mathcal{O})}$. In our scenario, the hypotheses are about the possible STRIPS action models that can be build given a set of operator headers, or in other words, given the $F_v(\xi)$ sets. Moreover, we assume that all the possible STRIPS action models are equiprobable so according to the Bayes rule $P(\mathcal{H}|\mathcal{O})$ is proportional to $P(\mathcal{O}|\mathcal{H})$. Finally, we assume that $P(\mathcal{O}|\mathcal{H})$ is given by the *observations STRIPS edit distance*, mapping distances into probabilities (the larger the distance the lower the likelihood), according to the following sigmoid function $P(\mathcal{O}|\Xi) = \frac{1}{1+e^{-\delta(\Xi,\mathcal{O})}}$.

With this regard, the probability distribution of the possible STRIPS models (within the set of operator headers, that is within the $F_v(\xi)$ sets) given a sequence of state observations, $P(\Xi|\mathcal{O})$, can be computed by:

1. The *observations STRIPS edit distance* is computed for every possible model.
2. The resulting distances are plugged into the previous $P(\Xi|\mathcal{O})$ equation to map distances into likelihoods.
3. Likelihoods are plugged into the Bayes' rule $P(\Xi|\mathcal{O})$ to obtain the normalized posterior probabilities. Since we are computing a probability distribution these probabilities must sum up 1.0.

4.3 STRIPS model evaluation and recognition with classical planning

The previous compilation can be extended to estimate the *observations STRIPS edit distance*. In this case the tuple $\Lambda = \langle \Psi, \Xi, \mathcal{O} \rangle$ does not represent a learning task but the edition of a given STRIPS action model Ξ to cover the sequence of observations \mathcal{O} . The extended compilation outputs a classical planning task $P'_\Lambda = \langle F_\Lambda, A'_\Lambda, I'_\Lambda, G_\Lambda \rangle$ where:

- F_Λ is defined as in the previous compilation.

- I'_Λ contains the fluents from F that encode s_0 (the first observation) and $mode_{prog}$ set to true. In addition, the given STRIPS action model Ξ is now encoded in the initial state. This means that fluents $pref_f(\xi)/del_f(\xi)/add_f(\xi)$ with $f \in F_v(\xi)$ hold in the initial state if they appear in the given STRIPS action model Ξ .
- G_Λ is defined as in the previous compilation.
- A'_Λ , comprises the three kinds of actions of A_Λ . The Actions for *applying* an already programmed operator schema and the actions for *validating* an observation $1 \leq i \leq n$ are defined exactly as in the previous compilation. The only difference are the actions for *programming* operator schema $\xi \in \Xi$ that must include now actions for adding a precondition and for removing a *negative* or *positive* effect.

The probability distribution $P(\Xi|\mathcal{O})$ would be exact if the classical planning problems resulting from our compilation are solved optimally (according to the number of edit actions), and will be approximate if they are solved with a satisfying classical planner or furthermore if what is solved is not P'_Λ but a relaxation of this task such as the *delete relaxation*.

5 Evaluation

This section evaluates our approach for learning STRIPS models starting from different amounts of input knowledge.

Reproducibility

We used IPC domains that satisfy the STRIPS requirement [Fox and Long, 2003], taken from the PLANING.DOMAINS repository [Muise, 2016]. For the learning of the STRIPS action models we used observations sequences of 25 states per domain. All experiments are run on an Intel Core i5 3.10 GHz x 4 with 4 GB of RAM.

MADAGASCAR is the classical planner we use to solve the instances that result from our compilations because its ability to deal with dead-ends [Rintanen, 2014]. In addition, MADAGASCAR can apply the actions for programming preconditions in a single planning step (in parallel) because these actions do not interact. Actions for programming action effects can also be applied in a single planning step reducing significantly the planning horizon.

The compilation source code, the evaluation scripts and the benchmarks are fully available at this anonymous repository <https://github.com/anonsub/strips-learning> so any experimental data reported in the paper can be reproduced.

Supervised evaluation of the learned models

For each domain the learned model is compared with the actual model and its quality is quantified with the *precision* and *recall* metrics. Precision gives a notion of *soundness* while recall gives a notion of the *completeness* of the learned models. $Precision = \frac{tp}{tp+fp}$, where tp is the number of true positives (predicates that correctly appear in the action model) and fp is the number of false positives (predicates appear in the learned action model that should not appear). $Recall = \frac{tp}{tp+fn}$ where fn is the number of false negatives

(predicates that should appear in the learned action model but are missing).

When the learning hypothesis space is low constrained, the learned actions can be reformulated and still be compliant with the inputs. For instance in the *blocksworld*, operator *stack* could be *learned* with the preconditions and effects of the *unstack* operator (and vice versa). Furthermore, in a given action the role of the parameters that share the same type can be interchanged making non trivial to compute *precision* and *recall* with respect to a reference model.

To address these issues we defined an evaluation method robust to action reformulation. Precision and recall are often combined using the *harmonic mean*. This expression is called the *F-measure* (or the balanced *F-score*) and is formally defined as $F = 2 \times \frac{Precision \times Recall}{Precision + Recall}$. Given a reference STRIPS action model Ξ^* and the learned STRIPS action model Ξ we define the bijective function $f_{P\&R} : \Xi \mapsto \Xi^*$ such that $f_{P\&R}$ maximizes the accumulated *F-measure*. With this mapping defined we can compute the *precision* and *recall* of a learned STRIPS action $\xi \in \Xi$ with respect to the action $f_{P\&R}(\xi) \in \Xi^*$ even if actions are reformulated in the learning process.

6 Conclusions

As far as we know, this is the first work on learning STRIPS action models from state observations, exclusively using classical planning and evaluated over a wide range of different domains. Recently, Stern and Juba 2017 proposed a classical planning compilation for learning action models but following the *finite domain* representation for the state variables and did not report experimental results since the compilation was not implemented.

The size of the compiled classical planning instances depends on the number of input examples. Generating *informative* examples for learning planning action models is an open issue. Planning actions include preconditions that are only satisfied by specific sequences of actions, often, with a low probability of being chosen by chance [Fern *et al.*, 2004]. The success of recent algorithms for exploring planning tasks [Francés *et al.*, 2017] motivates the development of novel techniques able to autonomously collect informative learning examples. The combination of such exploration techniques with our learning approach is an intriguing research direction that opens the door to the bootstrapping of planning action models.

References

- [Amir and Chang, 2008] Eyal Amir and Allen Chang. Learning partially observable deterministic action models. *Journal of Artificial Intelligence Research*, 33:349–402, 2008.
- [Bonet *et al.*, 2009] Blai Bonet, Héctor Palacios, and Héctor Geffner. Automatic derivation of memoryless policies and finite-state controllers using classical planners. In *ICAPS*, 2009.
- [Cresswell *et al.*, 2013] Stephen N Cresswell, Thomas L McCluskey, and Margaret M West. Acquiring planning

	Pre		Add		Del		P	R
	P	R	P	R	P	R		
blocks	0.57	0.44	0.63	0.56	0.57	0.44	0.59	0.48
driverlog	0.2	0.07	0.18	0.29	0.2	0.14	0.19	0.17
ferry	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
floor-tile	0.5	0.14	0.75	0.55	0.6	0.27	0.62	0.32
Grid	-	-	-	-	-	-	-	-
gripper-strips	0.25	0.17	0.25	0.25	0.25	0.25	0.25	0.22
hanoi	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
hiking	-	-	-	-	-	-	-	-
n-puzzle	-	-	-	-	-	-	-	-
parking	0.6	0.21	0.14	0.11	0.4	0.22	0.38	0.18
satellite	0.25	0.07	0.33	0.4	0.67	0.5	0.42	0.32
Sokoban	-	-	-	-	-	-	-	-
transport	0.33	0.1	0.33	0.4	0.0	0.0	0.22	0.17
zeno-travel	0.25	0.07	0.17	0.14	0.0	0.0	0.14	0.07
	-	-	-	-	-	-	-	-

Table 1: Precision and recall values obtained when learning from labels without computing the $f_{P\&R}$ mapping.

	Pre		Add		Del		P	R
	P	R	P	R	P	R		
blocks	0.71	0.56	0.75	0.67	0.71	0.56	0.73	0.59
driverlog	0.6	0.21	0.27	0.43	0.6	0.43	0.49	0.36
ferry	0.5	0.29	0.75	0.75	0.5	0.5	0.58	0.51
floor-tile	0.5	0.14	0.75	0.55	0.6	0.27	0.62	0.32
Grid	-	-	-	-	-	-	-	-
gripper-strips	1.0	0.67	1.0	1.0	1.0	1.0	1.0	0.89
hanoi	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
hiking	-	-	-	-	-	-	-	-
n-puzzle	-	-	-	-	-	-	-	-
parking	0.6	0.21	0.14	0.11	0.4	0.22	0.38	0.18
satellite	0.25	0.07	0.33	0.4	0.67	0.5	0.42	0.32
Sokoban	-	-	-	-	-	-	-	-
transport	1.0	0.3	0.67	0.8	0.67	0.4	0.78	0.5
zeno-travel	0.5	0.14	0.5	0.43	0.33	0.14	0.44	0.24
	-	-	-	-	-	-	-	-

Table 2: Precision and recall values obtained when learning from labels but computing the $f_{P\&R}$ mapping.

	Total time	Preprocess	Plan length
Blocks	1.40	0.00	70
Driverlog	1.50	0.00	89
Ferry	1.49	0.00	64
Floortile	351.38	0.11	156
Grid	-	-	-
Gripper	0.04	0.00	59
Hanoi	2.33	0.01	49
Hiking	-	-	-
Parking	-	-	-
Satellite	78.36	0.03	98
Sokoban	-	-	-
Transport	285.61	0.10	106
Zenotravel	6.20	0.46	71

Table 3: Planning results obtained when learning from labels.

- domain models using LOCM. *The Knowledge Engineering Review*, 28(02):195–213, 2013.
- [Fern *et al.*, 2004] Alan Fern, Sung Wook Yoon, and Robert Givan. Learning domain-specific control knowledge from random walks. In *ICAPS*, pages 191–199, 2004.
- [Fox and Long, 2003] Maria Fox and Derek Long. PDDL2.1: An extension to PDDL for expressing temporal planning domains. *J. Artif. Intell. Res.(JAIR)*, 20:61–124, 2003.
- [Francés *et al.*, 2017] Guillem Francés, Miquel Ramírez, Nir Lipovetzky, and Hector Geffner. Purely declarative action representations are overrated: Classical planning with simulators. In *International Joint Conference on Artificial Intelligence*, 2017.
- [Geffner and Bonet, 2013] Hector Geffner and Blai Bonet. A concise introduction to models and methods for automated planning, 2013.
- [Ghallab *et al.*, 2004] Malik Ghallab, Dana Nau, and Paolo Traverso. *Automated Planning: theory and practice*. Elsevier, 2004.
- [Kambhampati, 2007] Subbarao Kambhampati. Model-lite planning for the web age masses: The challenges of planning with incomplete and evolving domain models. In *Proceedings of the National Conference on Artificial Intelligence*, 2007.
- [McDermott *et al.*, 1998] Drew McDermott, Malik Ghallab, Adele Howe, Craig Knoblock, Ashwin Ram, Manuela Veloso, Daniel Weld, and David Wilkins. PDDL – The Planning Domain Definition Language, 1998.
- [Muise, 2016] Christian Muise. Planning. domains. *ICAPS system demonstration*, 2016.
- [Ramírez and Geffner, 2009] Miquel Ramírez and Hector Geffner. Plan recognition as planning. In *Proceedings of the 21st international joint conference on Artificial intelligence. Morgan Kaufmann Publishers Inc*, pages 1778–1783, 2009.
- [Ramírez, 2012] Miquel Ramírez. *Plan recognition as planning*. PhD thesis, Universitat Pompeu Fabra, 2012.
- [Rintanen, 2014] Jussi Rintanen. Madagascar: Scalable planning with sat. *Proceedings of the 8th International Planning Competition (IPC-2014)*, 2014.
- [Segovia-Aguas *et al.*, 2016] Javier Segovia-Aguas, Sergio Jiménez, and Anders Jonsson. Hierarchical finite state controllers for generalized planning. In *Proceedings of the Twenty-Fifth International Joint Conference on Artificial Intelligence*, pages 3235–3241. AAAI Press, 2016.
- [Segovia-Aguas *et al.*, 2017] Javier Segovia-Aguas, Sergio Jiménez, and Anders Jonsson. Generating context-free grammars using classical planning. In *International Joint Conference on Artificial Intelligence*, 2017.
- [Slaney and Thiébaux, 2001] John Slaney and Sylvie Thiébaux. Blocks world revisited. *Artificial Intelligence*, 125(1-2):119–153, 2001.
- [Stern and Juba, 2017] Roni Stern and Brendan Juba. Efficient, safe, and probably approximately complete learning of action models. *IJCAI*, 2017.
- [Yang *et al.*, 2007] Qiang Yang, Kangheng Wu, and Yunfei Jiang. Learning action models from plan examples using weighted max-sat. *Artificial Intelligence*, 171(2-3):107–143, 2007.
- [Zhuo and Kambhampati, 2013] Hankz Hankui Zhuo and Subbarao Kambhampati. Action-model acquisition from noisy plan traces. In *IJCAI*, pages 2444–2450, 2013.