

# Learning and Recognition of STRIPS Action Models from State Observations

Diego Aineto and Sergio Jiménez and Eva Onaindia

Departamento de Sistemas Informáticos y Computación

Universitat Politècnica de València.

Camino de Vera s/n. 46022 Valencia, Spain

{dieaigar,serjice,onaindia}@dsic.upv.es

## Abstract

This paper presents a classical planning compilation for learning STRIPS action models from state observations. The compilation approach does not require observing the precise actions that produced the state observations because such actions are determined by an off-the-shelf classical planner. In addition, the compilation is extensible to estimate the probability distribution of the possible STRIPS models given a sequence of state observations. This extension allows us to address model recognition tasks as well as to semantically evaluate the quality of learned STRIPS models.

## 1 Introduction

Besides *plan synthesis* [Ghallab *et al.*, 2004], planning action models are also useful for *plan/goal recognition* [Ramírez, 2012]. At these planning tasks, automated planners are required to reason about an action model that correctly and completely captures the possible world transitions [Geffner and Bonet, 2013]. Unfortunately, building planning action models is complex, even for planning experts, and this knowledge acquisition task is a bottleneck that limits the potential of *AI planning* [Kambhampati, 2007].

The Machine Learning of planning action models is a promising alternative to hand-coding them and nowadays, there exist sophisticated algorithms like AMAN [Zhuo and Kambhampati, 2013], ARMS [Yang *et al.*, 2007], LOCM [Cresswell *et al.*, 2013] or SLAF [Amir and Chang, 2008]. Motivated by recent advances on the synthesis of different kinds of generative models with classical planning [Bonet *et al.*, 2009; Segovia-Aguas *et al.*, 2016; 2017], this paper presents a novel approach for learning STRIPS action models that introduces the following contributions:

1. Is defined as a planning compilation. This fact opens the door to the *bootstrapping* of planning action models and allows us to report results over a wide range of planning domains.
2. Does not require observing the particular executed actions. An off-the-shelf classical planner determines these actions given the state observations.

3. Can estimate the probability distribution of the possible STRIPS models, given an observation sequence. This extension allow us to address model recognition tasks as well as to evaluate the quality of the learned models semantically, without comparing with a reference model.

## 2 Background

This section defines the planning models used in this work as well as the input (states and state invariants) and the output (an STRIPS action model) of the addressed learning task.

### 2.1 Classical planning

We use  $F$  to denote the set of *fluents* (propositional variables) describing a state. A *literal*  $l$  is a valuation of a fluent  $f \in F$ , i.e. either  $l = f$  or  $l = \neg f$ . A set of literals  $L$  represents a partial assignment of values to fluents (WLOG we assume that  $L$  does not assign conflicting values to any fluent). We use  $\mathcal{L}(F)$  to denote the set of all literal sets on  $F$ , i.e. all partial assignments of values to fluents.

A *state*  $s$  is a full assignment of values to fluents, i.e.  $|s| = |F|$ , so the size of the state space is  $2^{|F|}$ . Explicitly including negative literals  $\neg f$  in states simplifies subsequent definitions but often, we will abuse notation by defining a state  $s$  only in terms of the fluents that are true in  $s$ , as is common in STRIPS planning.

A *classical planning frame* is a tuple  $\Phi = \langle F, A \rangle$ , where  $F$  is a set of fluents and  $A$  is a set of actions. Each action  $a \in A$  comprises three sets of literals:

- $\text{pre}(a) \subseteq \mathcal{L}(F)$ , called *preconditions*, the literals that must hold for the action  $a \in A$  to be applicable.
- $\text{eff}^+(a) \subseteq \mathcal{L}(F)$ , called *positive effects*, that defines the fluents set to true by the application of the action  $a \in A$ .
- $\text{eff}^-(a) \subseteq \mathcal{L}(F)$ , called *negative effects*, that defines the fluents set to false by the action application.

We say that an action  $a \in A$  is *applicable* in a state  $s$  iff  $\text{pre}(a) \subseteq s$ . The result of applying  $a$  in  $s$  is the *successor state* denoted by  $\theta(s, a) = \{s \setminus \text{eff}^-(a)\} \cup \text{eff}^+(a)$ .

A *classical planning problem* is a tuple  $P = \langle F, A, I, G \rangle$ , where  $I$  is an initial state and  $G \subseteq \mathcal{L}(F)$  is a goal condition. A *plan* for  $P$  is an action sequence  $\pi = \langle a_1, \dots, a_n \rangle$  that induces the *state trajectory*  $\langle s_0, s_1, \dots, s_n \rangle$  such that  $s_0 = I$

and, for each  $1 \leq i \leq n$ ,  $a_i$  is applicable in  $s_{i-1}$  and generates the successor state  $s_i = \theta(s_{i-1}, a_i)$ . The *plan length* is denoted with  $|\pi| = n$ . A plan  $\pi$  *solves*  $P$  iff  $G \subseteq s_n$ , i.e. if the goal condition is satisfied at the last state reached after following the application of the plan  $\pi$  in the initial state  $I$ .

### Classical planning with conditional effects

Our approach for learning STRIPS action models is compiling the learning task into a classical planning task with conditional effects.

An action  $a \in A$  is now defined as a set of *preconditions*  $pre(a) \in \mathcal{L}(F)$  and a set of *conditional effects*  $cond(a)$ . Each conditional effect  $C \triangleright E \in cond(a)$  is composed of two sets of literals  $C \in \mathcal{L}(F)$ , the *condition*, and  $E \in \mathcal{L}(F)$ , the *effect*. An action  $a \in A$  is *applicable* in a state  $s$  if and only if  $pre(a) \subseteq s$ , and the *triggered effects* resulting from the action application are the effects whose conditions hold in  $s$ :

$$triggered(s, a) = \bigcup_{C \triangleright E \in cond(a), C \subseteq s} E,$$

The result of applying action  $a$  in state  $s$  is the *successor* state  $\theta(s, a) = \{s \setminus eff_c^-(s, a)\} \cup eff_c^+(s, a)$  where  $eff_c^-(s, a) \subseteq triggered(s, a)$  and  $eff_c^+(s, a) \subseteq triggered(s, a)$  are, respectively, the triggered *negative* and *positive* effects.

## 2.2 STRIPS action schemes

This work addresses the learning of PDDL action schemes that follow the STRIPS requirement [McDermott *et al.*, 1998; Fox and Long, 2003]. Figure 1 shows the *stack* action schema, coded in PDDL, from a four-operator *blocksworld* [Slaney and Thiébaux, 2001].

To formalize the output of the learning task, we assume that fluents  $F$  are instantiated from a set of *predicates*  $\Psi$ , as in PDDL. Each predicate  $p \in \Psi$  has an argument list of arity  $ar(p)$ . Given a set of *objects*  $\Omega$ , the set of fluents  $F$  is induced by assigning objects in  $\Omega$  to the arguments of predicates in  $\Psi$ , i.e.  $F = \{p(\omega) : p \in \Psi, \omega \in \Omega^{ar(p)}\}$  s.t.  $\Omega^k$  is the  $k$ -th Cartesian power of  $\Omega$ .

Let  $\Omega_v = \{v_i\}_{i=1}^{\max_{a \in A} ar(a)}$  be a new set of objects ( $\Omega \cap \Omega_v = \emptyset$ ), denoted as *variable names*, and that is bound by the maximum arity of an action in a given planning frame. For instance, in a three-block *blocksworld*  $\Omega = \{block_1, block_2, block_3\}$  while  $\Omega_v = \{v_1, v_2\}$  because the operators with the maximum arity, *stack* and *unstack*, have arity two.

Let us also define  $F_v$ , a new set of fluents  $F \cap F_v = \emptyset$ , that results from instantiating  $\Psi$  using only the objects in  $\Omega_v$  and that defines the elements that can appear in an action schema. For the *blocksworld*,  $F_v = \{handempty, holding(v_1), holding(v_2), clear(v_1), clear(v_2), ontable(v_1), ontable(v_2), on(v_1, v_1), on(v_1, v_2), on(v_2, v_1), on(v_2, v_2)\}$ .

We assume also that actions  $a \in A$  are instantiated from STRIPS operator schemes  $\xi = \langle head(\xi), pre(\xi), add(\xi), del(\xi) \rangle$  where:

- $head(\xi) = \langle name(\xi), pars(\xi) \rangle$ , is the operator *header* defined by its name and the corresponding

```
(:action stack
:parameters (?v1 ?v2 - object)
:precondition (and (holding ?v1) (clear ?v2))
:effect (and (not (holding ?v1)) (not (clear ?v2))
(handempty) (clear ?v1) (on ?v1 ?v2)))
```

Figure 1: STRIPS operator schema coding, in PDDL, the *stack* action from the *blocksworld*.

*variable names*,  $pars(\xi) = \{v_i\}_{i=1}^{ar(\xi)}$ . The headers of a four-operator *blocksworld* are *pickup*( $v_1$ ), *putdown*( $v_1$ ), *stack*( $v_1, v_2$ ) and *unstack*( $v_1, v_2$ ).

- The preconditions  $pre(\xi) \subseteq F_v$ , the negative effects  $del(\xi) \subseteq F_v$ , and the positive effects  $add(\xi) \subseteq F_v$  such that,  $del(\xi) \subseteq pre(\xi)$ ,  $del(\xi) \cap add(\xi) = \emptyset$  and  $pre(\xi) \cap add(\xi) = \emptyset$ .

Finally we define  $F_v(\xi) \subseteq F_v$  as the subset of elements that can appear in a given action schema  $\xi$ . For instance, for the *stack* action schema  $F_v(stack) = F_v$  while  $F_v(pickup) = \{handempty, holding(v_1), clear(v_1), ontable(v_1), on(v_1, v_1)\}$  only contains the fluents from  $F_v$  that do not involve  $v_2$ .

## 3 Learning STRIPS action models

Learning STRIPS action models from fully available input knowledge, i.e. from plans where every action in the plan is available as well as its corresponding *pre-* and *post-states*, is straightforward. In this case, STRIPS operator schemes are derived lifting the literals that change between the pre and post-state of the corresponding action executions. Preconditions are derived lifting the minimal set of literals that appears in all the pre-states of the corresponding actions.

We formalize a more challenging learning task, where less input knowledge is available. The addressed learning task corresponds to observing an agent acting in the world but watching only the results of its plan executions, the actual executed actions are unobserved. This learning task is defined as  $\Lambda = \langle \Psi, \Xi, \mathcal{O} \rangle$ :

- $\Psi$  is the set of predicates that define the abstract state space of a given classical planning frame.
- $\Xi$  is the set of empty operator schemes containing only the headers with the name and parameter of each operator schema  $\xi \in \Xi$ .
- $\mathcal{O} = \{s_0, s_1, \dots, s_n\}$  is a sequence of *state observations* obtained observing the execution of an *unobserved* plan  $\pi = \langle a_1, \dots, a_n \rangle$ .

A solution to  $\Lambda$  is a set of operator schema  $\Xi'$  compliant with the predicates in  $\Psi$ , the headers in  $\Xi$  and the sequence of state observations  $\mathcal{O}$ . A planning compilation is a suitable approach for addressing a  $\Lambda$  learning task because a solution must not only determine the STRIPS action model  $\Xi'$  but also, the *unobserved* plan  $\pi = \langle a_1, \dots, a_n \rangle$ , that can explain  $\mathcal{O}$ . Figure 2 shows an example of a  $\Lambda$  task for learning a STRIPS action model in the *blocksworld* from the sequence of five state observations that corresponds to inverting a 2-blocks tower.

```

;;; Predicates in  $\Psi$ 
(handempty) (holding ?o - object)
(clear ?o - object) (ontable ?o - object)
(on ?o1 - object ?o2 - object)

;;; Headers in  $\Xi$ 
(pickup v1) (putdown v1) (stack v1 v2) (unstack v1 v2)

;;; Observations in  $\mathcal{O}$ 
;;; observation #0
(clear block2) (on block2 block1) (ontable block1)
(handempty)

;;; observation #1
(holding block2) (clear block1) (ontable block1)

;;; observation #2
(clear block1) (ontable block1) (clear block2)
(ontable block2) (handempty)

;;; observation #3
(holding block1) (clear block2) (ontable block2)

;;; observation #4
(clear block1) (on block1 block2) (ontable block2)
(handempty)

```

Figure 2: Example of a task for learning a STRIPS action model in the blocksworld from a sequence of five state observations.

### 3.1 Learning with classical planning

Our approach for addressing the learning task, is compiling it into a classical planning task with conditional effects. The intuition behind the compilation is that a solution to the resulting classical planning task is a sequence of actions that:

1. Programs the STRIPS action model  $\Xi'$ . A solution plan starts with a *prefix* that, for each  $\xi \in \Xi$ , determines which fluents  $f \in F_v$  belong to its  $pre(\xi)$ ,  $del(\xi)$  and  $add(\xi)$  sets.
2. Validates the STRIPS action model  $\Xi'$  in  $\mathcal{O}$ . The solution plan continues with a *postfix* that produces the sequence of states in  $\mathcal{O}$  starting from  $s_0$ , and using the programmed action model  $\Xi'$ .

Given a learning task  $\Lambda = \langle \Psi, \Xi, \mathcal{O} \rangle$  the compilation outputs a classical planning task  $P_\Lambda = \langle F_\Lambda, A_\Lambda, I_\Lambda, G_\Lambda \rangle$ :

- $F_\Lambda$  contains:
  - Fluents  $F$  built instantiating the predicates  $\Psi$  with the objects appearing in the input observations  $\mathcal{O}$ .
  - Fluents representing the programmed action model  $pre_f(\xi)$ ,  $del_f(\xi)$  and  $add_f(\xi)$ , for every  $f \in F_v(\xi)$ . If a fluent  $pre_f(\xi)/del_f(\xi)/add_f(\xi)$  holds, it means that  $f$  is a precondition/negative effect/positive effect in the STRIPS operator schema  $\xi \in \Xi$ . For instance, the preconditions of the *stack* schema (Figure 1) are represented by fluents `pre_holding_stack_v1` and `pre_clear_stack_v2` set to *True*.
  - Fluent  $mode_{prog}$  indicating whether the operator schemes are programmed or validated (already programmed) and fluents  $\{test_i\}_{1 \leq i \leq n}$ , indicating the observation where the action model is validated.

- $I_\Lambda$  contains the fluents from  $F$  that encode  $s_0$  (the first observation) and every  $pre_f(\xi) \in F_\Lambda$  and  $mode_{prog}$  set to true. Our compilation assumes that initially operator schemes are programmed with every possible precondition, no negative effect and no positive effect.
- $G_\Lambda = \bigcup_{1 \leq i \leq n} \{test_i\}$ , indicates that the programmed action model is validated in all the input observations.
- $A_\Lambda$  comprises three kinds of actions:

1. Actions for *programming* operator schema  $\xi \in \Xi$ :
  - Actions for **removing** a *precondition*  $f \in F_v(\xi)$  from the action schema  $\xi \in \Xi$ .

$$\begin{aligned}
pre(\text{programPre}_f, \xi) &= \{\neg del_f(\xi), \neg add_f(\xi), \\
&\quad mode_{prog}, pre_f(\xi)\}, \\
cond(\text{programPre}_f, \xi) &= \{\emptyset\} \triangleright \{\neg pre_f(\xi)\}.
\end{aligned}$$

- Actions for **adding** a *negative or positive* effect  $f \in F_v(\xi)$  to the action schema  $\xi \in \Xi$ .

$$\begin{aligned}
pre(\text{programEff}_f, \xi) &= \{\neg del_f(\xi), \neg add_f(\xi), \\
&\quad mode_{prog}\}, \\
cond(\text{programEff}_f, \xi) &= \{pre_f(\xi)\} \triangleright \{del_f(\xi)\}, \\
&\quad \{\neg pre_f(\xi)\} \triangleright \{add_f(\xi)\}.
\end{aligned}$$

2. Actions for *applying* an already programmed operator schema  $\xi \in \Xi$  bound with the objects  $\omega \subseteq \Omega^{ar}(\xi)$ . We assume operators headers are known so the binding of the operator schema is done implicitly by order of appearance of the action parameters, i.e. variables  $vars(\xi)$  are bound to the objects in  $\omega$  appearing at the same position. Figure 3 shows the PDDL encoding of the action for applying a programmed operator *stack*.

$$\begin{aligned}
pre(\text{apply}_{\xi, \omega}) &= \{pre_f(\xi) \implies p(\omega)\}_{\forall p \in \Psi, f=p(vars(\xi))}, \\
cond(\text{apply}_{\xi, \omega}) &= \{del_f(\xi)\} \triangleright \{\neg p(\omega)\}_{\forall p \in \Psi, f=p(vars(\xi))}, \\
&\quad \{add_f(\xi)\} \triangleright \{p(\omega)\}_{\forall p \in \Psi, f=p(vars(\xi))}, \\
&\quad \{mode_{prog}\} \triangleright \{\neg mode_{prog}\}.
\end{aligned}$$

3. Actions for *validating* an observation  $1 \leq i \leq n$ .

$$\begin{aligned}
pre(\text{validate}_i) &= s_i \cup \{test_j\}_{j \in 1 \leq j < i} \\
&\quad \cup \{\neg test_j\}_{j \in t \leq j \leq n} \cup \{\neg mode_{prog}\}, \\
cond(\text{validate}_i) &= \{\emptyset\} \triangleright \{test_i\}.
\end{aligned}$$

### 3.2 Compilation properties

**Lemma 1.** *Soundness.* Any classical plan  $\pi$  that solves  $P_\Lambda$  induces an action model  $\Xi$  that solves the learning task  $\Lambda$ .

*Proof sketch.* The compilation forces that once the preconditions of an operator schema  $\xi \in \Xi$  are programmed, they cannot be altered. The same happens with the positive and negative effects (furthermore, effects can only be programmed after preconditions are programmed). Once operator schemes are programmed they can only be applied because of the  $mode_{prog}$  fluent. To solve  $P_\Lambda$ , goals  $\{test_i\}$ ,  $1 \leq i \leq n$  can only be achieved: executing an applicable

```

(:action apply_stack
:parameters (?o1 - object ?o2 - object)
:precondition
  (and (or (not (pre_on_stack_v1_v1)) (on ?o1 ?o1))
        (or (not (pre_on_stack_v1_v2)) (on ?o1 ?o2))
        (or (not (pre_on_stack_v2_v1)) (on ?o2 ?o1))
        (or (not (pre_on_stack_v2_v2)) (on ?o2 ?o2))
        (or (not (pre_ontable_stack_v1)) (ontable ?o1))
        (or (not (pre_ontable_stack_v2)) (ontable ?o2))
        (or (not (pre_clear_stack_v1)) (clear ?o1))
        (or (not (pre_clear_stack_v2)) (clear ?o2))
        (or (not (pre_holding_stack_v1)) (holding ?o1))
        (or (not (pre_holding_stack_v2)) (holding ?o2))
        (or (not (pre_handempty_stack)) (handempty)))
:effect
  (and (when (del_on_stack_v1_v1) (not (on ?o1 ?o1)))
        (when (del_on_stack_v1_v2) (not (on ?o1 ?o2)))
        (when (del_on_stack_v2_v1) (not (on ?o2 ?o1)))
        (when (del_on_stack_v2_v2) (not (on ?o2 ?o2)))
        (when (del_ontable_stack_v1) (not (ontable ?o1)))
        (when (del_ontable_stack_v2) (not (ontable ?o2)))
        (when (del_clear_stack_v1) (not (clear ?o1)))
        (when (del_clear_stack_v2) (not (clear ?o2)))
        (when (del_holding_stack_v1) (not (holding ?o1)))
        (when (del_holding_stack_v2) (not (holding ?o2)))
        (when (del_handempty_stack) (not (handempty)))
        (when (add_on_stack_v1_v1) (on ?o1 ?o1))
        (when (add_on_stack_v1_v2) (on ?o1 ?o2))
        (when (add_on_stack_v2_v1) (on ?o2 ?o1))
        (when (add_on_stack_v2_v2) (on ?o2 ?o2))
        (when (add_ontable_stack_v1) (ontable ?o1))
        (when (add_ontable_stack_v2) (ontable ?o2))
        (when (add_clear_stack_v1) (clear ?o1))
        (when (add_clear_stack_v2) (clear ?o2))
        (when (add_holding_stack_v1) (holding ?o1))
        (when (add_holding_stack_v2) (holding ?o2))
        (when (add_handempty_stack) (handempty))
        (when (modeProg) (not (modeProg)))))

```

Figure 3: Action for applying an already programmed schema *stack* as encoded in PDDL (implications coded as disjunctions).

sequence of programmed operator schemes that reaches the state  $s_i$ , starting from  $s_0$ . If this is achieved for every  $1 \leq i \leq n$ , it means that the programmed action model  $\Xi$  is compliant with the provided input knowledge and hence, solves  $\Lambda$ .  $\square$

**Lemma 2. Completeness.** Any STRIPS action model  $\Xi$  computable from  $\Lambda = \langle \Psi, \Xi, \mathcal{O} \rangle$  can be obtained by solving the corresponding classical planning task  $P_\Lambda$ .

*Proof sketch.* By definition,  $F_v(\xi) \subseteq F_\Lambda$  fully captures the set of elements that can appear in a STRIPS action schema  $\xi \in \Xi$ . Any possible STRIPS action schema built with the fluents in  $F_v$  can be computed because the compilation only constrains the application of  $\text{apply}_{\xi, \omega}$  actions iff they are not compliant with the  $\mathcal{O}$  sequence.  $\square$

## 4 Recognizing STRIPS action models

Inspired by the *Plan recognition as planning* approach [Ramirez and Geffner, 2009] we show that the previous compilation can be extended to estimate the probability distribution of the possible STRIPS models given a sequence of state observations. This extension allows us to address model recognition tasks as well as to semantically evaluate the quality of learned STRIPS models

### 4.1 The STRIPS edit distance

We assume that how well a given STRIPS action model  $\Xi$  explains a given sequence of observations  $\mathcal{O}$ , depends on the amount of edits that one has to introduce to  $\Xi$  to produce the sequence of observations  $\mathcal{O}$ . If the given model  $\Xi$  perfectly explains the observations  $\mathcal{O}$ , no edit have to be introduced and the probability of explaining the observations is  $P(\Xi|\mathcal{O}) = 1$ . On the other hand, with a probability of explaining the observations of  $P(\Xi|\mathcal{O}) = 0$ , is the model that requires the maximum number of edits to explain the sequence of observations  $\mathcal{O}$ .

To compute the  $P(\Xi|\mathcal{O})$  probability distribution, we need a metric that quantifies this amount of edits. To define such metric, we need to define first which are the possible edit operations that can be performed on a STRIPS action model. With the aim of keeping tractable the branching factor of the classical planning task resulting from our compilation, we only define two STRIPS edit operations:

- *Deletion.* A fluent  $\text{pre}_f(\xi)/\text{del}_f(\xi)/\text{add}_f(\xi)$  is removed from the operator schema  $\xi$
- *Insertion.* A fluent  $\text{pre}_f(\xi)/\text{del}_f(\xi)/\text{add}_f(\xi)$  is added to the operator schema  $\xi$

Now we are ready to define a metric that assesses then how dissimilar two given STRIPS action models are

**Definition 3.** The **STRIPS edit distance**, denoted as  $\delta(\Xi, \Xi')$ , is the minimum number of edit operations that allow to transform a given STRIPS action model  $\Xi$  into the given STRIPS action model  $\Xi'$ . This edit distance satisfies the metric axioms provided that the two operations has the same positive cost.

Since the size of the  $F_v$  set is bound, we can define the maximum number of edits that can be introduced to a given STRIPS action model.

**Definition 4.** The maximum **STRIPS edit distance** is an upperbound, denoted as  $\delta(\Xi, *) = \sum_{\xi \in \Xi} 3|F_v(\xi)|$ , that indicates the maximum distance required to transform a given STRIPS action model  $\Xi$  into any STRIPS action model  $\Xi'$ .

In more detail, for a given operator schema  $\xi$  the maximum number of edits that can be introduced to their precondition set is  $|F_v(\xi)|$ . With regard to action effects, the maximum number of edits that can be introduced to their precondition set is  $2|F_v(\xi)|$ .

Likewise we can define the edit distance with regard to a sequence of observations  $\mathcal{O}_i$

**Definition 5.**

### 4.2 Recognition and evaluation of STRIPS models

The probability distribution of the possible STRIPS models given a sequence of state observations allows to evaluate the quality of a learned STRIPS action models with regard to a sequence of observations that act as a test set.

### 4.3 Recognition with classical planning

## 5 Evaluation

This section evaluates our approach for learning STRIPS models starting from different amounts of input knowledge.

## Reproducibility

We used IPC domains that satisfy the STRIPS requirement [Fox and Long, 2003], taken from the PLANING.DOMAINS repository [Muise, 2016]. For the learning of the STRIPS action models we used observations sequences of 25 states per domain. All experiments are run on an Intel Core i5 3.10 GHz x 4 with 4 GB of RAM.

MADAGASCAR is the classical planner we use to solve the instances that result from our compilations because its ability to deal with dead-ends [Rintanen, 2014]. In addition, MADAGASCAR can apply the actions for programming preconditions in a single planning step (in parallel) because these actions do not interact. Actions for programming action effects can also be applied in a single planning step reducing significantly the planning horizon.

The compilation source code, the evaluation scripts and the benchmarks are fully available at this anonymous repository <https://github.com/anonsub/strips-learning> so any experimental data reported in the paper can be reproduced.

## Supervised evaluation of the learned models

For each domain the learned model is compared with the actual model and its quality is quantified with the *precision* and *recall* metrics. Precision gives a notion of *soundness* while recall gives a notion of the *completeness* of the learned models.  $Precision = \frac{tp}{tp+fp}$ , where  $tp$  is the number of true positives (predicates that correctly appear in the action model) and  $fp$  is the number of false positives (predicates appear in the learned action model that should not appear).  $Recall = \frac{tp}{tp+fn}$  where  $fn$  is the number of false negatives (predicates that should appear in the learned action model but are missing).

When the learning hypothesis space is low constrained, the learned actions can be reformulated and still be compliant with the inputs. For instance in the *blocksworld*, operator *stack* could be *learned* with the preconditions and effects of the *unstack* operator (and vice versa). Furthermore, in a given action the role of the parameters that share the same type can be interchanged making non trivial to compute *precision* and *recall* with respect to a reference model.

To address these issues we defined an evaluation method robust to action reformulation. Precision and recall are often combined using the *harmonic mean*. This expression is called the *F-measure* (or the balanced *F-score*) and is formally defined as  $F = 2 \times \frac{Precision \times Recall}{Precision + Recall}$ . Given a reference STRIPS action model  $\Xi^*$  and the learned STRIPS action model  $\Xi$  we define the bijective function  $f_{P\&R} : \Xi \mapsto \Xi^*$  such that  $f_{P\&R}$  maximizes the accumulated *F-measure*. With this mapping defined we can compute the *precision* and *recall* of a learned STRIPS action  $\xi \in \Xi$  with respect to the action  $f_{P\&R}(\xi) \in \Xi^*$  even if actions are reformulated in the learning process.

## 6 Conclusions

As far as we know, this is the first work on learning STRIPS action models from state observations, exclusively using classical planning and evaluated over a wide range of different domains.

	Pre		Add		Del		P	R
	P	R	P	R	P	R		
blocks	0.57	0.44	0.63	0.56	0.57	0.44	0.59	0.48
driverlog	0.2	0.07	0.18	0.29	0.2	0.14	0.19	0.17
ferry	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
floor-tile	0.5	0.14	0.75	0.55	0.6	0.27	0.62	0.32
Grid	-	-	-	-	-	-	-	-
gripper-strips	0.25	0.17	0.25	0.25	0.25	0.25	0.25	0.22
hanoi	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
hiking	-	-	-	-	-	-	-	-
n-puzzle	-	-	-	-	-	-	-	-
parking	0.6	0.21	0.14	0.11	0.4	0.22	0.38	0.18
satellite	0.25	0.07	0.33	0.4	0.67	0.5	0.42	0.32
Sokoban	-	-	-	-	-	-	-	-
transport	0.33	0.1	0.33	0.4	0.0	0.0	0.22	0.17
zeno-travel	0.25	0.07	0.17	0.14	0.0	0.0	0.14	0.07
	-	-	-	-	-	-	-	-

Table 1: Precision and recall values obtained when learning from labels without computing the  $f_{P\&R}$  mapping.

	Pre		Add		Del		P	R
	P	R	P	R	P	R		
blocks	0.71	0.56	0.75	0.67	0.71	0.56	0.73	0.59
driverlog	0.6	0.21	0.27	0.43	0.6	0.43	0.49	0.36
ferry	0.5	0.29	0.75	0.75	0.5	0.5	0.58	0.51
floor-tile	0.5	0.14	0.75	0.55	0.6	0.27	0.62	0.32
Grid	-	-	-	-	-	-	-	-
gripper-strips	1.0	0.67	1.0	1.0	1.0	1.0	1.0	0.89
hanoi	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
hiking	-	-	-	-	-	-	-	-
n-puzzle	-	-	-	-	-	-	-	-
parking	0.6	0.21	0.14	0.11	0.4	0.22	0.38	0.18
satellite	0.25	0.07	0.33	0.4	0.67	0.5	0.42	0.32
Sokoban	-	-	-	-	-	-	-	-
transport	1.0	0.3	0.67	0.8	0.67	0.4	0.78	0.5
zeno-travel	0.5	0.14	0.5	0.43	0.33	0.14	0.44	0.24
	-	-	-	-	-	-	-	-

Table 2: Precision and recall values obtained when learning from labels but computing the  $f_{P\&R}$  mapping.

	Total time	Preprocess	Plan length
Blocks	1.40	0.00	70
Driverlog	1.50	0.00	89
Ferry	1.49	0.00	64
Floortile	351.38	0.11	156
Grid	-	-	-
Gripper	0.04	0.00	59
Hanoi	2.33	0.01	49
Hiking	-	-	-
Parking	-	-	-
Satellite	78.36	0.03	98
Sokoban	-	-	-
Transport	285.61	0.10	106
Zenotravel	6.20	0.46	71

Table 3: Planning results obtained when learning from labels.

The empirical results show that since our approach is strongly based on inference can generate non-trivial models from very small data sets. In addition, the SAT-based planner MADAGASCAR is particularly suitable for the approach because its ability to deal with planning instances populated with dead-ends and because many actions for programming the STRIPS model can be done in parallel since they do not interact reducing significantly the planning horizon.

## References

- [Amir and Chang, 2008] Eyal Amir and Allen Chang. Learning partially observable deterministic action models. *Journal of Artificial Intelligence Research*, 33:349–402, 2008.
- [Bonet *et al.*, 2009] Blai Bonet, Héctor Palacios, and Héctor Geffner. Automatic derivation of memoryless policies and finite-state controllers using classical planners. In *ICAPS*, 2009.
- [Cresswell *et al.*, 2013] Stephen N Cresswell, Thomas L McCluskey, and Margaret M West. Acquiring planning domain models using LOCM. *The Knowledge Engineering Review*, 28(02):195–213, 2013.
- [Fox and Long, 2003] Maria Fox and Derek Long. PDDL2.1: An extension to PDDL for expressing temporal planning domains. *J. Artif. Intell. Res.(JAIR)*, 20:61–124, 2003.
- [Geffner and Bonet, 2013] Hector Geffner and Blai Bonet. A concise introduction to models and methods for automated planning, 2013.
- [Ghallab *et al.*, 2004] Malik Ghallab, Dana Nau, and Paolo Traverso. *Automated Planning: theory and practice*. Elsevier, 2004.
- [Kambhampati, 2007] Subbarao Kambhampati. Model-lite planning for the web age masses: The challenges of planning with incomplete and evolving domain models. In *Proceedings of the National Conference on Artificial Intelligence*, 2007.
- [McDermott *et al.*, 1998] Drew McDermott, Malik Ghallab, Adele Howe, Craig Knoblock, Ashwin Ram, Manuela Veloso, Daniel Weld, and David Wilkins. PDDL – The Planning Domain Definition Language, 1998.
- [Muisse, 2016] Christian Muise. Planning. domains. *ICAPS system demonstration*, 2016.
- [Ramirez and Geffner, 2009] Miquel Ramirez and Hector Geffner. Plan recognition as planning. In *Proceedings of the 21st international joint conference on Artificial intelligence*. Morgan Kaufmann Publishers Inc, pages 1778–1783, 2009.
- [Ramírez, 2012] Miquel Ramírez. *Plan recognition as planning*. PhD thesis, Universitat Pompeu Fabra, 2012.
- [Rintanen, 2014] Jussi Rintanen. Madagascar: Scalable planning with sat. *Proceedings of the 8th International Planning Competition (IPC-2014)*, 2014.
- [Segovia-Aguas *et al.*, 2016] Javier Segovia-Aguas, Sergio Jiménez, and Anders Jonsson. Hierarchical finite state controllers for generalized planning. In *Proceedings of the Twenty-Fifth International Joint Conference on Artificial Intelligence*, pages 3235–3241. AAAI Press, 2016.
- [Segovia-Aguas *et al.*, 2017] Javier Segovia-Aguas, Sergio Jiménez, and Anders Jonsson. Generating context-free grammars using classical planning. In *International Joint Conference on Artificial Intelligence*, 2017.
- [Slaney and Thiébaux, 2001] John Slaney and Sylvie Thiébaux. Blocks world revisited. *Artificial Intelligence*, 125(1-2):119–153, 2001.
- [Yang *et al.*, 2007] Qiang Yang, Kangheng Wu, and Yunfei Jiang. Learning action models from plan examples using weighted max-sat. *Artificial Intelligence*, 171(2-3):107–143, 2007.
- [Zhuo and Kambhampati, 2013] Hankz Hankui Zhuo and Subbarao Kambhampati. Action-model acquisition from noisy plan traces. In *IJCAI*, pages 2444–2450, 2013.