Learning and Recognition of STRIPS Action Models from State Observations
Action Model Recognition from State Observations via Planning

Abstract

This paper presents a classical planning compilation for learning STRIPS action models from state observations. The compilation approach does not require observing the precise actions that produced the state observations because such actions are determined by a classical planner. In addition, the compilation is extensible to (1) semantically evaluate the quality of learned action models and (2) estimate the probability distribution of the possible action models given a state observation sequence.

EVA

This paper presents a first novel approach to STRIPS action model recognition from state observations. Unlike other related works, our classical planning compilation approach not only allows learning a STRIPS action model from an observation sequence but it is also extendible to action model recognition. The learning task does not require observing the precise actions that produce the state observations because such actions are determined by a classical planner. We present a two-fold evaluation of the learned models: a syntactic and semantic assessment against a reference model and via a test dataset that calculates the probability of a model given the observation. Moreover, we show how these probabilistic values are used to estimate the probability of the action model recognition.

1 Introduction

Besides *plan synthesis* [Ghallab *et al.*, 2004], planning action models are also useful for *plan/goal recognition* [Ramírez, 2012]. At these planning tasks, automated planners are required to reason about an action model that correctly and completely captures the possible world transitions [Geffner and Bonet, 2013]. Unfortunately, building planning action models is complex, even for planning experts, and this knowledge acquisition task is a bottleneck that limits the potential of *AI planning* [Kambhampati, 2007].

The Machine Learning of planning action models is a promising alternative to hand-coding them and nowadays, there exist sophisticated algorithms like AMAN [Zhuo and Kambhampati, 2013], ARMS [Yang *et al.*, 2007], LOCM [Cresswell *et al.*, 2013] or SLAF [Amir and Chang, 2008]. Motivated by recent advances on the synthesis of different kinds of generative models with classical planning [Bonet *et al.*, 2009; Segovia-Aguas *et al.*, 2016; 2017], this paper presents a novel approach for learning STRIPS action models that introduces the following contributions:

1. Is defined as a classical planning compilation. This fact opens the door to the *bootstrapping* of planning action models and allows us to report results over a wide range of planning domains.

2. Does not require observing the particular executed actions. An off-the-shelf classical planner determines these actions given the state observations.

3. Is extensible to semantically evaluate the quality of learned STRIPS models as well as to estimate the probability distribution of the possible STRIPS models given a sequence of state observations.

2 Background

This section defines the planning models used in this work as well as the output (an STRIPS action model) of the addressed learning task.

2.1 Classical planning

We use F to denote the set of *fluents* (propositional variables) describing a state. A *literal* l is a valuation of a fluent $f \in F$, i.e. either $l = f$ or $l = \neg f$. A set of literals L represents a partial assignment of values to fluents (WLOG we assume that L does not assign conflicting values to any fluent). We use $\mathcal{L}(F)$ to denote the set of all literal sets on F, i.e. all partial assignments of values to fluents.

A *state* s is a full assignment of values to fluents, i.e. $|s| = |F|$, so the size of the state space is $2^{|F|}$. Explicitly including negative literals $\neg f$ in states simplifies subsequent definitions but often, we will abuse notation by defining a state s only in terms of the fluents that are true in s, as is common in STRIPS planning.

A *classical planning frame* is a tuple $\Phi = \langle F, A \rangle$, where F is a set of fluents and A is a set of actions. Each action $a \in A$ comprises three sets of literals:

- pre$(a) \subseteq \mathcal{L}(F)$, called *preconditions*, the literals that must hold for the action $a \in A$ to be applicable.

- $eff^+(a) \subseteq \mathcal{L}(F)$, called *positive effects*, that defines the fluents set to true by the application of the action $a \in A$.

- $eff^-(a) \subseteq \mathcal{L}(F)$, called *negative effects*, that defines the fluents set to false by the action application.

We say that an action $a \in A$ is *applicable* in a state s iff $pre(a) \subseteq s$. The result of applying a in s is the *successor state* denoted by $\theta(s, a) = \{s \setminus eff^-(a)) \cup eff^+(a)\}$.

A *classical planning problem* is a tuple $P = \langle F, A, I, G \rangle$, where I is an initial state and $G \subseteq \mathcal{L}(F)$ is a goal condition. A *plan* for P is an action sequence $\pi = \langle a_1, \ldots, a_n \rangle$ that induces the *state trajectory* $\langle s_0, s_1, \ldots, s_n \rangle$ such that $s_0 = I$ and, for each $1 \leq i \leq n$, a_i is applicable in s_{i-1} and generates the successor state $s_i = \theta(s_{i-1}, a_i)$. The *plan length* is denoted with $|\pi| = n$. A plan π *solves* P iff $G \subseteq s_n$, i.e. if the goal condition is satisfied at the last state reached after following the application of the plan π in the initial state I.

Classical planning with conditional effects

Our approach for learning STRIPS action models is compiling the leaning task into a classical planning task with conditional effects.

An action $a \in A$ with conditional effects is defined as a set of *preconditions* $pre(a) \in \mathcal{L}(F)$ and a set of *conditional effects* $cond(a)$. Each conditional effect $C \triangleright E \in cond(a)$ is composed of two sets of literals $C \in \mathcal{L}(F)$, the *condition*, and $E \in \mathcal{L}(F)$, the *effect*. An action $a \in A$ is *applicable* in a state s if and only if $pre(a) \subseteq s$, and the *triggered effects* resulting from the action application are the effects whose conditions hold in s:

$$triggered(s, a) = \bigcup_{C \triangleright E \in cond(a), C \subseteq s} E,$$

The result of applying action a in state s is the *successor* state $\theta(s, a) = \{s \setminus eff_c^-(s, a)) \cup eff_c^+(s, a)\}$ where $eff_c^-(s, a) \subseteq triggered(s, a)$ and $eff_c^+(s, a) \subseteq triggered(s, a)$ are, respectively, the triggered *negative* and *positive* effects.

2.2 STRIPS action schemes

This work addresses the learning of PDDL action schemes that follow the STRIPS requirement [McDermott *et al.*, 1998; Fox and Long, 2003]. Figure 1 shows the *stack* action schema, coded in PDDL, from a four-operator *blocksworld* [Slaney and Thiébaux, 2001].

To formalize the output of the learning task, we assume that fluents F are instantiated from a set of *predicates* Ψ, as in PDDL. Each predicate $p \in \Psi$ has an argument list of arity $ar(p)$. Given a set of *objects* Ω, the set of fluents F is induced by assigning objects in Ω to the arguments of predicates in Ψ, i.e. $F = \{p(\omega) : p \in \Psi, \omega \in \Omega^{ar(p)}\}$ s.t. Ω^k is the k-th Cartesian power of Ω.

Let $\Omega_v = \{v_i\}_{i=1}^{\max_{a \in A} ar(a)}$ be a new set of objects ($\Omega \cap \Omega_v = \emptyset$), denoted as *variable names*, and that is bound by the maximum arity of an action in a given planning frame. For instance, in a three-block blocksworld $\Omega = \{block_1, block_2, block_3\}$ while $\Omega_v = \{v_1, v_2\}$ because the operators with the maximum arity, stack and unstack, have arity two.

```
(:action stack
 :parameters (?v1 ?v2 - object)
 :precondition (and (holding ?v1) (clear ?v2))
 :effect (and (not (holding ?v1))
              (not (clear ?v2))
              (handempty) (clear ?v1)
              (on ?v1 ?v2)))
```

Figure 1: STRIPS operator schema coding, in PDDL, the *stack* action from the *blocksworld*.

Let us also define F_v, a new set of fluents $F \cap F_v = \emptyset$, that results from instantiating Ψ using only the objects in Ω_v and that defines the elements that can appear in an action schema. For the blocksworld, F_v={handempty, holding(v_1), holding(v_2), clear(v_1), clear(v_2), ontable(v_1), ontable(v_2), on(v_1, v_1), on(v_1, v_2), on(v_2, v_1), on(v_2, v_2)}.

We assume also that actions $a \in A$ are instantiated from STRIPS operator schemes $\xi = \langle head(\xi), pre(\xi), add(\xi), del(\xi) \rangle$ where:

- $head(\xi) = \langle name(\xi), pars(\xi) \rangle$, is the operator *header* defined by its name and the corresponding *variable names*, $pars(\xi) = \{v_i\}_{i=1}^{ar(\xi)}$. The headers of a four-operator blocksworld are pickup(v_1), putdown(v_1), stack(v_1, v_2) and unstack(v_1, v_2).

- The preconditions $pre(\xi) \subseteq F_v$, the negative effects $del(\xi) \subseteq F_v$, and the positive effects $add(\xi) \subseteq F_v$ such that, $del(\xi) \subseteq pre(\xi)$, $del(\xi) \cap add(\xi) = \emptyset$ and $pre(\xi) \cap add(\xi) = \emptyset$.

Finally we also define $F_v(\xi) \subseteq F_v$ as the subset of elements that can appear in a given action schema ξ. For instance, for the *stack* action schema $F_v(\texttt{stack}) = F_v$ while $F_v(\texttt{pickup})$={handempty, holding(v_1), clear(v_1), ontable(v_1), on(v_1, v_1)} only contains the fluents from F_v that do not involve v_2 because the action header contained the single parameter v_1.

3 Learning STRIPS action models

Learning STRIPS action models from fully available input knowledge, i.e. from plans where every action in the plan is available as well as its corresponding *pre-* and *post-states*, is straightforward. In this case, STRIPS operator schemes are derived lifting the literals that change between the pre and post-state of the corresponding action executions. Preconditions are derived lifting the minimal set of literals that appears in all the pre-states of the corresponding actions.

In this paper we address a more challenging learning task, where less input knowledge is available. The addressed learning task corresponds to observing an agent acting in the world but watching only the results of its plan executions, the actual executed actions are unobserved. This learning task is defined as $\Lambda = \langle \Psi, \Xi, \mathcal{O} \rangle$, where:

- Ψ is the set of predicates that define the abstract state space of a given classical planning frame.

- Ξ is the set of empty operator schemes, wherein each $\xi \in \Xi$ is only composed of $head(\xi)$.

- $\mathcal{O} = \langle s_0, s_1, \ldots, s_n \rangle$ is a sequence of *state observations* obtained observing the execution of an *unobserved* plan $\pi = \langle a_1, \ldots, a_n \rangle$.

A solution to Λ is a set of operator schema Ξ' compliant with the predicates in Ψ, the headers in Ξ and the given sequence of state observations $\mathcal{O}$. A planning compilation is a suitable approach for addressing a Λ learning task because a solution must not only determine the STRIPS action model Ξ' but also, the *unobserved* plan $\pi = \langle a_1, \ldots, a_n \rangle$, that explains $\mathcal{O}$. Figure 2 shows an example of a Λ task for learning a STRIPS action model in the blocksworld from the sequence of five state observations that corresponds to inverting a 2-blocks tower.

```
;;; Predicates in Ψ

(handempty) (holding ?o  - object)
(clear ?o - object) (ontable ?o - object)
(on ?o1 - object ?o2 - object)

;;; Headers in Ξ

(pickup v1) (putdown v1)
(stack v1 v2} (unstack v1 v2)

;;; Observations in O

;;; observation #0
(clear block2) (on block2 block1)
(ontable block1) (handempty)

;;; observation #1
(holding block2) (clear block1)
(ontable block1)

;;; observation #2
(clear block1) (ontable block1)
(clear block2) (ontable block2)
(handempty)

;;; observation #3
(holding block1) (clear block2)
(ontable block2)

;;; observation #4
(clear block1) (on block1 block2)
(ontable block2) (handempty)
```

Figure 2: Example of a Λ task for learning a STRIPS action model in the *blocksworld* from a sequence of five state observations.

3.1 Learning with classical planning

Our approach for addressing the learning task is compiling Λ into a classical planning task with conditional effects. The intuition behind the compilation is that a solution to the resulting classical planning task is a sequence of actions that:

1. Programs the STRIPS action model Ξ'. A solution plan starts with a *prefix* that, for each $\xi \in \Xi$, determines which fluents $f \in F_v(\xi)$ belong to its $pre(\xi)$, $del(\xi)$ and $add(\xi)$ sets.
2. Validates the STRIPS action model Ξ' in $\mathcal{O}$. The solution plan continues with a postfix that produces the given se-

quence of states $\langle s_0, s_1, \ldots, s_n \rangle$ using the programmed action model Ξ'.

Given a learning task $\Lambda = \langle \Psi, \Xi, \mathcal{O} \rangle$ the compilation outputs a classical planning task $P_\Lambda = \langle F_\Lambda, A_\Lambda, I_\Lambda, G_\Lambda \rangle$:

- F_Λ contains:
  - The set of fluents F that is built instantiating the predicates Ψ with the objects appearing in the input observations $\mathcal{O}$.
  - Fluents representing the programmed action model $pre_f(\xi)$, $del_f(\xi)$ and $add_f(\xi)$, for every $f \in F_v(\xi)$. If a fluent $pre_f(\xi)/del_f(\xi)/add_f(\xi)$ holds, it means that f is a precondition/negative effect/positive effect in the STRIPS operator schema $\xi \in \Xi$. For instance, the preconditions of the *stack* schema (Figure 1) are represented by the fluents pre_holding_stack_v_1 and pre_clear_stack_v_2 set to *True*.
  - Fluent $mode_{prog}$ indicating whether the operator schemes are programmed or validated (already programmed) and fluents $\{test_i\}_{1 \leq i \leq n}$, indicating the observation where the action model is validated.

- I_Λ contains the fluents from F that encode s_0 (the first observation) and $mode_{prog}$ set to true. In addition, our compilation assumes that initially operator schemes are programmed with every possible precondition, no negative effect and no positive effect. With this regard, the fluents $pre_f(\xi)$ hold at the initial state for every $f \in F_v(\xi)$.

- $G_\Lambda = \bigcup_{1 \leq i \leq n}\{test_i\}$, indicates that the programmed action model is validated in all the input observations.

- A_Λ comprises three kinds of actions:

  1. Actions for *programming* operator schema $\xi \in \Xi$:
     - Actions for **removing** a *precondition* $f \in F_v(\xi)$ from the action schema $\xi \in \Xi$.

$$\mathsf{pre}(\mathsf{programPre}_{f,\xi}) = \{\neg del_f(\xi), \neg add_f(\xi),$$
$$mode_{prog}, pre_f(\xi)\},$$
$$\mathsf{cond}(\mathsf{programPre}_{f,\xi}) = \{\emptyset\} \rhd \{\neg pre_f(\xi)\}.$$

     - Actions for **adding** a *negative* or *positive* effect $f \in F_v(\xi)$ to the action schema $\xi \in \Xi$.

$$\mathsf{pre}(\mathsf{programEff}_{f,\xi}) = \{\neg del_f(\xi), \neg add_f(\xi),$$
$$mode_{prog}\},$$
$$\mathsf{cond}(\mathsf{programEff}_{f,\xi}) = \{pre_f(\xi)\} \rhd \{del_f(\xi)\},$$
$$\{\neg pre_f(\xi)\} \rhd \{add_f(\xi)\}.$$

  2. Actions for *applying* an already programmed operator schema $\xi \in \Xi$ bound with the objects $\omega \subseteq \Omega^{ar(\xi)}$. We assume that the operators headers are known so the binding of the operator schema is done implicitly by order of appearance of the action parameters, i.e. variables $pars(\xi)$ are bound to the objects in ω appearing at the same position.

Figure 3 shows the PDDL encoding of the action for applying a programmed operator $stack$.

$$\mathsf{pre}(\mathsf{apply}_{\xi,\omega}) = \{pre_f(\xi) \implies p(\omega)\}_{\forall p \in \Psi, f=p(pars(\xi))},$$
$$\mathsf{cond}(\mathsf{apply}_{\xi,\omega}) = \{del_f(\xi)\} \rhd \{\neg p(\omega)\}_{\forall p \in \Psi, f=p(pars(\xi))},$$
$$\{add_f(\xi)\} \rhd \{p(\omega)\}_{\forall p \in \Psi, f=p(pars(\xi))},$$
$$\{mode_{prog}\} \rhd \{\neg mode_{prog}\}.$$

3. Actions for *validating* an observation $1 \leq i \leq n$.

$$\mathsf{pre}(\mathsf{validate}_i) = s_i \cup \{test_j\}_{j \in 1 \leq j < i}$$
$$\cup \{\neg test_j\}_{j \in i \leq j \leq n} \cup \{\neg mode_{prog}\},$$
$$\mathsf{cond}(\mathsf{validate}_i) = \{\emptyset\} \rhd \{test_i\}.$$

```
(:action apply_stack
 :parameters (?o1 - object ?o2 - object)
 :precondition
  (and (or (not (pre_on_stack_v1_v1)) (on ?o1 ?o1))
       (or (not (pre_on_stack_v1_v2)) (on ?o1 ?o2))
       (or (not (pre_on_stack_v2_v1)) (on ?o2 ?o1))
       (or (not (pre_on_stack_v2_v2)) (on ?o2 ?o2))
       (or (not (pre_ontable_stack_v1)) (ontable ?o1))
       (or (not (pre_ontable_stack_v2)) (ontable ?o2))
       (or (not (pre_clear_stack_v1)) (clear ?o1))
       (or (not (pre_clear_stack_v2)) (clear ?o2))
       (or (not (pre_holding_stack_v1)) (holding ?o1))
       (or (not (pre_holding_stack_v2)) (holding ?o2))
       (or (not (pre_handempty_stack)) (handempty)))
 :effect
  (and (when (del_on_stack_v1_v1) (not (on ?o1 ?o1)))
       (when (del_on_stack_v1_v2) (not (on ?o1 ?o2)))
       (when (del_on_stack_v2_v1) (not (on ?o2 ?o1)))
       (when (del_on_stack_v2_v2) (not (on ?o2 ?o2)))
       (when (del_ontable_stack_v1) (not (ontable ?o1)))
       (when (del_ontable_stack_v2) (not (ontable ?o2)))
       (when (del_clear_stack_v1) (not (clear ?o1)))
       (when (del_clear_stack_v2) (not (clear ?o2)))
       (when (del_holding_stack_v1) (not (holding ?o1)))
       (when (del_holding_stack_v2) (not (holding ?o2)))
       (when (del_handempty_stack) (not (handempty)))
       (when (add_on_stack_v1_v1) (on ?o1 ?o1))
       (when (add_on_stack_v1_v2) (on ?o1 ?o2))
       (when (add_on_stack_v2_v1) (on ?o2 ?o1))
       (when (add_on_stack_v2_v2) (on ?o2 ?o2))
       (when (add_ontable_stack_v1) (ontable ?o1))
       (when (add_ontable_stack_v2) (ontable ?o2))
       (when (add_clear_stack_v1) (clear ?o1))
       (when (add_clear_stack_v2) (clear ?o2))
       (when (add_holding_stack_v1) (holding ?o1))
       (when (add_holding_stack_v2) (holding ?o2))
       (when (add_handempty_stack) (handempty))
       (when (modeProg) (not (modeProg)))))
```

Figure 3: Action for applying an already programmed schema $stack$ as encoded in PDDL (implications coded as disjunctions).

3.2 Compilation properties

Lemma 1. *Soundness. Any classical plan π that solves P_Λ induces an action model Ξ' that solves the learning task Λ.*

Proof sketch. The compilation forces that once the preconditions of an operator schema $\xi \in \Xi'$ are programmed, they cannot be altered. The same happens with the positive and negative effects. Furthermore because of the preconditions of the programPre$_{f,\xi}$ actions, effects are only programmable after preconditions are programmed. Once the operator schemes Ξ' are programmed, they can only be applied because of the $mode_{prog}$ fluent. To solve P_Λ, the goals

$\{test_i\}$, $1 \leq i \leq n$ can only be achieved: executing an applicable sequence of programmed operator schemes that reaches every state $s_i \in \mathcal{O}$, starting from s_0 and following the sequence $1 \leq i \leq n$. Therefore if $test_n$ is achieved, it means that the programmed action model Ξ' is compliant with the provided input knowledge and hence, solves Λ. $\square$

Lemma 2. *Completeness. Any STRIPS action model Ξ' that solves a $\Lambda = \langle \Psi, \Xi, \mathcal{O} \rangle$ learning task, is computable solving the corresponding classical planning task P_Λ.*

Proof sketch. By definition, $F_v(\xi) \subseteq F_\Lambda$ fully captures the set of elements that can appear in a STRIPS action schema $\xi \in \Xi$. In addition, any possible STRIPS action schema Ξ' that can be built with the fluents in F_v can be computed with the P_Λ compilation. The only STRIPS action models that the compilation cannot compute are the ones that violate the state constraints defined by the $\mathcal{O}$ sequence, which constrain as well the solutions to the Λ learning task. $\square$

4 STRIPS action models recognition

Inspired by *plan recognition as planning* [Ramırez and Geffner, 2009], we show that the previous compilation can be extended to (1) semantically evaluate the quality of learned STRIPS models, without comparing with a reference model and (2), address STRIPS action model recognition tasks.

4.1 The edit distance

We assume that how well a given STRIPS action model Ξ explains a given sequence of observations $\mathcal{O}$, depends on the amount of *edition* that one has to introduce to Ξ to produce the sequence of observations $\mathcal{O}$. In the extreme, if the given model Ξ perfectly explains the observations $\mathcal{O}$, no *edit* to the model must be introduced.

Let us define the *edit operations* allowed to edit a given STRIPS action model. With the aim of keeping tractable the branching factor of the classical planning task that results from our compilation, we only define two STRIPS *edit operations*:

- *Deletion*. A fluent $pre_f(\xi)/del_f(\xi)/add_f(\xi)$, such that $f \in F_v(\xi)$, is removed from the operator schema $\xi \in \Xi$.
- *Insertion*. A fluent $pre_f(\xi)/del_f(\xi)/add_f(\xi)$, such that $f \in F_v(\xi)$, is added to the operator schema $\xi \in \Xi$.

With this defined, we can now formalize an edit distance metric that assesses how dissimilar two given STRIPS action models are. This distance satisfies the *metric axioms* provided that the two STRIPS *edit operations* have the same positive cost.

Definition 3. *Let Ξ and Ξ' be two action models, both built from the same set of possible elements F_v. The **edit distance**, denoted as $\delta(\Xi, \Xi')$, is the minimum number of edit operations to transform Ξ into Ξ'.*

Since the size of the F_v set is bound, the maximum number of edits that can be introduced to a given action model is bound as well. In more detail, for a given operator schema $\xi \in \Xi$ the maximum number of edits that can be introduced to their precondition set is $|F_v(\xi)|$. With regard to effects of ξ, the maximum number of edits that can be introduced are two times $|F_v(\xi)|$.

Definition 4. *Let Ξ be an action model built from the set of possible elements F_v. The **maximum edit distance**, defined as $\delta(\Xi, *) = \sum_{\xi \in \Xi} 3|F_v(\xi)|$, is an upper bound on the distance from Ξ to any* STRIPS *action model definable within F_v.*

Likewise, we can define a distance of a model Ξ to an observation sequence $\mathcal{O}$ on the basis of the *edit distance* between two models.

Definition 5. *Given an action model Ξ built from the set of possible elements F_v and a sequence of observations $\mathcal{O}$. The **observation edit distance**, denoted by $\delta(\Xi, \mathcal{O})$, is the minimal edit distance to transform Ξ into a model Ξ' which produces a plan $\pi = \langle a_1, \ldots, a_n \rangle$ such that π induces the observation sequence $\mathcal{O} = \langle s_0, s_1, \ldots, s_n \rangle$; i.e.,*

$$\delta(\Xi, \mathcal{O}) = \underset{\Xi' \to \mathcal{O}}{argmin}\, \delta(\Xi, \Xi')$$

The *observation edit distance* allow us to semantically evaluate the quality of a learned strips action model with respect to a sequence of observations that acts as a test set. The lower this distance is, the better the quality of the learned model. The semantic nature of this evaluation is robust to learning episodes where actions are reformulated and still compliant with the inputs (e.g. the *blocksworld* operator `stack` could be *learned* with the preconditions and effects of the `unstack` operator and vice versa or the roles of actions parameters with the same type could be interchanged).

4.2 Model recognition

Given a set of possible STRIPS models and a sequence of state observations the *recognition of* STRIPS *models* is the task of computing the probability distribution of these models according to the given observations. Here we show that the *observations* STRIPS *edit distance* is helpful to compute the probability distribution of the possible STRIPS models given a sequence of state observations. The main idea, taken from *plan recognition as planning* [Ramırez and Geffner, 2009], is to map these distances into likelihoods using the Bayes rule.

According to the Bayes rule, the probability of an hypothesis $\mathcal{H}$ given the observations $\mathcal{O}$ can be computed with $P(\mathcal{H}|\mathcal{O}) = \frac{P(\mathcal{O}|\mathcal{H})P(\mathcal{H})}{P(\mathcal{O})}$. In our scenario, the hypotheses are about the possible STRIPS action models that can be built given a set of operator headers, or in other words, given the $F_v(\xi)$ sets. Moreover, we assume that all the possible STRIPS action models are equiprobable so according to the Bayes rule $P(\mathcal{H}|\mathcal{O})$ is proportional to $P(\mathcal{O}|\mathcal{H})$. Finally, we assume that $P(\mathcal{O}|\mathcal{H})$ is given by the *observation edit distance*, mapping distances into probabilities (the larger the distance the lower the likelihood), according to the following expression $P(\mathcal{O}|\Xi) = 1 - \frac{\delta(\Xi, \mathcal{O})}{\delta(\Xi, *)}$.

Thus, $P(\Xi|\mathcal{O})$, the probability distribution of the possible STRIPS models (within the set of operator headers, that is within the $F_v(\xi)$ sets) given an observation sequence $\mathcal{O}$, is computed by:

1. Calculating the *observation edit distance* $\delta(\Xi, \mathcal{O})$ for every possible model Ξ.
2. Applying the resulting distances to the above formula $P(\mathcal{O}|\Xi)$ to map distances into likelihoods

3. Applying Bayes rule to obtain the normalized posterior probabilities $P(\Xi|\mathcal{O})$. The sum of the probabilities of all the models will be 1.

4.3 STRIPS model evaluation and recognition with classical planning

The previous compilation can be extended to estimate the *observations* STRIPS *edit distance*. In this case the tuple $\Lambda = \langle \Psi, \Xi, \mathcal{O} \rangle$ does not represent a learning task but the edition of a given STRIPS action model Ξ to cover the sequence of observations $\mathcal{O}$. The extended compilation outputs a classical planning task $P'_\Lambda = \langle F_\Lambda, A'_\Lambda, I'_\Lambda, G_\Lambda \rangle$ where:

- F_Λ is defined as in the previous compilation.

- I'_Λ contains the fluents from F that encode s_0 (the first observation) and $mode_{prog}$ set to true. In addition, the given STRIPS action model Ξ is now encoded in the initial state. This means that fluents $pre_f(\xi)/del_f(\xi)/add_f(\xi)$ with $f \in F_v(\xi)$ hold in the initial state if they appear in the given STRIPS action model Ξ.

- G_Λ is defined as in the previous compilation.

- A'_Λ, comprises the three kinds of actions of A_Λ. The Actions for *applying* an already programmed operator schema and the actions for *validating* an observation $1 \le i \le n$ are defined exactly as in the previous compilation. The only difference are the actions for *programming* operator schema $\xi \in \Xi$ that must include now actions for adding a precondition and for removing a *negative* or *positive* effect.

The probability distribution $P(\Xi|\mathcal{O})$ would be exact if the classical planning problems resulting from our compilation are solved optimally (according to the number of edit actions), and will be approximate if they are solved with a satisfying classical planner or furthermore if what is solved is not P'_Λ but a relaxation of this task such as the *delete relaxation*.

5 Evaluation

This section evaluates our approach for learning STRIPS models starting from different amounts of input knowledge.

Reproducibility

We used IPC domains that satisfy the STRIPS requirement [Fox and Long, 2003], taken from the PLANNING.DOMAINS repository [Muise, 2016]. For the learning of the STRIPS action models we used observations sequences of 25 states per domain. All experiments are run on an Intel Core i5 3.10 GHz x 4 with 4 GB of RAM.

MADAGASCAR is the classical planner we use to solve the instances that result from our compilations because its ability to deal with dead-ends [Rintanen, 2014]. In addition, MADAGASCAR can apply the actions for programming preconditions in a single planning step (in parallel) because these actions do not interact. Actions for programming action effects can also be applied in a single planning step reducing significantly the planning horizon.

The compilation source code, the evaluation scripts and the benchmarks are fully available at this anonymous repository *https://github.com/anonsub/strips-learning* so any experimental data reported in the paper can be reproduced.

Supervised evaluation of the learned models

For each domain the learned model is compared with the actual model and its quality is quantified with the *precision* and *recall* metrics. Precision gives a notion of *soundness* while recall gives a notion of the *completeness* of the learned models. $Precision = \frac{tp}{tp+fp}$, where tp is the number of true positives (predicates that correctly appear in the action model) and fp is the number of false positives (predicates appear in the learned action model that should not appear). $Recall = \frac{tp}{tp+fn}$ where fn is the number of false negatives (predicates that should appear in the learned action model but are missing).

When the learning hypothesis space is under constrained, the learned actions can be reformulated and still be compliant with the inputs. For instance in the *blocksworld*, operator `stack` could be *learned* with the preconditions and effects of the `unstack` operator (and vice versa). Furthermore, in a given action the role of the parameters that share the same type can be interchanged making non trivial to compute *precision* and *recall* with respect to a reference model.

To address these issues we defined an evaluation method robust to action reformulation. Precision and recall are often combined using the *harmonic mean*. This expression is called the *F-measure* (or the balanced *F-score*) and is formally defined as $F = 2 \times \frac{Precision \times Recall}{Precision+Recall}$. Given a reference STRIPS action model Ξ^* and the learned STRIPS action model Ξ we define the bijective function $f_{P\&R} : \Xi \mapsto \Xi^*$ such that $f_{P\&R}$ maximizes the accumulated *F-measure*. With this mapping defined we can compute the *precision* and *recall* of a learned STRIPS action $\xi \in \Xi$ with respect to the action $f_{P\&R}(\xi) \in \Xi^*$ even if actions are reformulated in the learning process.

6 Conclusions

As far as we know, this is the first work on learning STRIPS action models from state observations, exclusively using classical planning and evaluated over a wide range of different domains. Recently, Stern and Juba 2017 proposed a classical planning compilation for learning action models but following the *finite domain* representation for the state variables and did not report experimental results since the compilation was not implemented.

The size of the compiled classical planning instances depends on the number of input examples. Generating *informative* examples for learning planning action models is an open issue. Planning actions include preconditions that are only satisfied by specific sequences of actions, often, with a low probability of being chosen by chance [Fern *et al.*, 2004]. The success of recent algorithms for exploring planning tasks [Francés *et al.*, 2017] motivates the development of novel techniques able to autonomously collect informative learning examples. The combination of such exploration techniques with our learning approach is an intriguing re-

	Pre		Add		Del			
	P	R	P	R	P	R	P	R
blocks	0.57	0.44	0.63	0.56	0.57	0.44	0.59	0.48
driverlog	0.2	0.07	0.18	0.29	0.2	0.14	0.19	0.17
ferry	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
floor-tile	0.5	0.14	0.75	0.55	0.6	0.27	0.62	0.32
Grid	-	-	-	-	-	-	-	-
gripper-strips	0.25	0.17	0.25	0.25	0.25	0.25	0.25	0.22
hanoi	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
hiking	-	-	-	-	-	-	-	-
n-puzzle	-	-	-	-	-	-	-	-
parking	0.6	0.21	0.14	0.11	0.4	0.22	0.38	0.18
satellite	0.25	0.07	0.33	0.4	0.67	0.5	0.42	0.32
Sokoban	-	-	-	-	-	-	-	-
transport	0.33	0.1	0.33	0.4	0.0	0.0	0.22	0.17
zeno-travel	0.25	0.07	0.17	0.14	0.0	0.0	0.14	0.07
	-	-	-	-	-	-	-	-

Table 1: Precision and recall values obtained when learning from labels without computing the $f_{P\&R}$ mapping.

	Pre		Add		Del			
	P	R	P	R	P	R	P	R
blocks	0.71	0.56	0.75	0.67	0.71	0.56	0.73	0.59
driverlog	0.6	0.21	0.27	0.43	0.6	0.43	0.49	0.36
ferry	0.5	0.29	0.75	0.75	0.5	0.5	0.58	0.51
floor-tile	0.5	0.14	0.75	0.55	0.6	0.27	0.62	0.32
Grid	-	-	-	-	-	-	-	-
gripper-strips	1.0	0.67	1.0	1.0	1.0	1.0	1.0	0.89
hanoi	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
hiking	-	-	-	-	-	-	-	-
n-puzzle	-	-	-	-	-	-	-	-
parking	0.6	0.21	0.14	0.11	0.4	0.22	0.38	0.18
satellite	0.25	0.07	0.33	0.4	0.67	0.5	0.42	0.32
Sokoban	-	-	-	-	-	-	-	-
transport	1.0	0.3	0.67	0.8	0.67	0.4	0.78	0.5
zeno-travel	0.5	0.14	0.5	0.43	0.33	0.14	0.44	0.24
	-	-	-	-	-	-	-	-

Table 2: Precision and recall values obtained when learning from labels but computing the $f_{P\&R}$ mapping.

search direction that opens the door to the bootstrapping of planning action models.

References

[Amir and Chang, 2008] Eyal Amir and Allen Chang. Learning partially observable deterministic action models. *Journal of Artificial Intelligence Research*, 33:349–402, 2008.

[Bonet *et al.*, 2009] Blai Bonet, Héctor Palacios, and Héctor Geffner. Automatic derivation of memoryless policies and finite-state controllers using classical planners. In *ICAPS*, 2009.

[Cresswell *et al.*, 2013] Stephen N Cresswell, Thomas L McCluskey, and Margaret M West. Acquiring planning domain models using LOCM. *The Knowledge Engineering Review*, 28(02):195–213, 2013.

[Fern *et al.*, 2004] Alan Fern, Sung Wook Yoon, and Robert Givan. Learning domain-specific control knowledge from random walks. In *ICAPS*, pages 191–199, 2004.

[Fox and Long, 2003] Maria Fox and Derek Long. PDDL2.1: An extension to PDDL for expressing temporal planning domains. *J. Artif. Intell. Res.(JAIR)*, 20:61–124, 2003.

	Total time	Preprocess	Plan length
Blocks	1.40	0.00	70
Driverlog	1.50	0.00	89
Ferry	1.49	0.00	64
Floortile	351.38	0.11	156
Grid	-	-	-
Gripper	0.04	0.00	59
Hanoi	2.33	0.01	49
Hiking	-	-	-
Parking	-	-	-
Satellite	78.36	0.03	98
Sokoban	-	-	-
Transport	285.61	0.10	106
Zenotravel	6.20	0.46	71

Table 3: Planning results obtained when learning from labels.

[Francés *et al.*, 2017] Guillem Francés, Miquel Ramírez, Nir Lipovetzky, and Hector Geffner. Purely declarative action representations are overrated: Classical planning with simulators. In *International Joint Conference on Artificial Intelligence*, 2017.

[Geffner and Bonet, 2013] Hector Geffner and Blai Bonet. A concise introduction to models and methods for automated planning, 2013.

[Ghallab *et al.*, 2004] Malik Ghallab, Dana Nau, and Paolo Traverso. *Automated Planning: theory and practice*. Elsevier, 2004.

[Kambhampati, 2007] Subbarao Kambhampati. Model-lite planning for the web age masses: The challenges of planning with incomplete and evolving domain models. In *Proceedings of the National Conference on Artificial Intelligence*, 2007.

[McDermott *et al.*, 1998] Drew McDermott, Malik Ghallab, Adele Howe, Craig Knoblock, Ashwin Ram, Manuela Veloso, Daniel Weld, and David Wilkins. PDDL – The Planning Domain Definition Language, 1998.

[Muise, 2016] Christian Muise. Planning. domains. *ICAPS system demonstration*, 2016.

[Ramırez and Geffner, 2009] Miquel Ramırez and Hector Geffner. Plan recognition as planning. In *Proceedings of the 21st international joint conference on Artifical intelligence. Morgan Kaufmann Publishers Inc*, pages 1778–1783, 2009.

[Ramírez, 2012] Miquel Ramírez. *Plan recognition as planning*. PhD thesis, Universitat Pompeu Fabra, 2012.

[Rintanen, 2014] Jussi Rintanen. Madagascar: Scalable planning with sat. *Proceedings of the 8th International Planning Competition (IPC-2014)*, 2014.

[Segovia-Aguas *et al.*, 2016] Javier Segovia-Aguas, Sergio Jiménez, and Anders Jonsson. Hierarchical finite state controllers for generalized planning. In *Proceedings of the Twenty-Fifth International Joint Conference on Artificial Intelligence*, pages 3235–3241. AAAI Press, 2016.

[Segovia-Aguas *et al.*, 2017] Javier Segovia-Aguas, Sergio Jiménez, and Anders Jonsson. Generating context-free grammars using classical planning. In *International Joint Conference on Artificial Intelligence*, 2017.

[Slaney and Thiébaux, 2001] John Slaney and Sylvie Thiébaux. Blocks world revisited. *Artificial Intelligence*, 125(1-2):119–153, 2001.

[Stern and Juba, 2017] Roni Stern and Brendan Juba. Efficient, safe, and probably approximately complete learning of action models. *IJCAI*, 2017.

[Yang *et al.*, 2007] Qiang Yang, Kangheng Wu, and Yunfei Jiang. Learning action models from plan examples using weighted max-sat. *Artificial Intelligence*, 171(2-3):107–143, 2007.

[Zhuo and Kambhampati, 2013] Hankz Hankui Zhuo and Subbarao Kambhampati. Action-model acquisition from noisy plan traces. In *IJCAI*, pages 2444–2450, 2013.