

Computing the *least-commitment* action model from state observations

Diego Aineto¹, Sergio Jiménez¹, Eva Onaindia¹ and , Blai Bonet²

¹Departamento de Sistemas Informáticos y Computación. Universitat Politècnica de València. Valencia, Spain

²Departamento de Computación. Universidad Simón Bolívar. Caracas, Venezuela

{dieaigar,serjice,onaindia}@dsic.upv.es, bonet@usb.ve

Abstract

1 Introduction

Given a sequence of partially observed states, this paper formalizes the task of computing the *least-commitment* action model that is able to explain the given observation. This task is of interest because it allows the incremental learning of action models from arbitrary large sets of partial observations.

In addition, the paper introduces a new method to compute the *least-commitment* action model from a sequence of partially observed states. The method assumes that action models are specified as STRIPS action schema and it builds on top of off-the-shelf algorithms for *conformant planning*.

2 Background

This section formalizes the planning models we use in the paper as well as the kind of input observations for the computation of the *least-commitment* action model.

2.1 Classical planning with conditional effects

Let F be the set of *fluents* or *state variables* (propositional variables). A *literal* l is a valuation of a fluent $f \in F$, i.e. either $l = f$ or $l = \neg f$. L is a set of literals that represents a partial assignment of values to fluents, and $\mathcal{L}(F)$ is the set of all literals sets on F , i.e. all partial assignments of values to fluents. A *state* s is a full assignment of values to fluents. We explicitly include negative literals $\neg f$ in states s.t. $|s| = |F|$ and the size of the state space is $2^{|F|}$.

A *planning frame* is a tuple $\Phi = \langle F, A \rangle$, where F is a set of fluents and A is a set of *actions*. An action $a \in A$ is defined with *preconditions*, $\text{pre}(a) \in \mathcal{L}(F)$, *positive effects*, $\text{eff}^+(a) \in \mathcal{L}(F)$, and *negative effects* $\text{eff}^-(a) \in \mathcal{L}(F)$. The semantics of actions $a \in A$ is specified with two functions: $\rho(s, a)$ denotes whether action a is *applicable* in a state s and $\theta(s, a)$ denotes the *successor state* that results of applying action a in a state s . Then, $\rho(s, a)$ holds iff $\text{pre}(a) \subseteq s$. And the result of applying a in s is $\theta(s, a) = \{s \setminus \text{eff}^-(a)\} \cup \text{eff}^+(a)$.

A *classical planning problem* is a tuple $P = \langle F, A, I, G \rangle$, where I is the initial state and $G \in \mathcal{L}(F)$ is the set of goal

conditions over the state variables. A *plan* π is an action sequence $\pi = \langle a_1, \dots, a_n \rangle$, with $|\pi| = n$ denoting its *plan length*. The execution of π in the initial state I of P induces a *trajectory* $\tau(\pi, s_0) = \langle s_0, a_1, s_1, \dots, a_n, s_n \rangle$ such that $s_0 = I$ and, for each $1 \leq i \leq n$, it holds $\rho(s_{i-1}, a_i)$ and $s_i = \theta(s_{i-1}, a_i)$. A plan π solves P iff the induced *trajectory* $\tau(\pi, s_0)$ reaches a final state $G \subseteq s_n$. A solution plan is optimal iff its length is minimal.

An action $a_c \in A$ with conditional effects is defined as a set of preconditions $\text{pre}(a_c) \in \mathcal{L}(F)$ and a set of *conditional effects* $\text{cond}(a_c)$. Each conditional effect $C \triangleright E \in \text{cond}(a_c)$ is composed of two sets of literals: $C \in \mathcal{L}(F)$, the *condition*, and $E \in \mathcal{L}(F)$, the *effect*. An action a_c is applicable in a state s if $\rho(s, a_c)$ is true, and the *triggered effects* resulting from the action application are the effects whose conditions hold in s :

$$\text{triggered}(s, a_c) = \bigcup_{C \triangleright E \in \text{cond}(a_c), C \subseteq s} E,$$

The result of applying action a_c in state s is $\theta(s, a_c) = \{s \setminus \text{eff}_c^-(s, a) \cup \text{eff}_c^+(s, a)\}$, where $\text{eff}_c^-(s, a) \subseteq \text{triggered}(s, a)$ and $\text{eff}_c^+(s, a) \subseteq \text{triggered}(s, a)$ are, respectively, the triggered *negative* and *positive* effects.

2.2 The observation model

Given a classical planning problem $P = \langle F, A, I, G \rangle$, a plan π and a trajectory $\tau(\pi, s_0)$, we define the *observation of the trajectory* as a sequence of partial states coming from observing the execution of π in P . Formally, $\mathcal{O}(\tau) = \langle s_0^o, s_1^o, \dots, s_m^o \rangle$ where $s_0^o = I$.

A partially observable state s_i^o is one in which $|s_i^o| < |F|$; i.e., a state in which at least a fluent of F is not observable. Note that this definition also comprises the case $|s_i^o| = 0$, when the state is fully unobservable. Whatever the sequence of observed states of $\mathcal{O}(\tau)$ is, it must be consistent with the sequence of states of $\tau(\pi, s_0)$, meaning that $\forall i, s_i^o \subseteq s_i$. In practice, the number of observed states, m , ranges from 1 (the initial state, at least), to $|\pi| + 1$, and the observed intermediate states will comprise a number of fluents between $[1, |F|]$.

In other words, we assume there is a bijective monotone mapping between trajectories and observations [Ramírez and Geffner, 2009], thus also granting the inverse consistency relationship (the trajectory is a superset of the observation). Therefore, transiting between two consecutive ob-

served states in $\mathcal{O}(\tau)$ may require the execution of more than a single action ($\theta(s_i^o, \langle a_1, \dots, a_k \rangle) = s_{i+1}^o$, where $k \geq 1$ is unknown but finite. In other words, having $\mathcal{O}(\tau)$ does not imply knowing the actual length of π .

2.3 Conformant planning

Conformant planning is planning with incomplete information about the initial state, no sensing, and validating that goals are achieved with certainty (despite the uncertainty of the initial state) [Smith and Weld, 1998; Goldman and Boddy, 1996].

Syntactically, conformant planning problems are expressed in compact form through a set of state variables. A *conformant planning problem* can be defined as a tuple $P_c = \langle F, A, \Upsilon, G \rangle$ where F , A and G are the set of *fluents*, *actions* and *goals* (as previously defined for *classical planning*). Now Υ is a set of clauses over literals $l = f$ or $l = \neg f$ (for $f \in F$) that define the set of possible initial states.

A solution to a conformant planning problem is an action sequence that maps each possible initial state into a goal state. More precisely, an action sequence $\pi = \langle a_1, \dots, a_n \rangle$ is a *conformant plan* for P_c iff for each possible *trajectory* $\tau(\pi, s_0) = \langle s_0, a_1, s_1, \dots, a_n, s_n \rangle$, such that s_0 is a valuation of the fluents in F that satisfies Υ , then $\tau(\pi, s_0)$ reaches a final state $G \subseteq s_n$ where all goal conditions are met.

3 Computing the least-commitment action model from state observations

First, this section formalizes the notion of the *least-commitment* action model that is able to *explain* a sequence of partially observed states. Next, the section describes our approach to compute such model via *conformant planning*.

3.1 The least-commitment action model

The task of computing the *least-commitment* action model from state observations is defined as a $\langle P, \mathcal{O} \rangle$ pair:

- $P = \langle F, A[\cdot], I, G \rangle$ is a planning problem where $A[\cdot]$ is a set of actions s.t. for each $a \in A[\cdot]$, the semantics of a is unknown; i.e. the corresponding $\langle \rho, \theta \rangle$ functions are undefined.
- $\mathcal{O}(\tau)$ is a sequence of partial states coming from the observation of a trajectory $\tau(\pi, s_0)$ produced by the execution of certain unknown plan π that solves P .

Before formalizing the solution to this task, i.e. the *least-commitment* action model, we introduce several necessary definitions. We first start defining a *partially specified action model* inspired by the notion of *incomplete (annotated) model* [Sreedharan et al., 2018].

Definition 1 (Partially specified action model) *Given a set of actions $A[\cdot]$ and a set of fluents F then, a partially specified action model M is a set of possible models for the actions in $A[\cdot]$ such that: (1), any model $\mathcal{M} \in M$ defines the $\langle \rho, \theta \rangle$ functions of every action in $A[\cdot]$ and (2), for every $\mathcal{M} \in M$ the $\langle \rho, \theta \rangle$ functions are defined in the set of state variables F . (Note that if M is a singleton it represents a fully specified action model).*

Definition 2 (Explanation of an observation) *Given a fully specified action model M for the actions $A[\cdot]$ in P , and a sequence of partially observed states $\mathcal{O}(\tau)$, we say that the model explains the observation (denoted $\mathcal{M} \mapsto \mathcal{O}(\tau)$) iff there exists a plan π , consistent with $\mathcal{O}(\tau)$, that solves P . We say that π is the best explanation for $\mathcal{O}(\tau)$ iff, in addition, π is optimal.*

Now we are ready to define the *least-commitment* action model for an observation $\mathcal{O}(\tau)$.

Definition 3 (The least-commitment action model) *Given a partially specified action model M that represents the space of possible action models and a sequence of partially observed states $\mathcal{O}(\tau)$, then the least-commitment action model is another partially specified action model that represents the largest subset of models $M^* \subseteq M$ such that every model $\mathcal{M} \in M^*$ explains the input observation.*

3.2 The space of STRIPS action models

Despite previous definitions are general, this work focuses on the particular kind of action models that are specified as STRIPS action schemata.

A STRIPS *action schema* ξ is defined by four lists: A list of *parameters* $\text{pars}(\xi)$, and three list of predicates (namely $\text{pre}(\xi)$, $\text{del}(\xi)$ and $\text{add}(\xi)$) that shape the kind of fluents that can appear in the *preconditions*, *negative effects* and *positive effects* of the actions induced from that schema. Let be Ψ the set of *predicates* that shape the propositional state variables F , and a list of *parameters* $\text{pars}(\xi)$. The set of elements that can appear in $\text{pre}(\xi)$, $\text{del}(\xi)$ and $\text{add}(\xi)$ of the STRIPS action schema ξ is given by FOL interpretations of Ψ over the parameters $\text{pars}(\xi)$. We denote this set of FOL interpretations as $\mathcal{I}_{\Psi, \xi}$.

Despite any element of $\mathcal{I}_{\Psi, \xi}$ can *a priori* appear in the $\text{pre}(\xi)$, $\text{del}(\xi)$ and $\text{add}(\xi)$ of schema ξ , the space of possible STRIPS schemata is constrained by a set \mathcal{C} that includes:

1. *Syntactic constraints.* STRIPS constraints require $\text{del}(\xi) \subseteq \text{pre}(\xi)$, $\text{del}(\xi) \cap \text{add}(\xi) = \emptyset$ and $\text{pre}(\xi) \cap \text{add}(\xi) = \emptyset$. Considering exclusively these syntactic constraints, the size of the space of possible STRIPS schemata is given by $2^{2 \times |\mathcal{I}_{\Psi, \xi}|}$.
2. *Domain-specific constraints.* One can introduce domain-specific knowledge to constrain further the space of possible schemata. For instance, in the *blocksworld* one can argue that $\text{on}(v_1, v_1)$ and $\text{on}(v_2, v_2)$ will not appear in the $\text{pre}(\xi)$, $\text{del}(\xi)$ and $\text{add}(\xi)$ lists of an action schema ξ because, in this particular domain, only one block can be on top of another block. As a rule of thumb, *state invariants* constraining the possible states of a given planning domain belong to this second class of constraints [Fox and Long, 1998].

Definition 4 (Well-defined STRIPS action schemata)

*Given a set of predicates Ψ , a list of action parameters $\text{pars}(\xi)$, and set of FOL constraints \mathcal{C} , ξ is a **well-defined STRIPS action schema** iff its three lists $\text{pre}(\xi) \subseteq \mathcal{I}_{\Psi, \xi}$, $\text{del}(\xi) \subseteq \mathcal{I}_{\Psi, \xi}$ and $\text{add}(\xi) \subseteq \mathcal{I}_{\Psi, \xi}$ only contain elements in $\mathcal{I}_{\Psi, \xi}$ and they satisfy all the constraints in \mathcal{C} .*

```

(:action stack
  :parameters (?v1 ?v2)
  :precondition (and (holding ?v1) (clear ?v2))
  :effect (and (not (holding ?v1)) (not (clear ?v2))
              (clear ?v1) (handempty) (on ?v1 ?v2)))

(pre_holding_v1_stack) (pre_clear_v2_stack)
(eff_holding_v1_stack) (eff_clear_v2_stack)
(eff_clear_v1_stack) (eff_handempty_stack) (eff_on_v1_v2_stack)

```

Figure 1: PDDL encoding of the `stack(?v1, ?v2)` schema and our propositional representation for this same schema.

We say a planning model \mathcal{M} is *well-defined* if all its STRIPS action schemata are *well-defined*.

3. *Observation constraints.* A sequence of state observations $\mathcal{O}(\tau)$ constraints further the space of possible action schemata. This *semantic knowledge* included in the observations introduce a third type of constraints and can also be added to the set \mathcal{C} .

3.3 Computing the *least-commitment* model via conformant planning

Given the set of actions $A[\cdot]$ with unknown ρ and θ functions and a sequence of partial states $\mathcal{O}(\tau)$, we can build a classical planning problem $P = \langle F, A[\cdot], I, G \rangle$ such that $\mathcal{O}(\tau)$ represents the observation of a $\tau(\pi, I)$ trajectory that solves P . In this section we show that starting from $\mathcal{O}(\tau)$ and P we can build a *conformant planning problem* P_c whose solution induces the *least-commitment* action model for the input observation $\mathcal{O}(\tau)$.

In more detail, we build a *conformant planning problem* $P_c = \langle F_c, A_c, \Upsilon, G \rangle$ such that:

- The set of fluents F_c extends F with two new sets of fluents:
 - $\{test_j\}_{1 \leq j \leq m}$, indicating the state observation $s_j \in \mathcal{O}(\tau)$ where the action model is validated
 - The fluents $pre_e_ \xi$ and $eff_e_ \xi$ (where $e \in \mathcal{I}_{\Psi, \xi}$) implementing a propositional encoding of the *pre-conditions*, *negative*, and *positive* effects of an action schema ξ . Our encoding exploits the syntactic constraint of STRIPS so, if $pre_e_ \xi$ and $eff_e_ \xi$ holds it means that $e \in \mathcal{I}_{\Psi, \xi}$ is a negative effect in ξ and if $pre_e_ \xi$ does not hold but $eff_e_ \xi$ holds, it means that $e \in \mathcal{I}_{\Psi, \xi}$ is a positive effect in ξ . Figure 1 shows the PDDL encoding of the `stack(?v1, ?v2)` schema and our propositional representation for this same schema.
- The set of actions A_c contains now actions of three different kinds:
 - Actions for *committing* $pre_e_ \xi$ fluents to a positive/negative value (similar actions are also defined for *committing* $eff_e_ \xi$ fluents to a posi-

tive/negative value).

$$\begin{aligned}
 pre(commit_T_pre_e_ \xi) &= \{mode_{commit}\}, \\
 cond(commit_T_pre_e_ \xi) &= \{pre_e_ \xi\} \triangleright \{pre_e_ \xi\}, \\
 &\quad \{\neg pre_e_ \xi\} \triangleright \{pre_e_ \xi\}. \\
 pre(commit_ \perp_pre_e_ \xi) &= \{mode_{commit}\}, \\
 cond(commit_ \perp_pre_e_ \xi) &= \{pre_e_ \xi\} \triangleright \{\neg pre_e_ \xi\}, \\
 &\quad \{\neg pre_e_ \xi\} \triangleright \{\neg pre_e_ \xi\}.
 \end{aligned}$$

- Actions for *validating* that committed models explain the s_j observed states, $0 \leq j < m$.

$$\begin{aligned}
 pre(validate_j) &= s_j \cup \{test_{j-1}\}, \\
 cond(validate_j) &= \{\emptyset\} \triangleright \{\neg test_{j-1}, test_j\}, \\
 &\quad \{mode_{commit}\} \triangleright \{\neg mode_{commit}, mode_{val}\}.
 \end{aligned}$$

- *Editable* actions whose semantics is given by the value of the $pre_e_ \xi$, $eff_e_ \xi$ fluents at the current state. Figure 2 shows the PDDL encoding of an *editable* `stack(?v1, ?v2)` schema. Note that this *editable* schema when


```

(pre_holding_v1_stack) (pre_clear_v2_stack)
(eff_holding_v1_stack) (eff_clear_v2_stack)
(eff_clear_v1_stack) (eff_handempty_stack)
(eff_on_v1_v2_stack)

```

holds, it behaves exactly as the original PDDL schema defined in Figure 1. Formally, given an operator schema $\xi \in \mathcal{M}$ its *editable* version is:

$$\begin{aligned}
 pre(editable_\xi) &= \{pre_e_ \xi \implies e\}_{e \in \mathcal{I}_{\Psi, \xi}}, \\
 cond(editable_\xi) &= \{pre_e_ \xi, eff_e_ \xi\} \triangleright \{\neg e\}_{e \in \mathcal{I}_{\Psi, \xi}}, \\
 &\quad \{\neg pre_e_ \xi, eff_e_ \xi\} \triangleright \{e\}_{e \in \mathcal{I}_{\Psi, \xi}}.
 \end{aligned}$$

- The clauses in Υ comprises:

1. The *unit clauses* given by the fluents that hold in the initial state $I = s_0$ and $mode_{commit}$ set to true.
2. The clauses representing that the actual value of fluents $pre_e_ \xi, eff_e_ \xi$ is unknown. In other words, that any model from the STRIPS space of models can initially be part of the *least-commitment* action model. Formally, for every ξ and $e \in \mathcal{I}_{\Psi, \xi}$, then Υ includes these two clauses:

$$\begin{aligned}
 &pre_e_ \xi \text{ xor } \neg pre_e_ \xi. \\
 &eff_e_ \xi \text{ xor } \neg eff_e_ \xi.
 \end{aligned}$$

- The new goals are $G_c = \{test_m\}$.

3.4 Compilation properties

4 Evaluation

5 Conclusions

[Stern and Juba, 2017]

```

(:action stack
:parameters (?o1 - object ?o2 - object)
:precondition
  (and (or (not (pre_on_v1_v1_stack)) (on ?o1 ?o1))
        (or (not (pre_on_v1_v2_stack)) (on ?o1 ?o2))
        (or (not (pre_on_v2_v1_stack)) (on ?o2 ?o1))
        (or (not (pre_on_v2_v2_stack)) (on ?o2 ?o2))
        (or (not (pre_ontable_v1_stack)) (ontable ?o1))
        (or (not (pre_ontable_v2_stack)) (ontable ?o2))
        (or (not (pre_clear_v1_stack)) (clear ?o1))
        (or (not (pre_clear_v2_stack)) (clear ?o2))
        (or (not (pre_holding_v1_stack)) (holding ?o1))
        (or (not (pre_holding_v2_stack)) (holding ?o2))
        (or (not (pre_handempty_stack)) (handempty))))
:effect (and
  (when (and (pre_on_v1_v1_stack) (eff_on_v1_v1_stack)) (not (on ?o1 ?o1)))
  (when (and (pre_on_v1_v2_stack) (eff_on_v1_v2_stack)) (not (on ?o1 ?o2)))
  (when (and (pre_on_v2_v1_stack) (eff_on_v2_v1_stack)) (not (on ?o2 ?o1)))
  (when (and (pre_on_v2_v2_stack) (eff_on_v2_v2_stack)) (not (on ?o2 ?o2)))
  (when (and (pre_ontable_v1_stack) (eff_ontable_v1_stack)) (not (ontable ?o1)))
  (when (and (pre_ontable_v2_stack) (eff_ontable_v2_stack)) (not (ontable ?o2)))
  (when (and (pre_clear_v1_stack) (eff_clear_v1_stack)) (not (clear ?o1)))
  (when (and (pre_clear_v2_stack) (eff_clear_v2_stack)) (not (clear ?o2)))
  (when (and (pre_holding_v1_stack) (eff_holding_v1_stack)) (not (holding ?o1)))
  (when (and (pre_holding_v2_stack) (eff_holding_v2_stack)) (not (holding ?o2)))
  (when (and (pre_handempty_stack) (eff_handempty_stack)) (not (handempty)))
  (when (and (not (pre_on_v1_v1_stack)) (eff_on_v1_v1_stack)) (on ?o1 ?o1))
  (when (and (not (pre_on_v1_v2_stack)) (eff_on_v1_v2_stack)) (on ?o1 ?o2))
  (when (and (not (pre_on_v2_v1_stack)) (eff_on_v2_v1_stack)) (on ?o2 ?o1))
  (when (and (not (pre_on_v2_v2_stack)) (eff_on_v2_v2_stack)) (on ?o2 ?o2))
  (when (and (not (pre_ontable_v1_stack)) (eff_ontable_v1_stack)) (ontable ?o1))
  (when (and (not (pre_ontable_v2_stack)) (eff_ontable_v2_stack)) (ontable ?o2))
  (when (and (not (pre_clear_v1_stack)) (eff_clear_v1_stack)) (clear ?o1))
  (when (and (not (pre_clear_v2_stack)) (eff_clear_v2_stack)) (clear ?o2))
  (when (and (not (pre_holding_v1_stack)) (eff_holding_v1_stack)) (holding ?o1))
  (when (and (not (pre_holding_v2_stack)) (eff_holding_v2_stack)) (holding ?o2))
  (when (and (not (pre_handempty_stack)) (eff_handempty_stack)) (handempty))))

```

Figure 2: PDDL encoding of the editable version of the stack(?v1, ?v2) schema.

References

- [Fox and Long, 1998] Maria Fox and Derek Long. The automatic inference of state invariants in tim. *Journal of Artificial Intelligence Research*, 9:367–421, 1998.
- [Goldman and Boddy, 1996] Robert P Goldman and Mark S Boddy. Expressive planning and explicit knowledge. In *AIPS*, volume 96, pages 110–117, 1996.
- [Ramírez and Geffner, 2009] Miquel Ramírez and Hector Geffner. Plan recognition as planning. In *International Joint conference on Artificial Intelligence, (IJCAI-09)*, pages 1778–1783. AAAI Press, 2009.
- [Smith and Weld, 1998] David E Smith and Daniel S Weld. Conformant graphplan. In *AAAI/IAAI*, pages 889–896, 1998.
- [Sreedharan *et al.*, 2018] Sarath Sreedharan, Tathagata Chakraborti, and Subbarao Kambhampati. Handling model uncertainty and multiplicity in explanations via model reconciliation. In *International Conference on Automated Planning and Scheduling, (ICAPS-18)*, pages 518–526, 2018.
- [Stern and Juba, 2017] Roni Stern and Brendan Juba. Efficient, safe, and probably approximately complete learning of action models. In *International Joint Conference on Artificial Intelligence, (IJCAI-17)*, pages 4405–4411, 2017.