Computing *the least commitment* action model from state observations

Diego Aineto[1] , **Sergio Jiménez**[1] , **Eva Onaindia**[1] and , **Blai Bonet**[2]

[1]Departamento de Sistemas Informáticos y Computación. Universitat Politècnica de València. Valencia, Spain
[2]Departamento de Computación. Universidad Simón Bolívar. Caracas, Venezuela

{dieaigar,serjice,onaindia}@dsic.upv.es, bonet@usb.ve

Abstract

Given a sequence of partially observed states, this paper formalizes the task of computing *the least commitment* action model that can explain the given observation. This task is of interest because it allows the incremental learning of action models from arbitrary large sets of state observations.

In addition, the paper introduces a new method to compute *the least commitment* action model from a sequence of state observations. The method assumes that action models are specified as STRIPS action schema and it builds on top of off-the-shelf algorithms for *conformant planning*.

1 Introduction

Given a sequence of partially observed states, this paper formalizes the task of computing *the least commitment* action model that can explain the given observation. This task is of interest because it allows the incremental learning of action models from arbitrary large sets of state observations.

In addition, the paper introduces a new method to compute *the least commitment* action model from a sequence of state observations. The method assumes that action models are specified as STRIPS action schema and it builds on top of off-the-shelf algorithms for *conformant planning*.

2 Background

This section formalizes the model of *classical planning with conditional effects*, the *conformant planning* model, and the kind of input observations for the computation of *the least commitment* action model.

2.1 Classical planning with conditional effects

Let F be the set of *fluents* or *state variables* (propositional variables). A *literal* l is a valuation of a fluent $f \in F$, i.e. either $l = f$ or $l = \neg f$. L is a set of literals that represents a partial assignment of values to fluents, and $\mathcal{L}(F)$ is the set of all literals sets on F, i.e. all partial assignments of values to fluents. A *state* s is a full assignment of values to fluents. We explicitly include negative literals $\neg f$ in states s.t. $|s| = |F|$ and the size of the state space is $2^{|F|}$.

A *planning frame* is a tuple $\Phi = \langle F, A \rangle$, where F is a set of fluents and A is a set of *actions*. An action $a \in A$ is defined with *preconditions*, $\text{pre}(a) \in \mathcal{L}(F)$, *positive effects*, $\text{eff}^+(a) \in \mathcal{L}(F)$, and *negative effects* $\text{eff}^-(a) \in \mathcal{L}(F)$. The semantics of actions $a \in A$ is specified with two functions: $\rho(s, a)$ denotes whether action a is *applicable* in a state s and $\theta(s, a)$ denotes the *successor state* that results of applying action a in a state s. Then, $\rho(s, a)$ holds iff $\text{pre}(a) \subseteq s$. And the result of applying a in s is $\theta(s, a) = \{s \setminus \text{eff}^-(a)) \cup \text{eff}^+(a)\}$.

A *classical planning problem* is a tuple $P = \langle F, A, I, G \rangle$, where I is the initial state and $G \in \mathcal{L}(F)$ is the set of goal conditions over the state variables. A *plan* π is an action sequence $\pi = \langle a_1, \ldots, a_n \rangle$, with $|\pi| = n$ denoting its *plan length*. The execution of π in the initial state I of P induces a *trajectory* $\tau(\pi, s_0) = \langle s_0, a_1, s_1, \ldots, a_n, s_n \rangle$ such that $s_0 = I$ and, for each $1 \leq i \leq n$, it holds $\rho(s_{i-1}, a_i)$ and $s_i = \theta(s_{i-1}, a_i)$. A plan π solves P iff the induced *trajectory* $\tau(\pi, s_0)$ reaches a final state $G \subseteq s_n$. A solution plan is optimal iff its length is minimal.

An action $a_c \in A$ with conditional effects is defined as a set of preconditions $\text{pre}(a_c) \in \mathcal{L}(F)$ and a set of *conditional effects* $\text{cond}(a_c)$. Each conditional effect $C \triangleright E \in \text{cond}(a_c)$ is composed of two sets of literals: $C \in \mathcal{L}(F)$, the *condition*, and $E \in \mathcal{L}(F)$, the *effect*. An action a_c is applicable in a state s if $\rho(s, a_c)$ is true, and the *triggered effects* resulting from the action application are the effects whose conditions hold in s:

$$triggered(s, a_c) = \bigcup_{C \triangleright E \in \text{cond}(a_c), C \subseteq s} E,$$

The result of applying action a_c in state s is $\theta(s, a_c) = \{s \setminus \text{eff}_c^-(s, a)) \cup \text{eff}_c^+(s, a)\}$, where $\text{eff}_c^-(s, a) \subseteq triggered(s, a)$ and $\text{eff}_c^+(s, a) \subseteq triggered(s, a)$ are, respectively, the triggered *negative* and *positive* effects.

2.2 The observation model

Given a classical planning problem $P = \langle F, A, I, G \rangle$, a plan π and a trajectory $\tau(\pi, s_0)$, we define the *observation of the trajectory* as sequence of partial states that represents the observation from the execution of π in P. Formally, $\mathcal{O}(\tau) = \langle s_0^o, s_1^o \ldots, s_m^o \rangle$ where $s_0^o = I$.

A partially observable state s_i^o is one in which $|s_i^o| < |F|$; i.e., a state in which at least a fluent of F is not observable. Note that this definition also comprises the case $|s_i^o| = 0$, when the state is fully unobservable. Whatever the sequence of observed states of $\mathcal{O}(\tau)$ is, it must be consistent with the sequence of states of $\tau(\pi, s_0)$, meaning that $\forall i, s_i^o \subseteq s_i$. In practice, the number of observed states, m, ranges from 1 (the initial state, at least), to $|\pi|+1$, and the observed intermediate states will comprise a number of fluents between $[1, |F|]$.

We assume a bijective monotone mapping between trajectories and observations [Ramírez and Geffner, 2009], thus also granting the inverse consistency relationship (the trajectory is a superset of the observation). Therefore, transiting between two consecutive observed states in $\mathcal{O}(\tau)$

may require the execution of more than a single action $(\theta(s_i^o, \langle a_1, \ldots, a_k \rangle) = s_{i+1}^o$, where $k \geq 1$ is unknown but finite. In other words, having $\mathcal{O}(\tau)$ does not imply knowing the actual length of π.

2.3 Conformant planning

Conformant planning is planning with incomplete information about the initial state, no sensing, and where goals have to be achieved with certainty (despite this uncertainty of the initial state) [Smith and Weld, 1998; Goldman and Boddy, 1996].

Syntactically, conformant planning problems are expressed in compact form through a set of state variables. A *conformant planning problem* can be defined as a tuple $P_c = \langle F, A, \Upsilon, G \rangle$ where F, A and G are the set of *fluents*, *actions* and *goals* (as previously defined for *classical planning*). Now Υ is a set of clauses over literals $l = f$ or $l = \neg f$ (for $f \in F$) that define the set of possible initial states.

A solution to a conformant planning problem is an action sequence that maps each possible initial state into a goal state. More precisely, an action sequence $\pi = \langle a_1, \ldots, a_n \rangle$ is a *conformant plan* for P_c iff for each possible *trajectory* $\tau(\pi, s_0) = \langle s_0, a_1, s_1, \ldots, a_n, s_n \rangle$, such that s_0 is a valuation of the fluents in F that satisfies Υ, then the $\tau(\pi, s_0)$ trajectory reaches a final state $G \subseteq s_n$ in which all the goal conditions are met.

3 Computing *the least commitment* action model from state observations

This section formalizes the notion of *the least commitment action model* that is able to *explain* a sequence of state observations and our approach to compute such model via *conformant planning*.

3.1 The *least commitment* action model

The task of computing *the least commitment* action model is defined as a tuple $\langle P, M, \mathcal{O} \rangle$ where:

- $P = \langle F, A[\cdot], I, G \rangle$ is a planning problem where $A[\cdot]$ is a set of actions. For each $a \in A[\cdot]$, the semantics of a is unknown; i.e. the functions ρ and/or θ of a are undefined.

- M is the set of possible models for the actions in $A[\cdot]$. A model $\mathcal{M} \in M$ defines the semantics of every action in $A[\cdot]$. Planning models $\mathcal{M} \in M$ differ in the $\langle \rho, \theta \rangle$ functions of the actions but they all use the same set of state variables F.

- $\mathcal{O}(\tau)$ is a sequence of partial states coming from the observation of a trajectory $\tau(\pi, s_0)$ produced by the execution of certain unknown plan π that solves P.

The *solution* to this task is *the least commitment* action model but, before formalizing it, we introduce the following necessary definition.

Definition 1 (Explanation of an observation) *Given a model $\mathcal{M} \in M$ for the actions $A[\cdot]$ in P, and a sequence of partially observed states $\mathcal{O}(\tau)$, we say that the model explains the observation, that we denote by $\mathcal{M} \mapsto \mathcal{O}(\tau)$, iff*

there exists a plan π that solves P and that it is consistent with π. We say that π is the best explanation for $\mathcal{O}(\tau)$ iff, in additon, π is optimal.

Now we are ready to define *the least commitment* action model as the largest subset of models $M^* \subseteq M$ such that every model $\mathcal{M} \in M^*$ explains the input observation $\mathcal{O}(\tau)$.

3.2 The space of possible action models

In spite of previous definitions are general, this work focuses on action models that are specified as STRIPS action schemata. STRIPS action schemata provide a compact representation for specifying action models. *A* STRIPS *action schema ξ is defined by four lists*: A list of *parameters* $pars(\xi)$, and three list of predicates (namely $pre(\xi)$, $del(\xi)$ and $add(\xi)$) that shape the kind of fluents that can appear in the *preconditions*, *negative effects* and *positive effects* of the actions induced from that schema.

Let be Ψ the set of *predicates* that shape the propositional state variables F, and a list of *parameters* $pars(\xi)$. The set of elements that can appear in $pre(\xi)$, $del(\xi)$ and $add(\xi)$ of the STRIPS action schema ξ is given by FOL interpretations of Ψ over the parameters $pars(\xi)$. We denote this set of FOL interpretations as $\mathcal{I}_{\Psi, \xi}$.

Despite any element of $\mathcal{I}_{\Psi, \xi}$ can *a priori* appear in the $pre(\xi)$, $del(\xi)$ and $add(\xi)$ of schema ξ, the space of possible STRIPS schemata is constrained by a set $\mathcal{C}$ that includes:

1. *Syntactic constraints.* STRIPS constraints require $del(\xi) \subseteq pre(\xi)$, $del(\xi) \cap add(\xi) = \emptyset$ and $pre(\xi) \cap add(\xi) = \emptyset$. Considering exclusively these syntactic constraints, the size of the space of possible STRIPS schemata is given by $2^{2 \times |\mathcal{I}_{\Psi, \xi}|}$.

2. *Domain-specific constraints.* One can introduce domain-specific knowledge to constrain further the space of possible schemata. For instance, in the *blocksworld* one can argue that $on(v_1, v_1)$ and $on(v_2, v_2)$ will not appear in the $pre(\xi)$, $del(\xi)$ and $add(\xi)$ lists of an action schema ξ because of the semantic of the on predicate (i.e. only one block can be on top of another block). As a rule of thumb, *invariants* constraining the possible states of a given planning domain belong to this class of constraints [Fox and Long, 1998].

Definition 2 (Well-defined STRIPS action schemata)
Given a set of predicates Ψ, a list of action parameters $pars(\xi)$, and set of FOL constraints $\mathcal{C}$, ξ is a **well-defined STRIPS action schema** *iff its three lists $pre(\xi) \subseteq \mathcal{I}_{\Psi, \xi}$, $del(\xi) \subseteq \mathcal{I}_{\Psi, \xi}$ and $add(\xi) \subseteq \mathcal{I}_{\Psi, \xi}$ only contain elements in $\mathcal{I}_{\Psi, \xi}$ and they satisfy all the constraints in $\mathcal{C}$.*

We say a planning model $\mathcal{M}$ is *well-defined* if all its STRIPS action schemata are *well-defined*.

3. *Observation constraints.* A sequence of state observations $\mathcal{O}(\tau)$ constraints further the space of possible action schemata. This *semantic knowledge* included in the observations introduce a third type of constraints and can also be added to the set $\mathcal{C}$.

```
(:action stack
  :parameters (?v1 ?v2)
  :precondition (and (holding ?v1) (clear ?v2))
  :effect (and (not (holding ?v1)) (not (clear ?v2))
               (clear ?v1) (handempty) (on ?v1 ?v2)))


(pre_holding_v1_stack) (pre_clear_v2_stack)
(eff_holding_v1_stack) (eff_clear_v2_stack)
(eff_clear_v1_stack) (eff_handempty_stack) (eff_on_v1_v2_stack)
```

Figure 1: PDDL encoding of the `stack(?v1,?v2)` schema and our propositional representation for this same action.

3.3 Computing *the least commitment* model via conformant planning

Given a classical planning problem $P = \langle F, A[\cdot], I, G \rangle$ and a sequence of partial states $\mathcal{O}(\tau)$ coming from the observation of a $\tau(\pi, I)$ trajectory, we show here that we can build a *conformant planning problem* P_c whose solution is a plan that induces the *least commitment* model.

In more detail, our compilation defines a *conformant planning problem* $P_c = \langle F_c, A_c, \Upsilon, G \rangle$ such that:

- The set of fluents F_c extends F with two new sets of fluents:
  - $\{test_j\}_{1 \leq j \leq m}$, indicating the state observation $s_j \in \mathcal{O}(\tau)$ where the action model is validated
  - The fluents pre_e_ξ and eff_e_ξ such that $e \in \mathcal{I}_{\Psi,\xi}$ for a propositional encoding of the *preconditions*, *negative* and *positive* effects of schema ξ. Note that our encoding exploits the syntactic constraint of STRIPS. According to our encoding, if pre_e_ξ and eff_e_ξ holds it means that $e \in \mathcal{I}_{\Psi,\xi}$ is a negative effect in ξ and if pre_e_ξ does not hold but eff_e_ξ holds, it means that $e \in \mathcal{I}_{\Psi,\xi}$ is a positive effect in ξ. Figure 1 shows the PDDL encoding of the `stack(?v1,?v2)` schema and our propositional representation for this same action.

- The set of actions A_c contains now actions of three different kinds:
  - Actions for *committing* pre_e_ξ fluents to a positive and a negative value.

$$\text{pre}(\text{commit}\top\_\text{pre}\_e\_\xi) = \{mode_{commit}\},$$
$$\text{cond}(\text{commit}\top\_\text{pre}\_e\_\xi) = \{pre\_e\_\xi\} \rhd \{pre\_e\_\xi\},$$
$$\{\neg pre\_e\_\xi\} \rhd \{pre\_e\_\xi\}.$$
$$\text{pre}(\text{commit}\bot\_\text{pre}\_e\_\xi) = \{mode_{commit}\},$$
$$\text{cond}(\text{commit}\bot\_\text{pre}\_e\_\xi) = \{pre\_e\_\xi\} \rhd \{\neg pre\_e\_\xi\},$$
$$\{\neg pre\_e\_\xi\} \rhd \{\neg pre\_e\_\xi\}.$$

  In addition, similar actions are defined for *committing* eff_e_ξ fluents to a positive value and to a negative value.
  - Actions for *validating* the committed models at the s_j observed states, $0 \leq j < m$

$$\text{pre}(\text{validate}_j) = s_j \cup \{test_{j-1}\},$$
$$\text{cond}(\text{validate}_j) = \{\emptyset\} \rhd \{\neg test_{j-1}, test_j,$$
$$\{mode_{commit}\} \rhd \{\neg mode_{commit}, mode_{val}\}.$$

```
(:action stack
 :parameters (?o1 - object ?o2 - object)
 :precondition
   (and (or (not (pre_on_v1_v1_stack)) (on ?o1 ?o1))
        (or (not (pre_on_v1_v2_stack)) (on ?o1 ?o2))
        (or (not (pre_on_v2_v1_stack)) (on ?o2 ?o1))
        (or (not (pre_on_v2_v2_stack)) (on ?o2 ?o2))
        (or (not (pre_ontable_v1_stack)) (ontable ?o1))
        (or (not (pre_ontable_v2_stack)) (ontable ?o2))
        (or (not (pre_clear_v1_stack)) (clear ?o1))
        (or (not (pre_clear_v2_stack)) (clear ?o2))
        (or (not (pre_holding_v1_stack)) (holding ?o1))
        (or (not (pre_holding_v2_stack)) (holding ?o2))
        (or (not (pre_handempty_stack)) (handempty)))
 :effect (and
   (when (and (pre_on_v1_v1_stack)(eff_on_v1_v1_stack)) (not (on ?o1 ?o1)))
   (when (and (pre_on_v1_v2_stack)(eff_on_v1_v2_stack)) (not (on ?o1 ?o2)))
   (when (and (pre_on_v2_v1_stack)(eff_on_v2_v1_stack)) (not (on ?o2 ?o1)))
   (when (and (pre_on_v2_v2_stack)(eff_on_v2_v2_stack)) (not (on ?o2 ?o2)))
   (when (and (pre_ontable_v1_stack)(eff_ontable_v1_stack)) (not (ontable ?o1)))
   (when (and (pre_ontable_v2_stack)(eff_ontable_v2_stack)) (not (ontable ?o2)))
   (when (and (pre_clear_v1_stack)(eff_clear_v1_stack)) (not (clear ?o1)))
   (when (and (pre_clear_v2_stack)(eff_clear_v2_stack)) (not (clear ?o2)))
   (when (and (pre_holding_v1_stack)(eff_holding_v1_stack)) (not (holding ?o1)))
   (when (and (pre_holding_v2_stack)(eff_holding_v2_stack)) (not (holding ?o2)))
   (when (and (pre_handempty_stack)(eff_handempty_stack)) (not (handempty)))
   (when (and (not(pre_on_v1_v1_stack))(eff_on_v1_v1_stack)) (on ?o1 ?o1))
   (when (and (not(pre_on_v1_v2_stack))(eff_on_v1_v2_stack)) (on ?o1 ?o2))
   (when (and (not(pre_on_v2_v1_stack))(eff_on_v2_v1_stack)) (on ?o2 ?o1))
   (when (and (not(pre_on_v2_v2_stack))(eff_on_v2_v2_stack)) (on ?o2 ?o2))
   (when (and (not(pre_ontable_v1_stack))(eff_ontable_v1_stack)) (ontable ?o1))
   (when (and (not(pre_ontable_v2_stack))(eff_ontable_v2_stack)) (ontable ?o2))
   (when (and (not(pre_clear_v1_stack))(eff_clear_v1_stack)) (clear ?o1))
   (when (and (not(pre_clear_v2_stack))(eff_clear_v2_stack)) (clear ?o2))
   (when (and (not(pre_holding_v1_stack))(eff_holding_v1_stack)) (holding ?o1))
   (when (and (not(pre_holding_v2_stack))(eff_holding_v2_stack)) (holding ?o2))
   (when (and (not(pre_handempty_stack))(eff_handempty_stack)) (handempty))))
```

Figure 2: PDDL encoding of the editable version of the `stack(?v1,?v2)` schema.

  - Actions whose semantics is given by the value of the pre_e_ξ, eff_e_ξ fluents at the current state. Given an operator schema $\xi \in \mathcal{M}$ its *editable* version is formalized as:

$$\text{pre}(\text{editable}_\xi) = \{pre\_e\_\xi \implies e\}_{\forall e \in \mathcal{I}_{\Psi,\xi}}$$
$$\text{cond}(\text{editable}_\xi) = \{pre\_e\_\xi, eff\_e\_\xi\} \rhd \{\neg e\}_{\forall e \in \mathcal{I}_{\Psi,\xi}},$$
$$\{\neg pre\_e\_\xi, eff\_e\_\xi\} \rhd \{e\}_{\forall e \in \mathcal{I}_{\Psi,\xi}}.$$

  Figure 2 shows the PDDL encoding of the editable version of the `stack(?v1,?v2)` schema. Note that this editable schema, when the fluents `(pre_holding_v1_stack) (pre_clear_v2_stack) (eff_holding_v1_stack) (eff_clear_v2_stack) (eff_clear_v1_stack) (eff_handempty_stack) (eff_on_v1_v2_stack)` hold, it behaves exactly as defined in Figure 1.

- The clauses in Υ extend the set of fluents that hold in the initial state $I = s_0$ and $mode_{prog}$ set to true. In addition, Υ includes the two clauses pre_e_ξ xor $\neg pre\_e\_\xi$ and eff_e_ξ xor $\neg eff\_e\_\xi$ for every ξ and $e \in \mathcal{I}_{\Psi,\xi}$. This means that the actual value of these fluents is actually unknown, so any model from the STRIPS space of models can initially be part of the *the least commitment* action model.

- The new goals are $G_c = \{test_m\}$.

3.4 Compilation properties

4 Evaluation

5 Conclusions

References

[Fox and Long, 1998] Maria Fox and Derek Long. The automatic inference of state invariants in tim. *Journal of Artificial Intelligence Research*, 9:367–421, 1998.

[Goldman and Boddy, 1996] Robert P Goldman and Mark S Boddy. Expressive planning and explicit knowledge. In *AIPS*, volume 96, pages 110–117, 1996.

[Ramírez and Geffner, 2009] Miquel Ramírez and Hector Geffner. Plan recognition as planning. In *International Joint conference on Artifical Intelligence, (IJCAI-09)*, pages 1778–1783. AAAI Press, 2009.

[Smith and Weld, 1998] David E Smith and Daniel S Weld. Conformant graphplan. In *AAAI/IAAI*, pages 889–896, 1998.