Computing the *least-commitment* action model from state observations

Diego Aineto[1] , **Sergio Jiménez**[1] , **Eva Onaindia**[1] and , **Blai Bonet**[2]

[1]Departamento de Sistemas Informáticos y Computación. Universitat Politècnica de València. Valencia, Spain

[2]Departamento de Computación. Universidad Simón Bolívar. Caracas, Venezuela

{dieaigar,serjice,onaindia}@dsic.upv.es, bonet@usb.ve

Abstract

1 Introduction

Given an input sequence of partially observed states, this paper formalizes the task of computing the *least-commitment* action model that is able to *explain* the given observation. This task is of interest because it allows the incremental learning of action models from arbitrary large sets of partial observations.

In addition, the paper introduces a new method to compute the *least-commitment* action model for an input sequence of partially observed states. The method assumes that action models are specified as STRIPS action schemata and it is built on top of off-the-shelf algorithms for *classical planning*.

2 Background

This section formalizes the *planning models* we use in the paper as well as the kind of state *observations* that are given as input for computing the *least-commitment* action model.

2.1 Classical planning with conditional effects

Let F be the set of *fluents* or *state variables* (propositional variables) describing a state. A *literal* l is a valuation of a fluent $f \in F$; i.e. either $l = f$ or $l = \neg f$. A set of literals L represents a partial assignment of values to fluents (without loss of generality, we will assume that L does not contain conflicting values). Given L, let $\neg L = \{\neg l : l \in L\}$ be its complement. We use $\mathcal{L}(F)$ to denote the set of all literal sets on F; i.e. all partial assignments of values to fluents. A *state* s is a full assignment of values to fluents; $|s| = |F|$.

A *classical planning frame* is a tuple $\Phi = \langle F, A \rangle$, where F is a set of fluents and A is a set of *actions*. Each classical planning action $a \in A$ has a precondition $\mathsf{pre}(a) \in \mathcal{L}(F)$ and a set of effects $\mathsf{eff}(a) \in \mathcal{L}(F)$. The semantics of actions $a \in A$ is specified with two functions: $\rho(s, a)$ denotes whether action a is *applicable* in a state s and $\theta(s, a)$ denotes the *successor state* that results of applying action a in a state s. Then, $\rho(s, a)$ holds iff $\mathsf{pre}(a) \subseteq s$, i.e. if its precondition holds in s. The result of executing an applicable action $a \in A$ in a state s is a new state $\theta(s, a) = (s \setminus \neg\mathsf{eff}(a)) \cup \mathsf{eff}(a)$. Subtracting the complement of $\mathsf{eff}(a)$ from s ensures that $\theta(s, a)$

remains a well-defined state. The subset of action effects that assign a positive value to a state fluent is called *positive effects* and denoted by $\mathsf{eff}^+(a) \in \mathsf{eff}(a)$ while $\mathsf{eff}^-(a) \in \mathsf{eff}(a)$ denotes the *negative effects* of an action $a \in A$.

A *classical planning problem* is a tuple $P = \langle F, A, I, G \rangle$, where I is the initial state and $G \in \mathcal{L}(F)$ is the set of goal conditions over the state variables. A *plan* π is an action sequence $\pi = \langle a_1, \ldots, a_n \rangle$, with $|\pi| = n$ denoting its *plan length*. The execution of π on the initial state I of P induces a *trajectory* $\tau(\pi, s_0) = \langle s_0, a_1, s_1, \ldots, a_n, s_n \rangle$ such that $s_0 = I$ and, for each $1 \le i \le n$, it holds $\rho(s_{i-1}, a_i)$ and $s_i = \theta(s_{i-1}, a_i)$. A plan π solves P iff the induced *trajectory* $\tau(\pi, s_0)$ reaches a final state $G \subseteq s_n$, where all goal conditions are met. A solution plan is optimal iff its length is minimal.

An action $a_c \in A$ with conditional effects is defined as a set of preconditions $\mathsf{pre}(a_c) \in \mathcal{L}(F)$ and a set of *conditional effects* $\mathsf{cond}(a_c)$. Each conditional effect $C \rhd E \in \mathsf{cond}(a_c)$ is composed of two sets of literals: $C \in \mathcal{L}(F)$, the *condition*, and $E \in \mathcal{L}(F)$, the *effect*. An action a_c is applicable in a state s if $\rho(s, a_c)$ is true, and the result of applying action a_c in state s is $\theta(s, a_c) = \{s \setminus \neg\mathsf{eff}_c(s, a) \cup \mathsf{eff}_c(s, a)\}$ where $\mathsf{eff}_c(s, a)$ are the *triggered effects* resulting from the action application (conditional effects whose conditions hold in s):

$$\mathsf{eff}_c(s, a) = \bigcup_{C \rhd E \in \mathsf{cond}(a_c), C \subseteq s} E,$$

2.2 The observation model

Given a classical planning problem $P = \langle F, A, I, G \rangle$, a plan π and a trajectory $\tau(\pi, s_0)$, we define the *observation of the trajectory* as a sequence of partial states that results from observing the execution of π on I. Formally, $\mathcal{O}(\tau) = \langle s_0^o, s_1^o \ldots, s_m^o \rangle$ where $s_0^o = I$.

A *partial state* s_i^o, $0 < i < m$, is one in which $|s_i^o| < |F|$; i.e., a state in which at least a fluent of F is not observable. Note that this definition also comprises the case $|s_i^o| = 0$, when the state is fully unobservable. Whatever the sequence of observed states of $\mathcal{O}(\tau)$ is, it must be *consistent* with the sequence of states of $\tau(\pi, s_0)$, meaning that $\forall i, s_i^o \subseteq s_i$. In practice, the number of observed states m, ranges from 1 (the initial state, at least), to $|\pi| + 1$, and the observed intermediate states will comprise a number of fluents between $[1, |F|]$.

We are assuming then that there is a *bijective mono-tone mapping* between trajectories and observations [Ramírez and Geffner, 2009], thus also granting the inverse consistency relationship (the trajectory is a superset of the observation). Therefore, transiting between two consecutive observed states in $\mathcal{O}(\tau)$ may require the execution of more than a single action $(\theta(s_i^o, \langle a_1, \ldots, a_k \rangle) = s_{i+1}^o$, where $k \geq 1$ is unknown but finite. In other words, having $\mathcal{O}(\tau)$ does not imply knowing the actual length of π.

Definition 1 (Plan explanation) *Given a* classical planning problem P *and an observation* $\mathcal{O}(\tau)$*, a plan* π *explains* $\mathcal{O}(\tau)$ *(denoted* $\pi \mapsto \mathcal{O}(\tau)$*) iff* π *is a solution for* P *that is* consistent *with the state trajectory constraints imposed by the sequence of partial states* $\mathcal{O}(\tau)$.

If π is also optimal, we say that π is the *best explanation* for the input observation $\mathcal{O}(\tau)$.

3 Computing the *least-commitment* action model from observations

First, this section formalizes the notion of the *least-commitment* action model that is able to *explain* an observation. Next, the section describes our approach to compute such model.

3.1 The *least-commitment* action model

The task of computing the *least-commitment* action model from a sequence of state observations is defined as $\langle \Phi, \mathcal{O}(\tau) \rangle$:

- $\Phi = \langle F, A[\cdot] \rangle$ is a *classical planning frame* where the semantics of each action $a \in A[\cdot]$ is unknown; i.e. the corresponding $\langle \rho, \theta \rangle$ functions are undefined. We say that an *action model* $\mathcal{M}$ is a definition of the $\langle \rho, \theta \rangle$ functions of every action in $A[\cdot]$.

- $\mathcal{O}(\tau)$ is a sequence of partial states that results from the partial observation of a trajectory $\tau(\pi, s_0)$ within the *classical planning frame* Φ.

Given a *classical planning frame* $\Phi = \langle F, A[\cdot] \rangle$ and an observation $\mathcal{O}(\tau) = \langle s_0^o, s_1^o \ldots, s_m^o \rangle$ let $P_\mathcal{O}$ be the classical planning problem $P_\mathcal{O} = \langle F, A[\cdot], s_0^o, s_m^o \rangle$.

Definition 2 (Model explanation) *A model* $\mathcal{M}$ *explains an observation* $\mathcal{O}(\tau)$ *iff, when the* $\langle \rho, \theta \rangle$ *functions of the actions in* $P_\mathcal{O}$ *are given by* $\mathcal{M}$*, there exists a solution plan for* $P_\mathcal{O}$ *that* explains $\mathcal{O}(\tau)$.

Definition 3 (The *least-commitment* action model) *Given a* $\langle \Phi, \mathcal{O}(\tau) \rangle$ *task (and let* M *be the set of action models that represents the full space of possible action models for the actions in* $A[\cdot] \in \Phi$*), the* least-commitment *action model is the largest subset of models* $M^* \subseteq M$ *such that every model* $\mathcal{M} \in M^*$ *explains* the input observation.

This work focuses on the particular task of computing the *least-commitment* action model when action models are specified as STRIPS action schemata.

```
(:action stack
    :parameters (?v1 ?v2)
    :precondition (and (holding ?v1) (clear ?v2))
    :effect (and (not (holding ?v1)) (not (clear ?v2))
                 (clear ?v1) (handempty) (on ?v1 ?v2)))


(pre_holding_v1_stack) (pre_clear_v2_stack)
(eff_holding_v1_stack) (eff_clear_v2_stack)
(eff_clear_v1_stack) (eff_handempty_stack) (eff_on_v1_v2_stack)
```

Figure 1: PDDL encoding of the stack(?v1,?v2) schema and our propositional representation for this same schema.

3.2 The space of STRIPS action models

A STRIPS *action schema* ξ is defined by four lists: A list of *parameters* $pars(\xi)$, and three list of predicates (namely $pre(\xi)$, $del(\xi)$ and $add(\xi)$) that shape the kind of fluents that can appear in the *preconditions*, *negative effects* and *positive effects* of the actions induced from that schema. Let be Ψ the set of *predicates* that shape the propositional state variables F, and a list of *parameters* $pars(\xi)$. The set of elements that can appear in $pre(\xi)$, $del(\xi)$ and $add(\xi)$ of the STRIPS action schema ξ is given by FOL interpretations of Ψ over the parameters $pars(\xi)$. We denote this set of FOL interpretations as $\mathcal{I}_{\Psi,\xi}$. For instance, in the *blocksworld* the $\mathcal{I}_{\Psi,\xi}$ set contain eleven elements for the stack(v_1, v_2) schemata, $\mathcal{I}_{\Psi,stack}=\{$handempty, holding($v_1$), holding($v_2$), clear($v_1$), clear($v_2$), ontable($v_1$), ontable($v_2$), on($v_1, v_1$), on($v_1, v_2$), on($v_2, v_1$), on($v_2, v_2$)$\}$.

Despite any element of $\mathcal{I}_{\Psi,\xi}$ can *a priori* appear in the $pre(\xi)$, $del(\xi)$ and $add(\xi)$ of schema ξ, the space of possible STRIPS schemata is bounded by constraints of three kinds:

1. *Syntactic constraints*. STRIPS constraints require $del(\xi) \subseteq pre(\xi)$, $del(\xi) \cap add(\xi) = \emptyset$ and $pre(\xi) \cap add(\xi) = \emptyset$. Considering exclusively these syntactic constraints, the size of the space of possible STRIPS schemata is given by $2^{2 \times |\mathcal{I}_{\Psi,\xi}|}$. *Typing constraints* are also of this kind [McDermott *et al.*, 1998].

2. *Domain-specific constraints*. One can introduce domain-specific knowledge to constrain further the space of possible schemata. For instance, in the *blocksworld* one can argue that on(v_1, v_1) and on(v_2, v_2) will not appear in the $pre(\xi)$, $del(\xi)$ and $add(\xi)$ lists of an action schema ξ because, in this specific domain, a block cannot be on top of itself. *State invariants* are also constraints of this kind [Fox and Long, 1998].

3. *Observation constraints*. An observations $\mathcal{O}(\tau)$ depicts *semantic knowledge* that constraints further the space of possible action schemata.

In this work we introduce a propositional encoding of the *preconditions*, *negative*, and *positive* effects of a STRIPS action schema ξ using only fluents of two kinds pre_e_ξ and eff_e_ξ (where $e \in \mathcal{I}_{\Psi,\xi}$). This encoding exploits the syntactic constraints of STRIPS so is more compact that the one previuosly proposed by Aineto *et al.* 2018. In more detail, if pre_e_ξ and eff_e_ξ holds it means that $e \in \mathcal{I}_{\Psi,\xi}$ is a negative effect in ξ while if pre_e_ξ does not hold but eff_e_ξ

```
(Kpre_holding_v1_stack) (Kpre_clear_v2_stack)
(K¬pre_holding_v2_stack) (K¬pre_clear_v1_stack)
(K¬pre_handempty_stack)
(K¬pre_ontable_v1_stack) (K¬pre_ontable_v2_stack)
(K¬pre_on_v1_v2_stack) (K¬pre_on_v2_v1_stack)
(K¬pre_on_v1_v1_stack) (K¬pre_on_v2_v2_stack)

(K¬eff_holding_v2_stack) (K¬eff_clear_v2_stack)
(K¬eff_clear_v2_stack) (K¬eff_on_v2_v2_stack)
(K¬eff_on_v1_v1_stack) (K¬eff_on_v2_v1_stack)
```

Figure 2: *Partially specified* STRIPS model that represents of a set of action models for the `stack(?v1,?v2)` schema.

holds, it means that $e \in \mathcal{I}_{\Psi,\xi}$ is a positive effect in ξ. Figure 1 shows the PDDL encoding of the `stack(?v1,?v2)` schema and our propositional representation for this same schema.

3.3 *Partially specified* STRIPS action models

A set of action models can be defined *explicitely*, by enumerating all the models that belong to the set or *implictely*, by enumerating the constraints that must satisfy any model that belongs to the set. Inspired by the notion of *incomplete (annotated) model* [Sreedharan *et al.*, 2018], we introduce here *partially specified* STRIPS action models, a formalism for the *implicit* representation of a set of STRIPS schema.

We show now that extending to the *knwoledge level* our propositional encoding of STRIPS action schemes, we can compactly represent a set of STRIPS schema. The extension defines, for each proposition pre_e_ξ, with two propositions $K\mathsf{pre}\_e\_\xi$ and $K\neg\mathsf{pre}\_e\_\xi$, meaning that is *known* that $e \in \mathcal{I}_{\Psi,\xi}$ is a precondition of ξ and that is *known* that $e \in \mathcal{I}_{\Psi,\xi}$ is not a precondition of ξ. Likewise for each eff_e_ξ proposition we define these two, $K\mathsf{eff}\_e\_\xi$ and $K\neg\mathsf{eff}\_e\_\xi$, meaning that is *known* that $e \in \mathcal{I}_{\Psi,\xi}$ is an effect of ξ and that is *known* that $e \in \mathcal{I}_{\Psi,\xi}$ is not an effect of ξ. With respect to this formalization, the *full space* of possible STRIPS schemas for ξ is compactly represented as $\bigcap_{e \in \mathcal{I}_{\Psi,\xi}} \neg K\mathsf{pre}\_e\_\xi \wedge \neg K\neg\mathsf{pre}\_e\_\xi \wedge \neg K\mathsf{eff}\_e\_\xi \wedge \neg K\neg\mathsf{eff}\_e\_\xi$.

Figure 2 shows a *partially specified* STRIPS model that represents of a set of action models for the `stack(?v1,?v2)` schema. This *partially specified* STRIPS model represents a set of 2^5 models since (1), the actual five effects of the action are *unkwnon*, (2) preconditions are known to appear as in the actual definition of the `stack(?v1,?v2)` schema for the blocksworld and (3), it is *known* that the preconditions and effects that are not part of the actual definition of the `stack(?v1,?v2)` schema are not part of this *partially specified* STRIPS model.

Definition 4 (Partially specified model explanation)

3.4 Computing the *least-commitment* model via classical planning

Inspired by the *classical planning compilation* K_0 for conformant planning [Palacios and Geffner, 2009], this section shows that we can build a *classical planning problem* $P = \langle F', A', I', G' \rangle$ whose solution induces the *least-commitment* action model for an observation $\mathcal{O}(\tau)$:

- The set of fluents F' extends F with two new sets of fluents:

- $\{test_j\}_{1 \leq j \leq m}$, indicating the state observation $s_j \in \mathcal{O}(\tau)$ where the action model is validated
- The *knowledge level* fluents $K\mathsf{pre}\_e\_\xi$, $K\neg\mathsf{pre}\_e\_\xi$, $K\mathsf{eff}\_e\_\xi$ and $K\neg\mathsf{eff}\_e\_\xi$ encoding the space of possible *partially specified* action models.

- The set of actions A' contains now actions of three different kinds:
  - Actions for *committing* pre_e_ξ to a positive/negative value. Similar actions are also defined for *committing* eff_e_ξ to a positive/negative value but the value of eff_e_ξ can only be committed once the value of the corresponding pre_e_ξ is committed (i.e. once either $K\mathsf{pre}\_e\_\xi$ or $K\neg\mathsf{pre}\_e\_\xi$ holds in the current state).

$$\mathsf{pre}(\mathsf{commit}\top\_\mathsf{pre}\_e\_\xi) = \{mode_{commit}, \\ \neg K\mathsf{pre}\_e\_\xi, \neg K\neg\mathsf{pre}\_e\_\xi\},$$
$$\mathsf{cond}(\mathsf{commit}\top\_\mathsf{pre}\_e\_\xi) = \{\emptyset\} \triangleright \{K\mathsf{pre}\_e\_\xi\}.$$

$$\mathsf{pre}(\mathsf{commit}\bot\_\mathsf{pre}\_e\_\xi) = \{mode_{commit}, \\ \neg K\mathsf{pre}\_e\_\xi, \neg K\neg\mathsf{pre}\_e\_\xi\},$$
$$\mathsf{cond}(\mathsf{commit}\bot\_\mathsf{pre}\_e\_\xi) = \{\emptyset\} \triangleright \{K\neg\mathsf{pre}\_e\_\xi\}.$$

- Actions for *validating* that committed models explain the s_j observed states, $0 \leq j < m$.
$$\mathsf{pre}(\mathsf{validate}_j) = s_j \cup \{test_{j-1}\},$$
$$\mathsf{cond}(\mathsf{validate}_j) = \{\emptyset\} \triangleright \{\neg test_{j-1}, test_j\}, \\ \{mode_{commit}\} \triangleright \{\neg mode_{commit}, mode_{val}\}.$$

- *Editable* actions whose semantics is given by the value of the *knowledge level* fluents ($K\mathsf{pre}\_e\_\xi$, $K\neg\mathsf{pre}\_e\_\xi$, $K\mathsf{eff}\_e\_\xi$ and $K\neg\mathsf{eff}\_e\_\xi$) at the current state. Figure 3 shows the PDDL encoding of an *editable* `stack(?v1,?v2)` schema. This editable schema behaves exactly as the original PDDL schema defined in Figure 1 when the set of fluents `(Kpre_holding_v1_stack) (Kpre_clear_v2_stack) (Keff_holding_v1_stack) (Keff_clear_v2_stack) (Keff_clear_v1_stack) (Keff_handempty_stack) (Keff_on_v1_v2_stack)` hold at the current state as well as the remaining $K\neg\mathsf{pre}\_e\_\xi$ and $K\neg\mathsf{eff}\_e\_\xi$ for all the preconditions and effects that are not part of the actual `stack(?v1,?v2)` schema. Formally, given an operator schema $\xi \in \mathcal{M}$ its *editable* version is:

$$\mathsf{pre}(\mathsf{editable}_\xi) = \{\neg K\neg\mathsf{pre}\_e\_\xi \implies e\}_{\forall e \in \mathcal{I}_{\Psi,\xi}}$$
$$\mathsf{cond}(\mathsf{editable}_\xi) = \{K\mathsf{pre}\_e\_\xi, \neg K\neg\mathsf{eff}\_e\_\xi\} \triangleright \{\neg e\}_{\forall e \in \mathcal{I}_{\Psi,\xi}}, \\ \{K\neg\mathsf{pre}\_e\_\xi, \neg K\neg\mathsf{eff}\_e\_\xi\} \triangleright \{e\}_{\forall e \in \mathcal{I}_{\Psi,\xi}}.$$

- The new initial state $I' = I \cup \{mode_{commit}\}$ while the new goals are $G' = s_m \cup \{test_m\}$.

3.5 Compilation properties

4 Evaluation

5 Conclusions

Related work [Stern and Juba, 2017].

```
(:action stack
 :parameters (?o1 - object ?o2 - object)
 :precondition
   (and (or (Knotpre_on_v1_v1_stack) (on ?o1 ?o1))
        (or (Knotpre_on_v1_v2_stack) (on ?o1 ?o2))
        (or (Knotpre_on_v2_v1_stack) (on ?o2 ?o1))
        (or (Knotpre_on_v2_v2_stack) (on ?o2 ?o2))
        (or (Knotpre_ontable_v1_stack) (ontable ?o1))
        (or (Knotpre_ontable_v2_stack) (ontable ?o2))
        (or (Knotpre_clear_v1_stack) (clear ?o1))
        (or (Knotpre_clear_v2_stack) (clear ?o2))
        (or (Knotpre_holding_v1_stack) (holding ?o1))
        (or (Knotpre_holding_v2_stack) (holding ?o2))
        (or (Knotpre_handempty_stack) (handempty)))
 :effect (and
   (when (and (Kpre_on_v1_v1_stack)(not(Knot_eff_on_v1_v1_stack))) (not (on ?o1 ?o1)))
   (when (and (Kpre_on_v1_v2_stack)(not(Knot_eff_on_v1_v2_stack))) (not (on ?o1 ?o2)))
   (when (and (Kpre_on_v2_v1_stack)(not(Knot_eff_on_v2_v1_stack))) (not (on ?o2 ?o1)))
   (when (and (Kpre_on_v2_v2_stack)(not(Knot_eff_on_v2_v2_stack))) (not (on ?o2 ?o2)))
   (when (and (Kpre_ontable_v1_stack)(not(Knot_eff_ontable_v1_stack))) (not (ontable ?o1)))
   (when (and (Kpre_ontable_v2_stack)(not(Knot_eff_ontable_v2_stack))) (not (ontable ?o2)))
   (when (and (Kpre_clear_v1_stack)(not(Knot_eff_clear_v1_stack))) (not (clear ?o1)))
   (when (and (Kpre_clear_v2_stack)(not(Knot_eff_clear_v2_stack))) (not (clear ?o2)))
   (when (and (Kpre_holding_v1_stack)(not(Knot_eff_holding_v1_stack))) (not (holding ?o1)))
   (when (and (Kpre_holding_v2_stack)(not(Knot_eff_holding_v2_stack))) (not (holding ?o2)))
   (when (and (Kpre_handempty_stack)(not(Knot_eff_handempty_stack))) (not (handempty)))
   (when (and (Knot_pre_on_v1_v1_stack)(not(Knot_eff_on_v1_v1_stack))) (on ?o1 ?o1))
   (when (and (Knot_pre_on_v1_v2_stack)(not(Knot_eff_on_v1_v2_stack))) (on ?o1 ?o2))
   (when (and (Knot_pre_on_v2_v1_stack)(not(Knot_eff_on_v2_v1_stack))) (on ?o2 ?o1))
   (when (and (Knot_pre_on_v2_v2_stack)(not(Knot_eff_on_v2_v2_stack))) (on ?o2 ?o2))
   (when (and (Knot_pre_ontable_v1_stack)(not(Knot_eff_ontable_v1_stack))) (ontable ?o1))
   (when (and (Knot_pre_ontable_v2_stack)(not(Knot_eff_ontable_v2_stack))) (ontable ?o2))
   (when (and (Knot_pre_clear_v1_stack)(not(Knot_eff_clear_v1_stack))) (clear ?o1))
   (when (and (Knot_pre_clear_v2_stack)(not(Knot_eff_clear_v2_stack))) (clear ?o2))
   (when (and (Knot_pre_holding_v1_stack)(not(Knot_eff_holding_v1_stack))) (holding ?o1))
   (when (and (Knot_pre_holding_v2_stack)(not(Knot_eff_holding_v2_stack))) (holding ?o2))
   (when (and (Knot_pre_handempty_stack)(not(Knot_eff_handempty_stack))) (handempty))))
```

Figure 3: PDDL encoding of the editable version of the `stack(?v1,?v2)` schema.

References

[Aineto *et al.*, 2018] Diego Aineto, Sergio Jiménez, and Eva Onaindia. Learning STRIPS action models with classical planning. In *International Conference on Automated Planning and Scheduling, (ICAPS-18)*, pages 399–407. AAAI Press, 2018.

[Fox and Long, 1998] Maria Fox and Derek Long. The automatic inference of state invariants in tim. *Journal of Artificial Intelligence Research*, 9:367–421, 1998.

[McDermott *et al.*, 1998] Drew McDermott, Malik Ghallab, Adele Howe, Craig Knoblock, Ashwin Ram, Manuela Veloso, Daniel Weld, and David Wilkins. PDDL – The Planning Domain Definition Language, 1998.

[Palacios and Geffner, 2009] Héctor Palacios and Héctor Geffner. Compiling uncertainty away in conformant planning problems with bounded width. *Journal of Artificial Intelligence Research*, 35:623–675, 2009.

[Ramírez and Geffner, 2009] Miquel Ramírez and Hector Geffner. Plan recognition as planning. In *International Joint conference on Artifical Intelligence, (IJCAI-09)*, pages 1778–1783. AAAI Press, 2009.

[Sreedharan *et al.*, 2018] Sarath Sreedharan, Tathagata Chakraborti, and Subbarao Kambhampati. Handling model uncertainty and multiplicity in explanations via model reconciliation. In *International Conference on Automated Planning and Scheduling, (ICAPS-18)*, pages 518–526, 2018.

[Stern and Juba, 2017] Roni Stern and Brendan Juba. Efficient, safe, and probably approximately complete learning of action models. In *International Joint Conference on Artificial Intelligence, (IJCAI-17)*, pages 4405–4411, 2017.