

# Computing the *least-commitment* action model from state observations

Diego Aineto<sup>1</sup>, Sergio Jiménez<sup>1</sup>, Eva Onaindia<sup>1</sup> and , Blai Bonet<sup>2</sup>

<sup>1</sup>Departamento de Sistemas Informáticos y Computación. Universitat Politècnica de València. Valencia, Spain

<sup>2</sup>Departamento de Computación. Universidad Simón Bolívar. Caracas, Venezuela

{dieaigar,serjice,onaindia}@dsic.upv.es, bonet@usb.ve

## Abstract

## 1 Introduction

Given an input sequence of partially observed states, this paper formalizes the task of computing the *least-commitment* action model that is able to *explain* that input observation.

This task is of interest because, if we can keep the set of all the action models that are consistent with the observation, i.e. the *version space* [Mitchell, 1982], then each new observation has either no effect on the *version space* or eliminates some models from it. This allows to implement an incremental pruning of the *version space* and hence, to learn planning action models from arbitrary large sets of partial observations (one has never to go back and re-process old observations).

In addition, the paper introduces a new method to compute the *least-commitment* action model for an input sequence of partially observed states. The method assumes that action models are specified as STRIPS action schemata and it is built on top of off-the-shelf algorithms for *classical planning*.

## 2 Background

This section formalizes the *planning models* we use in the paper as well as the kind of state *observations* that are given as input for computing the *least-commitment* action model.

### 2.1 Classical planning with conditional effects

Let  $F$  be the set of *fluents* or *state variables* (propositional variables) describing a state. A *literal*  $l$  is a valuation of a fluent  $f \in F$ ; i.e. either  $l = f$  or  $l = \neg f$ . A set of literals  $L$  represents a partial assignment of values to fluents (without loss of generality, we will assume that  $L$  does not contain conflicting values). Given  $L$ , let  $\neg L = \{\neg l : l \in L\}$  be its complement. We use  $\mathcal{L}(F)$  to denote the set of all literal sets on  $F$ ; i.e. all partial assignments of values to fluents. A *state*  $s$  is a full assignment of values to fluents;  $|s| = |F|$ .

A *classical planning frame* is a tuple  $\Phi = \langle F, A \rangle$ , where  $F$  is a set of fluents and  $A$  is a set of *actions*. Each classical planning action  $a \in A$  has a precondition  $\text{pre}(a) \in \mathcal{L}(F)$ , a set of effects  $\text{eff}(a) \in \mathcal{L}(F)$ , and a positive action cost  $c(a)$ . The semantics of actions  $a \in A$  is specified with

two functions:  $\rho(s, a)$  denotes whether action  $a$  is *applicable* in a state  $s$  and  $\theta(s, a)$  denotes the *successor state* that results of applying action  $a$  in a state  $s$ . Then,  $\rho(s, a)$  holds iff  $\text{pre}(a) \subseteq s$ , i.e. if its precondition holds in  $s$ . The result of executing an applicable action  $a \in A$  in a state  $s$  is a new state  $\theta(s, a) = (s \setminus \neg \text{eff}(a)) \cup \text{eff}(a)$ . Subtracting the complement of  $\text{eff}(a)$  from  $s$  ensures that  $\theta(s, a)$  remains a well-defined state. The subset of action effects that assign a positive value to a state fluent is called *positive effects* and denoted by  $\text{eff}^+(a) \in \text{eff}(a)$  while  $\text{eff}^-(a) \in \text{eff}(a)$  denotes the *negative effects* of an action  $a \in A$ .

A *classical planning problem* is a tuple  $P = \langle F, A, I, G \rangle$ , where  $I$  is the initial state and  $G \in \mathcal{L}(F)$  is the set of goal conditions over the state variables. A *plan*  $\pi$  is an action sequence  $\pi = \langle a_1, \dots, a_n \rangle$ , with  $|\pi| = n$  denoting its *plan length*. The execution of  $\pi$  on the initial state  $I$  of  $P$  induces a *trajectory*  $\tau(\pi, s_0) = \langle s_0, a_1, s_1, \dots, a_n, s_n \rangle$  such that  $s_0 = I$  and, for each  $1 \leq i \leq n$ , it holds  $\rho(s_{i-1}, a_i)$  and  $s_i = \theta(s_{i-1}, a_i)$ . A plan  $\pi$  solves  $P$  iff the induced *trajectory*  $\tau(\pi, s_0)$  reaches a final state  $G \subseteq s_n$ , where all goal conditions are met. A solution plan is *optimal* iff it minimizes the sum of action costs.

An action  $a_c \in A$  with conditional effects is defined as a set of preconditions  $\text{pre}(a_c) \in \mathcal{L}(F)$  and a set of *conditional effects*  $\text{cond}(a_c)$ . Each conditional effect  $C \triangleright E \in \text{cond}(a_c)$  is composed of two sets of literals:  $C \in \mathcal{L}(F)$ , the *condition*, and  $E \in \mathcal{L}(F)$ , the *effect*. An action  $a_c$  is applicable in a state  $s$  if  $\rho(s, a_c)$  is true, and the result of applying action  $a_c$  in state  $s$  is  $\theta(s, a_c) = \{s \setminus \neg \text{eff}_c(s, a) \cup \text{eff}_c(s, a)\}$  where  $\text{eff}_c(s, a)$  are the *triggered effects* resulting from the action application (conditional effects whose conditions hold in  $s$ ):

$$\text{eff}_c(s, a) = \bigcup_{C \triangleright E \in \text{cond}(a_c), C \subseteq s} E,$$

### 2.2 The observation model

Given a classical planning problem  $P = \langle F, A, I, G \rangle$ , a plan  $\pi$  and a trajectory  $\tau(\pi, s_0)$ , we define the *observation of the trajectory* as a sequence of partial states that results from observing the execution of  $\pi$  on  $I$ . Formally,  $\mathcal{O}(\tau) = \langle s_0^o, s_1^o, \dots, s_m^o \rangle$  where  $s_0^o = I$ .

A *partial state*  $s_i^o$ ,  $0 < i < m$ , is one in which  $|s_i^o| < |F|$ ; i.e., a state in which at least a fluent of  $F$  is not observable. Note that this definition also comprises the case  $|s_i^o| = 0$ ,

when the state is fully unobservable. Whatever the sequence of observed states of  $\mathcal{O}(\tau)$  is, it must be *consistent* with the sequence of states of  $\tau(\pi, s_0)$ , meaning that  $\forall i, s_i^o \subseteq s_i$ . In practice, the number of observed states  $m$ , ranges from 1 (the initial state, at least), to  $|\pi| + 1$ , and the observed intermediate states will comprise a number of fluents between  $[1, |F|]$ .

We are assuming then that there is a *bijective monotone mapping* between trajectories and observations [Ramírez and Geffner, 2009], thus also granting the inverse consistency relationship (the trajectory is a superset of the observation). Therefore, transiting between two consecutive observed states in  $\mathcal{O}(\tau)$  may require the execution of more than a single action ( $\theta(s_i^o, \langle a_1, \dots, a_k \rangle) = s_{i+1}^o$ , where  $k \geq 1$  is unknown but finite. In other words, having  $\mathcal{O}(\tau)$  does not imply knowing the actual length of  $\pi$ .

**Definition 1 (Plan explanation)** *Given a classical planning problem  $P$  and an observation  $\mathcal{O}(\tau)$ , a plan  $\pi$  explains  $\mathcal{O}(\tau)$  (denoted  $\pi \mapsto \mathcal{O}(\tau)$ ) iff  $\pi$  is a solution for  $P$  that is consistent with the state trajectory constraints imposed by the sequence of partial states  $\mathcal{O}(\tau)$ .*

If  $\pi$  is also optimal, we say that  $\pi$  is the *best explanation* for the input observation  $\mathcal{O}(\tau)$ .

### 2.3 Conformant planning

*Conformant planning* is planning with incomplete information about the initial state, no sensing, and validating that goals are achieved with certainty (despite the uncertainty of the initial state) [Goldman and Boddy, 1996; Smith and Weld, 1998; Bonet and Geffner, 2000].

Syntactically, conformant planning problems are expressed in compact form through a set of state variables. A *conformant planning problem* can be defined as a tuple  $P_c = \langle F, A, \Upsilon, G \rangle$  where  $F$ ,  $A$  and  $G$  are the set of *fluents*, *actions* and *goals* (as previously defined for *classical planning*). Now  $\Upsilon$  is a set of clauses over literals  $l = f$  or  $l = \neg f$  (for  $f \in F$ ) that define the set of possible initial states.

A solution to a conformant planning problem is an action sequence that maps each possible initial state into a goal state. More precisely, an action sequence  $\pi = \langle a_1, \dots, a_n \rangle$  is a *conformant plan* for  $P_c$  iff, for each possible trajectory  $\tau(\pi, s_0) = \langle s_0, a_1, s_1, \dots, a_n, s_n \rangle$  s.t.  $s_0$  is a valuation of the fluents in  $F$  that satisfies  $\Upsilon$ , then  $\tau(\pi, s_0)$  reaches a final state  $G \subseteq s_n$ .

## 3 The least-commitment action model

The task of computing the *least-commitment* action model from a sequence of state observations is defined as  $\langle \Phi, \mathcal{O}(\tau) \rangle$ :

- $\Phi = \langle F, A[\cdot] \rangle$  is a *classical planning frame* where the semantics of each action  $a \in A[\cdot]$  is unknown; i.e. the corresponding  $\langle \rho, \theta \rangle$  functions are undefined. We say that an *action model*  $\mathcal{M}$  is a definition of the  $\langle \rho, \theta \rangle$  functions of every action in  $A[\cdot]$ .
- $\mathcal{O}(\tau)$  is a sequence of partial states that results from the partial observation of a trajectory  $\tau(\pi, s_0)$  within the *classical planning frame*  $\Phi$ .

Given a *classical planning frame*  $\Phi = \langle F, A[\cdot] \rangle$  and an observation  $\mathcal{O}(\tau) = \langle s_0^o, s_1^o, \dots, s_m^o \rangle$  let  $P_{\mathcal{O}}$  be the classical planning problem  $P_{\mathcal{O}} = \langle F, A[\cdot], s_0^o, s_m^o \rangle$ .

**Definition 2 (Model explanation)** *A model  $\mathcal{M}$  explains an observation  $\mathcal{O}(\tau)$  iff, when the  $\langle \rho, \theta \rangle$  functions of the actions in  $P_{\mathcal{O}}$  are given by  $\mathcal{M}$ , there exists a solution plan for  $P_{\mathcal{O}}$  that explains  $\mathcal{O}(\tau)$ .*

**Definition 3 (The least-commitment action model)** *Given a  $\langle \Phi, \mathcal{O}(\tau) \rangle$  task (and let  $M$  be the set of action models that represents the full space of possible action models for the actions in  $A[\cdot] \in \Phi$ ), the least-commitment action model is the largest subset of models  $M^* \subseteq M$  such that every model  $\mathcal{M} \in M^*$  explains the input observation.*

This work focuses on the particular task of computing the *least-commitment* action model when action models are specified as STRIPS action schemata.

### 3.1 The space of STRIPS action models

A STRIPS *action schema*  $\xi$  is defined by four lists: A list of *parameters*  $pars(\xi)$ , and three list of predicates (namely  $pre(\xi)$ ,  $del(\xi)$  and  $add(\xi)$ ) that shape the kind of fluents that can appear in the *preconditions*, *negative effects* and *positive effects* of the actions induced from that schema. Let be  $\Psi$  the set of *predicates* that shape the propositional state variables  $F$ , and a list of *parameters*  $pars(\xi)$ . The set of elements that can appear in  $pre(\xi)$ ,  $del(\xi)$  and  $add(\xi)$  of the STRIPS action schema  $\xi$  is given by FOL interpretations of  $\Psi$  over the parameters  $pars(\xi)$ . We denote this set of FOL interpretations as  $\mathcal{I}_{\Psi, \xi}$ . For instance, in the *blocksworld* the  $\mathcal{I}_{\Psi, \xi}$  set contain eleven elements for the  $stack(v_1, v_2)$  schemata,  $\mathcal{I}_{\Psi, stack} = \{handempty, holding(v_1), holding(v_2), clear(v_1), clear(v_2), ontable(v_1), ontable(v_2), on(v_1, v_1), on(v_1, v_2), on(v_2, v_1), on(v_2, v_2)\}$ .

Despite any element of  $\mathcal{I}_{\Psi, \xi}$  can *a priori* appear in the  $pre(\xi)$ ,  $del(\xi)$  and  $add(\xi)$  of schema  $\xi$ , the space of possible STRIPS schemata is bounded by constraints of three kinds:

1. *Syntactic constraints.* STRIPS constraints require  $del(\xi) \subseteq pre(\xi)$ ,  $del(\xi) \cap add(\xi) = \emptyset$  and  $pre(\xi) \cap add(\xi) = \emptyset$ . Considering exclusively these syntactic constraints, the size of the space of possible STRIPS schemata is given by  $2^{2 \times |\mathcal{I}_{\Psi, \xi}|}$ . *Typing constraints* are also of this kind [McDermott et al., 1998].
2. *Domain-specific constraints.* One can introduce domain-specific knowledge to constrain further the space of possible schemata. For instance, in the *blocksworld* one can argue that  $on(v_1, v_1)$  and  $on(v_2, v_2)$  will not appear in the  $pre(\xi)$ ,  $del(\xi)$  and  $add(\xi)$  lists of an action schema  $\xi$  because, in this specific domain, a block cannot be on top of itself. *State invariants* are also constraints of this kind [Fox and Long, 1998].
3. *Observation constraints.* An observations  $\mathcal{O}(\tau)$  depicts *semantic knowledge* that constraints further the space of possible action schemata.

```

(:action stack
 :parameters (?v1 ?v2)
 :precondition (and (holding ?v1) (clear ?v2))
 :effect (and (not (holding ?v1)) (not (clear ?v2))
              (clear ?v1) (handempty) (on ?v1 ?v2)))

(pre_holding_v1_stack) (pre_clear_v2_stack)
(eff_holding_v1_stack) (eff_clear_v2_stack)
(eff_clear_v1_stack) (eff_handempty_stack) (eff_on_v1_v2_stack)

```

Figure 1: PDDL encoding of the `stack(?v1, ?v2)` schema and our propositional representation for this same schema.

In this work we introduce a propositional encoding of the *preconditions*, *negative*, and *positive* effects of a STRIPS action schema  $\xi$  using only fluents of two kinds  $\text{pre}_e\xi$  and  $\text{eff}_e\xi$  (where  $e \in \mathcal{I}_{\Psi, \xi}$ ). This encoding exploits the syntactic constraints of STRIPS so is more compact than the one previously proposed by Aineto *et al.* 2018. In more detail, if  $\text{pre}_e\xi$  and  $\text{eff}_e\xi$  holds it means that  $e \in \mathcal{I}_{\Psi, \xi}$  is a negative effect in  $\xi$  while if  $\text{pre}_e\xi$  does not hold but  $\text{eff}_e\xi$  holds, it means that  $e \in \mathcal{I}_{\Psi, \xi}$  is a positive effect in  $\xi$ . Figure 1 shows the PDDL encoding of the `stack(?v1, ?v2)` schema and our propositional representation for this same schema.

### 3.2 Partially specified STRIPS action models

A set of action models can be defined *explicitly*, enumerating all the models that belong to the set or *implicitly*, enumerating the constraints that must satisfy any model that belongs to the set. Inspired by the notion of *incomplete (annotated) model* [Sreedharan *et al.*, 2018], we introduce here *partially specified* STRIPS action models, a formalism for the *implicit* representation of a set of STRIPS schema.

We show now that extending to the *knowledge level* our propositional encoding of STRIPS action schemes, we can compactly represent a set of STRIPS schema. The extension defines, for each proposition  $\text{pre}_e\xi$ , two propositions  $K\text{pre}_e\xi$  and  $K\neg\text{pre}_e\xi$ , meaning that is *known* that  $e \in \mathcal{I}_{\Psi, \xi}$  is a precondition of  $\xi$  and that is *known* that  $e \in \mathcal{I}_{\Psi, \xi}$  is not a precondition of  $\xi$ . Likewise for each  $\text{eff}_e\xi$  proposition we define these two,  $K\text{eff}_e\xi$  and  $K\neg\text{eff}_e\xi$ , meaning that is *known* that  $e \in \mathcal{I}_{\Psi, \xi}$  is an effect of  $\xi$  and that is *known* that  $e \in \mathcal{I}_{\Psi, \xi}$  is not an effect of  $\xi$ .

**Definition 4 (Partially specified model)** A partially specified action schema  $\xi[\cdot]$  is an arbitrary formula over the  $K\text{pre}_e\xi$ ,  $K\neg\text{pre}_e\xi$ ,  $K\text{eff}_e\xi$  and  $K\neg\text{eff}_e\xi$  propositions ( $e \in \mathcal{I}_{\Psi, \xi}$ ) whose valuations express the enumeration of the set of action models  $M_\xi$ .

With respect to this formalization, the *full space* of possible STRIPS schemas for  $\xi$  is compactly represented by the *partially specified action schema*  $\bigcap_{e \in \mathcal{I}_{\Psi, \xi}} \neg K\text{pre}_e\xi \wedge \neg K\neg\text{pre}_e\xi \wedge \neg K\text{eff}_e\xi \wedge \neg K\neg\text{eff}_e\xi$ . Figure 2 shows a *partially specified* STRIPS model that represents of a set of action models for the `stack(?v1, ?v2)` schema. This *partially specified* STRIPS model compactly represents a set of  $2^5$  models since (1), the actual five effects of the action are *unknown*, (2) preconditions are known to appear as in the actual definition of the `stack(?v1, ?v2)` schema for the *blocksworld* and (3), it is *known* that the preconditions and effects that are not part of the actual definition of the

```

(Kpre_holding_v1_stack) (Kpre_clear_v2_stack)
(Knegpre_holding_v2_stack) (Knegpre_clear_v1_stack)
(Knegpre_handempty_stack)
(Knegpre_ontable_v1_stack) (Knegpre_ontable_v2_stack)
(Knegpre_on_v1_v2_stack) (Knegpre_on_v2_v1_stack)
(Knegpre_on_v1_v1_stack) (Knegpre_on_v2_v2_stack)

(Knegeff_holding_v2_stack) (Knegeff_clear_v2_stack)
(Knegeff_clear_v2_stack) (Knegeff_on_v2_v2_stack)
(Knegeff_on_v1_v1_stack) (Knegeff_on_v2_v1_stack)

```

Figure 2: *Partially specified* STRIPS model that represents of a set of action models for the `stack(?v1, ?v2)` schema.

`stack(?v1, ?v2)` schema are not part of this *partially specified* STRIPS model.

**Definition 5 (Partially specified model explanation)** A *partially specified action schema*  $\xi[\cdot]$  (that defines the set of action models  $M_\xi$ ) explains an observation  $\mathcal{O}(\tau)$ , iff every model  $\mathcal{M} \in M_\xi$  explains the input observation.

## 4 Incremental learning of action models from observations

This section shows how to implement an incremental learning of action models by pruning action models that are not consistent with the input observations.

### 4.1 Computing the least-commitment model via conformant planning

The compilation for learning STRIPS action models defined by Aineto *et al.* 2018 can be used to compute a *partially specified model* that explains a given observation by producing a conformant planning task instead of a classical planning task. The initial state  $\Upsilon$  of the conformant planning task does not code an *initial* action model (typically *empty*) but the full space of action models. In more detail, the clauses in  $\Upsilon$  comprises:

1. The *unit clauses* given by the fluents that hold in the initial state  $I = s_0$ .
2. The clauses representing that the actual value of fluents  $\text{pre}_e\xi$ ,  $\text{eff}_e\xi$  is unknown. In other words, that any model from the STRIPS space of models (following the previously mentioned *syntactic constraints*) can initially be part of the *least-commitment* action model. Formally, for every  $\xi$  and  $e \in \mathcal{I}_{\Psi, \xi}$ , then  $\Upsilon$  includes these two clauses:

- $\text{pre}_e\xi \vee \neg\text{pre}_e\xi$ .
- $\text{eff}_e\xi \vee \neg\text{eff}_e\xi$ .

One can also add here clauses that encode *domain-specific constraints* (as mentioned in the previous section) to make the conformant planning problem easier to be solved for a specific domain.

A conformant plan that solves this task induces a *partially specified model* that *explains* the input observation. To get the *least-commitment* action model, we have to compute the *optimal* plan for the resulting conformant planning task when: The actions that *program* a given precondition (or effect) of an action schema have an associated cost of 1. These actions

represent now a reduction of one unit in the *uncertainty* of the initial *partially specified model*. In other words, these actions implement the different possible *immediate specification* of the current *partially specified model*. The remaining actions have a cost of 0.

## 4.2 Computing the *least-commitment* model via classical planning

Inspired by the *classical planning compilation*  $K_0$  for conformant planning [Palacios and Geffner, 2009], this section shows that we can build a *classical planning problem*  $P = \langle F', A', I', G' \rangle$  whose optimal solution induces the *least-commitment* action model for an observation  $\mathcal{O}(\tau)$ :

- The set of fluents  $F'$  extends  $F$  with two new sets of fluents:
  - $\{test_j\}_{1 \leq j \leq m}$ , indicating the state observation  $s_j \in \mathcal{O}(\tau)$  where the action model is validated
  - The *knowledge level* fluents  $Kpre.e.\xi$ ,  $K\neg pre.e.\xi$ ,  $Keff.e.\xi$  and  $K\neg eff.e.\xi$  encoding the space of possible *partially specified* action models.
- The set of actions  $A'$  contains now actions of three different kinds:
  - Actions for *committing*  $pre.e.\xi$  to a positive/negative value. Similar actions are also defined for *committing*  $eff.e.\xi$  to a positive/negative value but the value of  $eff.e.\xi$  can only be committed once the value of the corresponding  $pre.e.\xi$  is committed (i.e. once either  $Kpre.e.\xi$  or  $K\neg pre.e.\xi$  holds in the current state).

$$\begin{aligned} pre(commit \top pre.e.\xi) &= \{mode_{commit}, \\ &\quad \neg Kpre.e.\xi, \neg K\neg pre.e.\xi\}, \\ cond(commit \top pre.e.\xi) &= \{\emptyset\} \triangleright \{Kpre.e.\xi\}. \end{aligned}$$

$$\begin{aligned} pre(commit \perp pre.e.\xi) &= \{mode_{commit}, \\ &\quad \neg Kpre.e.\xi, \neg K\neg pre.e.\xi\}, \\ cond(commit \perp pre.e.\xi) &= \{\emptyset\} \triangleright \{K\neg pre.e.\xi\}. \end{aligned}$$

- Actions for *validating* that committed models explain the  $s_j$  observed states,  $0 \leq j < m$ .

$$\begin{aligned} pre(validate_j) &= s_j \cup \{test_{j-1}\}, \\ cond(validate_j) &= \{\emptyset\} \triangleright \{\neg test_{j-1}, test_j\}, \\ &\quad \{mode_{commit}\} \triangleright \{\neg mode_{commit}, mode_{commit}\}. \end{aligned}$$

- *Editable* actions whose semantics is given by the value of the *knowledge level* fluents ( $Kpre.e.\xi$ ,  $K\neg pre.e.\xi$ ,  $Keff.e.\xi$  and  $K\neg eff.e.\xi$ ) at the current state. Figure 3 shows the PDDL encoding of an *editable* `stack(?v1, ?v2)` schema. This editable schema behaves exactly as the original PDDL schema defined in Figure 1 when the set of fluents ( $Kpre.holding.v1.stack$ ) ( $Kpre.clear.v2.stack$ ) ( $Keff.holding.v1.stack$ ) ( $Keff.clear.v2.stack$ ) ( $Keff.clear.v1.stack$ ) ( $Keff.handempty.stack$ ) ( $Keff.on.v1.v2.stack$ ) hold at the current state as well as the remaining  $K\neg pre.e.\xi$  and  $K\neg eff.e.\xi$  for

```
(:action stack
:parameters (?o1 - object ?o2 - object)
:precondition
  (and (or (Knotpre_on_v1_v1_stack) (on ?o1 ?o1))
        (or (Knotpre_on_v1_v2_stack) (on ?o1 ?o2))
        (or (Knotpre_on_v2_v1_stack) (on ?o2 ?o1))
        (or (Knotpre_on_v2_v2_stack) (on ?o2 ?o2))
        (or (Knotpre_ontable_v1_stack) (ontable ?o1))
        (or (Knotpre_ontable_v2_stack) (ontable ?o2))
        (or (Knotpre_clear_v1_stack) (clear ?o1))
        (or (Knotpre_clear_v2_stack) (clear ?o2))
        (or (Knotpre_holding_v1_stack) (holding ?o1))
        (or (Knotpre_holding_v2_stack) (holding ?o2))
        (or (Knotpre_handempty_stack) (handempty)))
:effect (and
  (when (and (Kpre_on_v1_v1_stack) (not (Knot_eff_on_v1_v1_stack))) (not (on ?o1 ?o1))
    (when (and (Kpre_on_v1_v2_stack) (not (Knot_eff_on_v1_v2_stack))) (not (on ?o1 ?o2))
      (when (and (Kpre_on_v2_v1_stack) (not (Knot_eff_on_v2_v1_stack))) (not (on ?o2 ?o1))
        (when (and (Kpre_on_v2_v2_stack) (not (Knot_eff_on_v2_v2_stack))) (not (on ?o2 ?o2))
          (when (and (Kpre_ontable_v1_stack) (not (Knot_eff_ontable_v1_stack))) (not (ontable ?o1))
            (when (and (Kpre_ontable_v2_stack) (not (Knot_eff_ontable_v2_stack))) (not (ontable ?o2))
              (when (and (Kpre_clear_v1_stack) (not (Knot_eff_clear_v1_stack))) (not (clear ?o1))
                (when (and (Kpre_clear_v2_stack) (not (Knot_eff_clear_v2_stack))) (not (clear ?o2))
                  (when (and (Kpre_holding_v1_stack) (not (Knot_eff_holding_v1_stack))) (not (holding ?o1))
                    (when (and (Kpre_holding_v2_stack) (not (Knot_eff_holding_v2_stack))) (not (holding ?o2))
                      (when (and (Kpre_handempty_stack) (not (Knot_eff_handempty_stack))) (not (handempty))
                        (when (and (Knot_pre_on_v1_v1_stack) (not (Knot_eff_on_v1_v1_stack))) (on ?o1 ?o1))
                        (when (and (Knot_pre_on_v1_v2_stack) (not (Knot_eff_on_v1_v2_stack))) (on ?o1 ?o2))
                        (when (and (Knot_pre_on_v2_v1_stack) (not (Knot_eff_on_v2_v1_stack))) (on ?o2 ?o1))
                        (when (and (Knot_pre_on_v2_v2_stack) (not (Knot_eff_on_v2_v2_stack))) (on ?o2 ?o2))
                        (when (and (Knot_pre_ontable_v1_stack) (not (Knot_eff_ontable_v1_stack))) (ontable ?o1))
                        (when (and (Knot_pre_ontable_v2_stack) (not (Knot_eff_ontable_v2_stack))) (ontable ?o2))
                        (when (and (Knot_pre_clear_v1_stack) (not (Knot_eff_clear_v1_stack))) (clear ?o1))
                        (when (and (Knot_pre_clear_v2_stack) (not (Knot_eff_clear_v2_stack))) (clear ?o2))
                        (when (and (Knot_pre_holding_v1_stack) (not (Knot_eff_holding_v1_stack))) (holding ?o1))
                        (when (and (Knot_pre_holding_v2_stack) (not (Knot_eff_holding_v2_stack))) (holding ?o2))
                        (when (and (Knot_pre_handempty_stack) (not (Knot_eff_handempty_stack))) (handempty))
```

Figure 3: PDDL encoding of the editable version of the `stack(?v1, ?v2)` schema.

all the preconditions and effects that are not part of the actual `stack(?v1, ?v2)` schema. Formally, given an operator schema  $\xi \in \mathcal{M}$  its *editable* version is:

$$\begin{aligned} pre(editable_\xi) &= \{\neg K\neg pre.e.\xi \implies e\}_{e \in \mathcal{I}_{\Psi, \xi}} \\ cond(editable_\xi) &= \{Kpre.e.\xi, \neg K\neg eff.e.\xi\} \triangleright \{\neg e\}_{e \in \mathcal{I}_{\Psi, \xi}}, \\ &\quad \{K\neg pre.e.\xi, \neg K\neg eff.e.\xi\} \triangleright \{e\}_{e \in \mathcal{I}_{\Psi, \xi}}. \end{aligned}$$

- The new initial state  $I' = I \cup \{mode_{commit}\}$  while the new goals are  $G' = s_m \cup \{test_m\}$ .

## 4.3 Compilation properties

## 4.4 Learning algorithm

Given a *classical planning frame*  $\Phi = \langle F, A[\cdot] \rangle$  where the semantics of each action  $a \in A[\cdot]$  is unknown and a set of state observations  $\mathcal{O}_1, \dots, \mathcal{O}_T$ . Our learning algorithm for the incremental learning of action models from state observations is defined as follows:

1. Build the *least-commitment* task  $\langle \Phi, \mathcal{O}_t \rangle$  to process observation  $\mathcal{O}_t$ ,  $1 \leq t \leq T$ .
2. Build the classical planning task  $P = \langle F', A', I', G' \rangle$  to solve the previous task.
3. Extract the *partially specified* model from the previous solution.
4. Continue until, the *version space* is a singleton or no more observations left

Assuming a set of independent observations  $\{\mathcal{O}_1, \dots, \mathcal{O}_t\}$  we need to see at least  $\frac{\ln(1-\delta)+\ln n}{\epsilon}$  to be sure that the prediction error is at most  $\epsilon$  with probability  $(1 - \delta)$ .

## 5 Evaluation

## 6 Conclusions

Related work [Stern and Juba, 2017].

## References

- [Aineto *et al.*, 2018] Diego Aineto, Sergio Jiménez, and Eva Onaindia. Learning STRIPS action models with classical planning. In *International Conference on Automated Planning and Scheduling, (ICAPS-18)*, pages 399–407. AAAI Press, 2018.
- [Bonet and Geffner, 2000] Blai Bonet and Hector Geffner. Planning with incomplete information as heuristic search in belief space. In *Proceedings of the Fifth International Conference on Artificial Intelligence Planning Systems*, pages 52–61. AAAI Press, 2000.
- [Fox and Long, 1998] Maria Fox and Derek Long. The automatic inference of state invariants in tim. *Journal of Artificial Intelligence Research*, 9:367–421, 1998.
- [Goldman and Boddy, 1996] Robert P Goldman and Mark S Boddy. Expressive planning and explicit knowledge. In *AIPS*, volume 96, pages 110–117, 1996.
- [McDermott *et al.*, 1998] Drew McDermott, Malik Ghallab, Adele Howe, Craig Knoblock, Ashwin Ram, Manuela Veloso, Daniel Weld, and David Wilkins. PDDL – The Planning Domain Definition Language, 1998.
- [Mitchell, 1982] Tom M Mitchell. Generalization as search. *Artificial intelligence*, 18(2):203–226, 1982.
- [Palacios and Geffner, 2009] Héctor Palacios and Héctor Geffner. Compiling uncertainty away in conformant planning problems with bounded width. *Journal of Artificial Intelligence Research*, 35:623–675, 2009.
- [Ramírez and Geffner, 2009] Miquel Ramírez and Hector Geffner. Plan recognition as planning. In *International Joint conference on Artificial Intelligence, (IJCAI-09)*, pages 1778–1783. AAAI Press, 2009.
- [Smith and Weld, 1998] David E Smith and Daniel S Weld. Conformant graphplan. In *AAAI/IAAI*, pages 889–896, 1998.
- [Sreedharan *et al.*, 2018] Sarath Sreedharan, Tathagata Chakraborti, and Subbarao Kambhampati. Handling model uncertainty and multiplicity in explanations via model reconciliation. In *International Conference on Automated Planning and Scheduling, (ICAPS-18)*, pages 518–526, 2018.
- [Stern and Juba, 2017] Roni Stern and Brendan Juba. Efficient, safe, and probably approximately complete learning of action models. In *International Joint Conference on Artificial Intelligence, (IJCAI-17)*, pages 4405–4411, 2017.