

Computing the *least-commitment* action model from state observations

Diego Aineto¹, Sergio Jiménez¹, Eva Onaindia¹ and , Blai Bonet²

¹Departamento de Sistemas Informáticos y Computación. Universitat Politècnica de València. Valencia, Spain

²Departamento de Computación. Universidad Simón Bolívar. Caracas, Venezuela

{dieaigar,serjice,onaindia}@dsic.upv.es, bonet@usb.ve

Abstract

1 Introduction

Given a sequence of partially observed states, this paper formalizes the task of computing the *least-commitment* action model that is able to *explain* the given observation. This task is of interest because it allows the incremental learning of action models from arbitrary large sets of partial observations.

In addition, the paper introduces a new method to compute the *least-commitment* action model from a sequence of partially observed states. The method assumes that action models are specified as STRIPS action schemata and it is built on top of off-the-shelf algorithms for *conformant planning*.

2 Background

This section formalizes the planning models we use in the paper as well as the kind of observations that are given as input for the computation of the *least-commitment* action model.

2.1 Classical planning with conditional effects

Let F be the set of *fluents* or *state variables* (propositional variables). A *literal* l is a valuation of a fluent $f \in F$, i.e. either $l = f$ or $l = \neg f$. L is a set of literals that represents a partial assignment of values to fluents, and $\mathcal{L}(F)$ is the set of all literals sets on F , i.e. all partial assignments of values to fluents. A *state* s is a full assignment of values to fluents. We explicitly include negative literals $\neg f$ in states s.t. $|s| = |F|$ and the size of the state space is $2^{|F|}$.

A *classical planning frame* is a tuple $\Phi = \langle F, A \rangle$, where F is a set of fluents and A is a set of *actions*. An action $a \in A$ is defined with *preconditions*, $\text{pre}(a) \in \mathcal{L}(F)$, *positive effects*, $\text{eff}^+(a) \in \mathcal{L}(F)$, and *negative effects* $\text{eff}^-(a) \in \mathcal{L}(F)$. The semantics of actions $a \in A$ is specified with two functions: $\rho(s, a)$ denotes whether action a is *applicable* in a state s and $\theta(s, a)$ denotes the *successor state* that results of applying action a in a state s . Then, $\rho(s, a)$ holds iff $\text{pre}(a) \subseteq s$. And the result of applying a in s is $\theta(s, a) = \{s \setminus \text{eff}^-(a)\} \cup \text{eff}^+(a)$.

A *classical planning problem* is a tuple $P = \langle F, A, I, G \rangle$, where I is the initial state and $G \in \mathcal{L}(F)$ is the set of goal

conditions over the state variables. A *plan* π is an action sequence $\pi = \langle a_1, \dots, a_n \rangle$, with $|\pi| = n$ denoting its *plan length*. The execution of π on the initial state I of P induces a *trajectory* $\tau(\pi, s_0) = \langle s_0, a_1, s_1, \dots, a_n, s_n \rangle$ such that $s_0 = I$ and, for each $1 \leq i \leq n$, it holds $\rho(s_{i-1}, a_i)$ and $s_i = \theta(s_{i-1}, a_i)$. A plan π solves P iff the induced *trajectory* $\tau(\pi, s_0)$ reaches a final state $G \subseteq s_n$, where all goal conditions are met. A solution plan is optimal iff its length is minimal.

An action $a_c \in A$ with conditional effects is defined as a set of preconditions $\text{pre}(a_c) \in \mathcal{L}(F)$ and a set of *conditional effects* $\text{cond}(a_c)$. Each conditional effect $C \triangleright E \in \text{cond}(a_c)$ is composed of two sets of literals: $C \in \mathcal{L}(F)$, the *condition*, and $E \in \mathcal{L}(F)$, the *effect*. An action a_c is applicable in a state s if $\rho(s, a_c)$ is true, and the *triggered effects* resulting from the action application are the effects whose conditions hold in s :

$$\text{triggered}(s, a_c) = \bigcup_{C \triangleright E \in \text{cond}(a_c), C \subseteq s} E,$$

The result of applying action a_c in state s is $\theta(s, a_c) = \{s \setminus \text{eff}_c^-(s, a)\} \cup \text{eff}_c^+(s, a)\}$, where $\text{eff}_c^-(s, a) \subseteq \text{triggered}(s, a)$ and $\text{eff}_c^+(s, a) \subseteq \text{triggered}(s, a)$ are, respectively, the triggered *negative* and *positive* effects.

2.2 The observation model

Given a classical planning problem $P = \langle F, A, I, G \rangle$, a plan π and a trajectory $\tau(\pi, s_0)$, we define the *observation of the trajectory* as a sequence of partial states that results from observing the execution of π on I . Formally, $\mathcal{O}(\tau) = \langle s_0^o, s_1^o, \dots, s_m^o \rangle$ where $s_0^o = I$.

A partially observable state s_i^o , $0 < i < m$, is one in which $|s_i^o| < |F|$; i.e., a state in which at least a fluent of F is not observable. Note that this definition also comprises the case $|s_i^o| = 0$, when the state is fully unobservable. Whatever the sequence of observed states of $\mathcal{O}(\tau)$ is, it must be consistent with the sequence of states of $\tau(\pi, s_0)$, meaning that $\forall i, s_i^o \subseteq s_i$. In practice, the number of observed states m , ranges from 1 (the initial state, at least), to $|\pi| + 1$, and the observed intermediate states will comprise a number of fluents between $[1, |F|]$.

In other words, we assume there is a *bijective monotone mapping* between trajectories and observations [Ramírez and Geffner, 2009], thus also granting the inverse consistency relationship (the trajectory is a superset of the observation). Therefore, transiting between two consecutive observed states in $\mathcal{O}(\tau)$ may require the execution of more than a single action ($\theta(s_i^o, \langle a_1, \dots, a_k \rangle) = s_{i+1}^o$, where $k \geq 1$ is unknown but finite. In other words, having $\mathcal{O}(\tau)$ does not imply knowing the actual length of π .

Definition 1 (Explaining a $\mathcal{O}(\tau)$ observation) *Given a classical planning problem P and a sequence of partially observed states $\mathcal{O}(\tau)$, we say that a plan π explains the observation (denoted $\pi \mapsto \mathcal{O}(\tau)$) iff π is a solution for P that is consistent with $\mathcal{O}(\tau)$. If π is also optimal, we say that π is the best explanation for $\mathcal{O}(\tau)$.*

2.3 Conformant planning

Conformant planning is planning with incomplete information about the initial state, no sensing, and validating that goals are achieved with certainty (despite the uncertainty of the initial state) [Goldman and Boddy, 1996; Smith and Weld, 1998; Bonet and Geffner, 2000].

Syntactically, conformant planning problems are expressed in compact form through a set of state variables. A *conformant planning problem* can be defined as a tuple $P_c = \langle F, A, \Upsilon, G \rangle$ where F , A and G are the set of *fluents*, *actions* and *goals* (as previously defined for *classical planning*). Now Υ is a set of clauses over literals $l = f$ or $l = \neg f$ (for $f \in F$) that define the set of possible initial states.

A solution to a conformant planning problem is an action sequence that maps each possible initial state into a goal state. More precisely, an action sequence $\pi = \langle a_1, \dots, a_n \rangle$ is a *conformant plan* for P_c iff, for each possible trajectory $\tau(\pi, s_0) = \langle s_0, a_1, s_1, \dots, a_n, s_n \rangle$ s.t. s_0 is a valuation of the fluents in F that satisfies Υ , then $\tau(\pi, s_0)$ reaches a final state $G \subseteq s_n$.

3 Computing the least-commitment action model from state observations

First, this section formalizes the notion of the *least-commitment* action model that is able to *explain* a sequence of partially observed states. Next, the section describes our approach to compute such model via *conformant planning*.

3.1 The least-commitment action model

The task of computing the *least-commitment* action model from a sequence of state observations is defined as $\langle \Phi, \mathcal{O}(\tau) \rangle$:

- $\Phi = \langle F, A[\cdot] \rangle$ is a *classical planning frame* where $A[\cdot]$ is a set of actions s.t. the semantics of each $a \in A[\cdot]$ is unknown; i.e. the corresponding $\langle \rho, \theta \rangle$ functions are undefined.
- $\mathcal{O}(\tau)$ is a sequence of partial states that results from the partial observation of a trajectory $\tau(\pi, s_0)$ that is defined within the *classical planning frame* Φ .

Before formalizing the solution to this task, i.e. the *least-commitment* action model, we introduce several necessary definitions. Given a *classical planning frame*

$\Phi = \langle F, A[\cdot] \rangle$ and a sequence of partial states $\mathcal{O}(\tau) = \langle s_0^o, s_1^o, \dots, s_m^o \rangle$, we can build the classical planning problem $P_{\mathcal{O}} = \langle F, A[\cdot], s_0^o, s_m^o \rangle$. We say that an action model \mathcal{M} is a definition of the $\langle \rho, \theta \rangle$ functions of every action in $A[\cdot]$. Further we say that a model \mathcal{M} *explains* a sequence of observations $\mathcal{O}(\tau)$ iff, when the $\langle \rho, \theta \rangle$ functions of the actions in $P_{\mathcal{O}}$ are given by \mathcal{M} , there exists a solution plan for $P_{\mathcal{O}}$ that explains $\mathcal{O}(\tau)$.

Now, we define a *partially specified action model* inspired by the notion of *incomplete (annotated) model* [Sreedharan et al., 2018].

Definition 2 (Partially specified action model) *Given a classical planning frame $\Phi = \langle F, A[\cdot] \rangle$ then, a partially specified action model M is a set of possible models for the actions in $A[\cdot]$ such that: (1), any model $\mathcal{M} \in M$ is a definition of the $\langle \rho, \theta \rangle$ functions of every action in $A[\cdot]$ and (2), for every $\mathcal{M} \in M$ the $\langle \rho, \theta \rangle$ functions are defined in the set of state variables F . (Note that if M is a singleton it represents a fully specified action model).*

Now we are ready to define the *least-commitment* action model for an observation $\mathcal{O}(\tau)$.

Definition 3 (The least-commitment action model) *Given a $\langle \Phi, \mathcal{O}(\tau) \rangle$ task and the partially specified action model M that represents the full space of possible action models for the actions in $A[\cdot] \in \Phi$, then the least-commitment action model is another partially specified action model that represents the largest subset of models $M^* \subseteq M$ such that every model $\mathcal{M} \in M^*$ explains the input observation.*

3.2 The space of STRIPS action models

Despite previous definitions are general, this work focuses on the particular kind of action models that are specified as STRIPS action schemata.

A STRIPS *action schema* ξ is defined by four lists: A list of *parameters* $\text{pars}(\xi)$, and three list of predicates (namely $\text{pre}(\xi)$, $\text{del}(\xi)$ and $\text{add}(\xi)$) that shape the kind of fluents that can appear in the *preconditions*, *negative effects* and *positive effects* of the actions induced from that schema. Let be Ψ the set of *predicates* that shape the propositional state variables F , and a list of *parameters* $\text{pars}(\xi)$. The set of elements that can appear in $\text{pre}(\xi)$, $\text{del}(\xi)$ and $\text{add}(\xi)$ of the STRIPS action schema ξ is given by FOL interpretations of Ψ over the parameters $\text{pars}(\xi)$. We denote this set of FOL interpretations as $\mathcal{I}_{\Psi, \xi}$. For instance, in the *blocksworld* the $\mathcal{I}_{\Psi, \xi}$ set contain eleven elements for the $\text{stack}(v_1, v_2)$ schemata, $\mathcal{I}_{\Psi, \text{stack}} = \{\text{handempty}, \text{holding}(v_1), \text{holding}(v_2), \text{clear}(v_1), \text{clear}(v_2), \text{ontable}(v_1), \text{ontable}(v_2), \text{on}(v_1, v_1), \text{on}(v_1, v_2), \text{on}(v_2, v_1), \text{on}(v_2, v_2)\}$.

Despite any element of $\mathcal{I}_{\Psi, \xi}$ can *a priori* appear in the $\text{pre}(\xi)$, $\text{del}(\xi)$ and $\text{add}(\xi)$ of schema ξ , the space of possible STRIPS schemata is bounded by constraints of three kinds:

1. *Syntactic constraints.* STRIPS constraints require $\text{del}(\xi) \subseteq \text{pre}(\xi)$, $\text{del}(\xi) \cap \text{add}(\xi) = \emptyset$ and $\text{pre}(\xi) \cap \text{add}(\xi) = \emptyset$. Considering exclusively these syntactic constraints, the size of the space of possible STRIPS schemata is given by $2^{2 \times |\mathcal{I}_{\Psi, \xi}|}$.

```

(:action stack
:parameters (?v1 ?v2)
:precondition (and (holding ?v1) (clear ?v2))
:effect (and (not (holding ?v1)) (not (clear ?v2))
             (clear ?v1) (handempty) (on ?v1 ?v2)))

(pre_holding_v1_stack) (pre_clear_v2_stack)
(eff_holding_v1_stack) (eff_clear_v2_stack)
(eff_clear_v1_stack) (eff_handempty_stack) (eff_on_v1_v2_stack)

```

Figure 1: PDDL encoding of the `stack(?v1, ?v2)` schema and our propositional representation for this same schema.

2. *Domain-specific constraints.* One can introduce domain-specific knowledge to constrain further the space of possible schemata. For instance, in the *blocksworld* one can argue that $\text{on}(v_1, v_1)$ and $\text{on}(v_2, v_2)$ will not appear in the $\text{pre}(\xi)$, $\text{del}(\xi)$ and $\text{add}(\xi)$ lists of an action schema ξ because, in this specific domain, a block cannot be on top of itself. As a rule of thumb, *state invariants* constraining the possible states of a given planning domain belong to this second class of constraints [Fox and Long, 1998].
3. *Observation constraints.* A sequence of state observations $\mathcal{O}(\tau)$ depicts *semantic knowledge* that constraints further the space of possible action schemata.

3.3 Computing the least-commitment model via conformant planning

Given a task $\langle \Phi, \mathcal{O}(\tau) \rangle$, this section shows that we can build a *conformant planning problem* P_c whose solution induces the *least-commitment* action model for the input observation $\mathcal{O}(\tau)$. In more detail, we build a *conformant planning problem* $P_c = \langle F_c, A_c, \Upsilon, G_c \rangle$ such that:

- The set of fluents F_c extends F with two new sets of fluents:
 - $\{test_j\}_{1 \leq j \leq m}$, indicating the state observation $s_j \in \mathcal{O}(\tau)$ where the action model is validated
 - Fluents $\text{pre}_e.\xi$ and $\text{eff}_e.\xi$ (where $e \in \mathcal{I}_{\Psi, \xi}$) implementing a propositional encoding of the *preconditions*, *negative*, and *positive* effects of an action schema ξ . **Our encoding exploits the syntactic constraint of STRIPS** so, if $\text{pre}_e.\xi$ and $\text{eff}_e.\xi$ holds it means that $e \in \mathcal{I}_{\Psi, \xi}$ is a negative effect in ξ while if $\text{pre}_e.\xi$ does not hold but $\text{eff}_e.\xi$ holds, it means that $e \in \mathcal{I}_{\Psi, \xi}$ is a positive effect in ξ . Figure 1 shows the PDDL encoding of the `stack(?v1, ?v2)` schema and our propositional representation for this same schema.
- The set of actions A_c contains now actions of three different kinds:
 - Actions for *committing* $\text{pre}_e.\xi$ fluents to a positive/negative value (similar actions are also defined for *committing* $\text{eff}_e.\xi$ fluents to a posi-

tive/negative value).

$$\begin{aligned} \text{pre}(\text{commit}_\top.\text{pre}_e.\xi) &= \{mode_{commit}\}, \\ \text{cond}(\text{commit}_\top.\text{pre}_e.\xi) &= \{\text{pre}_e.\xi\} \triangleright \{\text{pre}_e.\xi\}, \\ &\quad \{\neg \text{pre}_e.\xi\} \triangleright \{\text{pre}_e.\xi\}. \end{aligned}$$

$$\begin{aligned} \text{pre}(\text{commit}_\perp.\text{pre}_e.\xi) &= \{mode_{commit}\}, \\ \text{cond}(\text{commit}_\perp.\text{pre}_e.\xi) &= \{\text{pre}_e.\xi\} \triangleright \{\neg \text{pre}_e.\xi\}, \\ &\quad \{\neg \text{pre}_e.\xi\} \triangleright \{\neg \text{pre}_e.\xi\}. \end{aligned}$$

- Actions for *validating* that committed models explain the s_j observed states, $0 \leq j < m$.

$$\begin{aligned} \text{pre}(\text{validate}_j) &= s_j \cup \{test_{j-1}\}, \\ \text{cond}(\text{validate}_j) &= \{\emptyset\} \triangleright \{\neg test_{j-1}, test_j, \\ &\quad \{mode_{commit}\} \triangleright \{\neg mode_{commit}, mode_{val}\}\}. \end{aligned}$$

- *Editable* actions whose semantics is given by the value of $\text{pre}_e.\xi$, $\text{eff}_e.\xi$ fluents at the current state. Figure 2 shows the PDDL encoding of an *editable* `stack(?v1, ?v2)` schema. Note that this *editable* schema when the set of fluents $(\text{pre_holding_v1_stack})$ $(\text{pre_clear_v2_stack})$ $(\text{eff_holding_v1_stack})$ $(\text{eff_clear_v2_stack})$ $(\text{eff_clear_v1_stack})$ $(\text{eff_handempty_stack})$ $(\text{eff_on_v1_v2_stack})$ hold at the current state, then it behaves exactly as the original PDDL schema defined in Figure 1. Formally, given an operator schema $\xi \in \mathcal{M}$ its *editable* version is:

$$\begin{aligned} \text{pre}(\text{editable}_\xi) &= \{\text{pre}_e.\xi \implies e\}_{e \in \mathcal{I}_{\Psi, \xi}}, \\ \text{cond}(\text{editable}_\xi) &= \{\text{pre}_e.\xi, \text{eff}_e.\xi\} \triangleright \{\neg e\}_{e \in \mathcal{I}_{\Psi, \xi}}, \\ &\quad \{\neg \text{pre}_e.\xi, \text{eff}_e.\xi\} \triangleright \{e\}_{e \in \mathcal{I}_{\Psi, \xi}}. \end{aligned}$$

- The clauses in Υ comprises:

1. The *unit clauses* given by the fluents that hold in the initial state $I = s_0$ and $mode_{commit}$ set to true.
2. The clauses representing that the actual value of fluents $\text{pre}_e.\xi$, $\text{eff}_e.\xi$ is unknown. In other words, that any model from the STRIPS space of models (following the previously mentioned syntactic constraints) can initially be part of the *least-commitment* action model. Formally, for every ξ and $e \in \mathcal{I}_{\Psi, \xi}$, then Υ includes these two clauses:
 - $\text{pre}_e.\xi \vee \neg \text{pre}_e.\xi$.
 - $\text{eff}_e.\xi \vee \neg \text{eff}_e.\xi$.

One can also add here clauses that encode *domain-specific constraints* (as mentioned in the previous section) to make the conformant planning problem easier for a specific domain.

- The new goals are $G_c = s_m \cup \{test_m\}$.

3.4 Optimization of the compilation

In fact, we do not need to compute the *least-commitment* action model (i.e. solve the P_c conformant planning problem)

```

(:action stack
:parameters (?o1 - object ?o2 - object)
:precondition
  (and (or (not (pre_on_v1_v1_stack)) (on ?o1 ?o1))
        (or (not (pre_on_v1_v2_stack)) (on ?o1 ?o2))
        (or (not (pre_on_v2_v1_stack)) (on ?o2 ?o1))
        (or (not (pre_on_v2_v2_stack)) (on ?o2 ?o2))
        (or (not (pre_ontable_v1_stack)) (ontable ?o1))
        (or (not (pre_ontable_v2_stack)) (ontable ?o2))
        (or (not (pre_clear_v1_stack)) (clear ?o1))
        (or (not (pre_clear_v2_stack)) (clear ?o2))
        (or (not (pre_holding_v1_stack)) (holding ?o1))
        (or (not (pre_holding_v2_stack)) (holding ?o2))
        (or (not (pre_handepty_stack)) (handepty)))
:effect (and
  (when (and (pre_on_v1_v1_stack) (eff_on_v1_v1_stack)) (not (on ?o1 ?o1)))
  (when (and (pre_on_v1_v2_stack) (eff_on_v1_v2_stack)) (not (on ?o1 ?o2)))
  (when (and (pre_on_v2_v1_stack) (eff_on_v2_v1_stack)) (not (on ?o2 ?o1)))
  (when (and (pre_on_v2_v2_stack) (eff_on_v2_v2_stack)) (not (on ?o2 ?o2)))
  (when (and (pre_ontable_v1_stack) (eff_ontable_v1_stack)) (not (ontable ?o1)))
  (when (and (pre_ontable_v2_stack) (eff_ontable_v2_stack)) (not (ontable ?o2)))
  (when (and (pre_clear_v1_stack) (eff_clear_v1_stack)) (not (clear ?o1)))
  (when (and (pre_clear_v2_stack) (eff_clear_v2_stack)) (not (clear ?o2)))
  (when (and (pre_holding_v1_stack) (eff_holding_v1_stack)) (not (holding ?o1)))
  (when (and (pre_holding_v2_stack) (eff_holding_v2_stack)) (not (holding ?o2)))
  (when (and (pre_handepty_stack) (eff_handepty_stack)) (not (handepty)))
  (when (and (not (pre_on_v1_v1_stack)) (eff_on_v1_v1_stack)) (on ?o1 ?o1))
  (when (and (not (pre_on_v1_v2_stack)) (eff_on_v1_v2_stack)) (on ?o1 ?o2))
  (when (and (not (pre_on_v2_v1_stack)) (eff_on_v2_v1_stack)) (on ?o2 ?o1))
  (when (and (not (pre_on_v2_v2_stack)) (eff_on_v2_v2_stack)) (on ?o2 ?o2))
  (when (and (not (pre_ontable_v1_stack)) (eff_ontable_v1_stack)) (ontable ?o1))
  (when (and (not (pre_ontable_v2_stack)) (eff_ontable_v2_stack)) (ontable ?o2))
  (when (and (not (pre_clear_v1_stack)) (eff_clear_v1_stack)) (clear ?o1))
  (when (and (not (pre_clear_v2_stack)) (eff_clear_v2_stack)) (clear ?o2))
  (when (and (not (pre_holding_v1_stack)) (eff_holding_v1_stack)) (holding ?o1))
  (when (and (not (pre_holding_v2_stack)) (eff_holding_v2_stack)) (holding ?o2))
  (when (and (not (pre_handepty_stack)) (eff_handepty_stack)) (handepty)))

```

Figure 2: PDDL encoding of the editable version of the stack (?v1, ?v2) schema.

from scratch for all the $\langle s_0^o, s_1^o \dots, s_m^o \rangle$ observations. A more efficient approach is to compute the *least-commitment* action model from scratch only for the first observation, that is for $\langle s_0^o, s_1^o \rangle$. Then, the obtained *least-commitment* action model can be used as input for improving the computation of the two first observations $\langle s_0^o, s_1^o, s_2^o \rangle$. In other words, including the commits found in the previous step as new *unit clauses* in set of the *initial clauses* Υ . This process is repeated until the *least-commitment* action model is computed for the full sequence of state observations $\langle s_0^o, s_1^o \dots, s_m^o \rangle$.

3.5 Compilation properties

4 Evaluation

5 Conclusions

Related work [Stern and Juba, 2017].

References

- [Bonet and Geffner, 2000] Blai Bonet and Hector Geffner. Planning with incomplete information as heuristic search in belief space. In *Proceedings of the Fifth International Conference on Artificial Intelligence Planning Systems*, pages 52–61. AAAI Press, 2000.
- [Fox and Long, 1998] Maria Fox and Derek Long. The automatic inference of state invariants in tim. *Journal of Artificial Intelligence Research*, 9:367–421, 1998.
- [Goldman and Boddy, 1996] Robert P Goldman and Mark S Boddy. Expressive planning and explicit knowledge. In *AIPS*, volume 96, pages 110–117, 1996.

[Ramírez and Geffner, 2009] Miquel Ramírez and Hector Geffner. Plan recognition as planning. In *International Joint conference on Artificial Intelligence, (IJCAI-09)*, pages 1778–1783. AAAI Press, 2009.

[Smith and Weld, 1998] David E Smith and Daniel S Weld. Conformant graphplan. In *AAAI/IAAI*, pages 889–896, 1998.

[Sreedharan *et al.*, 2018] Sarath Sreedharan, Tathagata Chakraborti, and Subbarao Kambhampati. Handling model uncertainty and multiplicity in explanations via model reconciliation. In *International Conference on Automated Planning and Scheduling, (ICAPS-18)*, pages 518–526, 2018.

[Stern and Juba, 2017] Roni Stern and Brendan Juba. Efficient, safe, and probably approximately complete learning of action models. In *International Joint Conference on Artificial Intelligence, (IJCAI-17)*, pages 4405–4411, 2017.