

# Computing the *least-commitment* action model from state observations

Diego Aineto<sup>1</sup>, Sergio Jiménez<sup>1</sup>, Eva Onaindia<sup>1</sup> and , Blai Bonet<sup>2</sup>

<sup>1</sup>Departamento de Sistemas Informáticos y Computación. Universitat Politècnica de València. Valencia, Spain

<sup>2</sup>Departamento de Computación. Universidad Simón Bolívar. Caracas, Venezuela

{dieaigar,serjice,onaindia}@dsic.upv.es, bonet@usb.ve

## Abstract

## 1 Introduction

Given a sequence of partially observed states, this paper formalizes the task of computing the *least-commitment* action model that is able to *explain* the given observation. This task is of interest because it allows the incremental learning of action models from arbitrary large sets of partial observations.

In addition, the paper introduces a new method to compute the *least-commitment* action model from a sequence of partially observed states. The method assumes that action models are specified as STRIPS action schemata and it is built on top of off-the-shelf algorithms for *conformant planning*.

## 2 Background

This section formalizes the planning models we use in the paper as well as the kind of observations that are given as input for the computation of the *least-commitment* action model.

### 2.1 Classical planning with conditional effects

Let  $F$  be the set of *fluents* or *state variables* (propositional variables). A *literal*  $l$  is a valuation of a fluent  $f \in F$ , i.e. either  $l = f$  or  $l = \neg f$ .  $L$  is a set of literals that represents a partial assignment of values to fluents, and  $\mathcal{L}(F)$  is the set of all literals sets on  $F$ , i.e. all partial assignments of values to fluents. A *state*  $s$  is a full assignment of values to fluents. We explicitly include negative literals  $\neg f$  in states s.t.  $|s| = |F|$  and the size of the state space is  $2^{|F|}$ .

A *classical planning frame* is a tuple  $\Phi = \langle F, A \rangle$ , where  $F$  is a set of fluents and  $A$  is a set of *actions*. An action  $a \in A$  is defined with *preconditions*,  $\text{pre}(a) \in \mathcal{L}(F)$ , *positive effects*,  $\text{eff}^+(a) \in \mathcal{L}(F)$ , and *negative effects*  $\text{eff}^-(a) \in \mathcal{L}(F)$ . The semantics of actions  $a \in A$  is specified with two functions:  $\rho(s, a)$  denotes whether action  $a$  is *applicable* in a state  $s$  and  $\theta(s, a)$  denotes the *successor state* that results of applying action  $a$  in a state  $s$ . Then,  $\rho(s, a)$  holds iff  $\text{pre}(a) \subseteq s$ . And the result of applying  $a$  in  $s$  is  $\theta(s, a) = \{s \setminus \text{eff}^-(a)\} \cup \text{eff}^+(a)$ .

A *classical planning problem* is a tuple  $P = \langle F, A, I, G \rangle$ , where  $I$  is the initial state and  $G \in \mathcal{L}(F)$  is the set of goal

conditions over the state variables. A *plan*  $\pi$  is an action sequence  $\pi = \langle a_1, \dots, a_n \rangle$ , with  $|\pi| = n$  denoting its *plan length*. The execution of  $\pi$  on the initial state  $I$  of  $P$  induces a *trajectory*  $\tau(\pi, s_0) = \langle s_0, a_1, s_1, \dots, a_n, s_n \rangle$  such that  $s_0 = I$  and, for each  $1 \leq i \leq n$ , it holds  $\rho(s_{i-1}, a_i)$  and  $s_i = \theta(s_{i-1}, a_i)$ . A plan  $\pi$  solves  $P$  iff the induced *trajectory*  $\tau(\pi, s_0)$  reaches a final state  $G \subseteq s_n$ , where all goal conditions are met. A solution plan is optimal iff its length is minimal.

An action  $a_c \in A$  with conditional effects is defined as a set of preconditions  $\text{pre}(a_c) \in \mathcal{L}(F)$  and a set of *conditional effects*  $\text{cond}(a_c)$ . Each conditional effect  $C \triangleright E \in \text{cond}(a_c)$  is composed of two sets of literals:  $C \in \mathcal{L}(F)$ , the *condition*, and  $E \in \mathcal{L}(F)$ , the *effect*. An action  $a_c$  is applicable in a state  $s$  if  $\rho(s, a_c)$  is true, and the *triggered effects* resulting from the action application are the effects whose conditions hold in  $s$ :

$$\text{triggered}(s, a_c) = \bigcup_{C \triangleright E \in \text{cond}(a_c), C \subseteq s} E,$$

The result of applying action  $a_c$  in state  $s$  is  $\theta(s, a_c) = \{s \setminus \text{eff}_c^-(s, a) \cup \text{eff}_c^+(s, a)\}$ , where  $\text{eff}_c^-(s, a) \subseteq \text{triggered}(s, a)$  and  $\text{eff}_c^+(s, a) \subseteq \text{triggered}(s, a)$  are, respectively, the triggered *negative* and *positive* effects.

### 2.2 The observation model

Given a classical planning problem  $P = \langle F, A, I, G \rangle$ , a plan  $\pi$  and a trajectory  $\tau(\pi, s_0)$ , we define the *observation of the trajectory* as a sequence of partial states that results from observing the execution of  $\pi$  on  $I$ . Formally,  $\mathcal{O}(\tau) = \langle s_0^o, s_1^o, \dots, s_m^o \rangle$  where  $s_0^o = I$ .

A partially observable state  $s_i^o$ ,  $0 < i < m$ , is one in which  $|s_i^o| < |F|$ ; i.e., a state in which at least a fluent of  $F$  is not observable. Note that this definition also comprises the case  $|s_i^o| = 0$ , when the state is fully unobservable. Whatever the sequence of observed states of  $\mathcal{O}(\tau)$  is, it must be consistent with the sequence of states of  $\tau(\pi, s_0)$ , meaning that  $\forall i, s_i^o \subseteq s_i$ . In practice, the number of observed states  $m$ , ranges from 1 (the initial state, at least), to  $|\pi| + 1$ , and the observed intermediate states will comprise a number of fluents between  $[1, |F|]$ .

In other words, we assume there is a *bijective monotone mapping* between trajectories and observations [?], thus also granting the inverse consistency relationship (the trajectory is

a superset of the observation). Therefore, transiting between two consecutive observed states in  $\mathcal{O}(\tau)$  may require the execution of more than a single action ( $\theta(s_i^o, \langle a_1, \dots, a_k \rangle) = s_{i+1}^o$ , where  $k \geq 1$  is unknown but finite. In other words, having  $\mathcal{O}(\tau)$  does not imply knowing the actual length of  $\pi$ ).

**Definition 1 (Explaining a  $\mathcal{O}(\tau)$  observation)** *Given a classical planning problem  $P$  and a sequence of partially observed states  $\mathcal{O}(\tau)$ , we say that a plan  $\pi$  explains the observation (denoted  $\pi \mapsto \mathcal{O}(\tau)$ ) iff  $\pi$  is a solution for  $P$  that is consistent with  $\mathcal{O}(\tau)$ . If  $\pi$  is also optimal, we say that  $\pi$  is the best explanation for  $\mathcal{O}(\tau)$ .*

### 2.3 Conformant planning

*Conformant planning* is planning with incomplete information about the initial state, no sensing, and validating that goals are achieved with certainty (despite the uncertainty of the initial state) [?; ?; ?].

Syntactically, conformant planning problems are expressed in compact form through a set of state variables. A *conformant planning problem* can be defined as a tuple  $P_c = \langle F, A, \Upsilon, G \rangle$  where  $F$ ,  $A$  and  $G$  are the set of fluents, actions and goals (as previously defined for *classical planning*). Now  $\Upsilon$  is a set of clauses over literals  $l = f$  or  $l = \neg f$  (for  $f \in F$ ) that define the set of possible initial states.

A solution to a conformant planning problem is an action sequence that maps each possible initial state into a goal state. More precisely, an action sequence  $\pi = \langle a_1, \dots, a_n \rangle$  is a *conformant plan* for  $P_c$  iff, for each possible trajectory  $\tau(\pi, s_0) = \langle s_0, a_1, s_1, \dots, a_n, s_n \rangle$  s.t.  $s_0$  is a valuation of the fluents in  $F$  that satisfies  $\Upsilon$ , then  $\tau(\pi, s_0)$  reaches a final state  $G \subseteq s_n$ .

## 3 Computing the least-commitment action model from state observations

First, this section formalizes the notion of the *least-commitment* action model that is able to *explain* a sequence of partially observed states. Next, the section describes our approach to compute such model via *conformant planning*.

### 3.1 The least-commitment action model

The task of computing the *least-commitment* action model from a sequence of state observations is defined as  $\langle \Phi, \mathcal{O}(\tau) \rangle$ :

- $\Phi = \langle F, A[\cdot] \rangle$  is a *classical planning frame* where  $A[\cdot]$  is a set of actions s.t. the semantics of each  $a \in A[\cdot]$  is unknown; i.e. the corresponding  $\langle \rho, \theta \rangle$  functions are undefined.
- $\mathcal{O}(\tau)$  is a sequence of partial states that results from the partial observation of a trajectory  $\tau(\pi, s_0)$  that is defined within the *classical planning frame*  $\Phi$ .

Before formalizing the solution to this task, i.e. the *least-commitment* action model, we introduce several necessary definitions. Given a *classical planning frame*  $\Phi = \langle F, A[\cdot] \rangle$  and a sequence of partial states  $\mathcal{O}(\tau) = \langle s_0^o, s_1^o, \dots, s_m^o \rangle$ , we can build the classical planning problem  $P_{\mathcal{O}} = \langle F, A[\cdot], s_0^o, s_m^o \rangle$ . We say that an action model  $\mathcal{M}$  is a

definition of the  $\langle \rho, \theta \rangle$  functions of every action in  $A[\cdot]$ . Further we say that a model  $\mathcal{M}$  *explains* a sequence of observations  $\mathcal{O}(\tau)$  iff, when the  $\langle \rho, \theta \rangle$  functions of the actions in  $P_{\mathcal{O}}$  are given by  $\mathcal{M}$ , there exists a solution plan for  $P_{\mathcal{O}}$  that explains  $\mathcal{O}(\tau)$ .

Now, we define a *partially specified action model* inspired by the notion of *incomplete (annotated) model* [?].

**Definition 2 (Partially specified action model)** *Given a classical planning frame  $\Phi = \langle F, A[\cdot] \rangle$  then, a partially specified action model  $M$  is a set of possible models for the actions in  $A[\cdot]$  such that: (1), any model  $\mathcal{M} \in M$  is a definition of the  $\langle \rho, \theta \rangle$  functions of every action in  $A[\cdot]$  and (2), for every  $\mathcal{M} \in M$  the  $\langle \rho, \theta \rangle$  functions are defined in the set of state variables  $F$ . (Note that if  $M$  is a singleton it represents a fully specified action model).*

Now we are ready to define the *least-commitment* action model for an observation  $\mathcal{O}(\tau)$ .

**Definition 3 (The least-commitment action model)** *Given a  $\langle \Phi, \mathcal{O}(\tau) \rangle$  task and the partially specified action model  $M$  that represents the full space of possible action models for the actions in  $A[\cdot] \in \Phi$ , then the least-commitment action model is another partially specified action model that represents the largest subset of models  $M^* \subseteq M$  such that every model  $\mathcal{M} \in M^*$  explains the input observation.*

### 3.2 The space of STRIPS action models

Despite previous definitions are general, this work focuses on the particular kind of action models that are specified as STRIPS action schemata.

A STRIPS *action schema*  $\xi$  is defined by four lists: A list of *parameters*  $pars(\xi)$ , and three list of predicates (namely  $pre(\xi)$ ,  $del(\xi)$  and  $add(\xi)$ ) that shape the kind of fluents that can appear in the *preconditions*, *negative effects* and *positive effects* of the actions induced from that schema. Let be  $\Psi$  the set of *predicates* that shape the propositional state variables  $F$ , and a list of *parameters*  $pars(\xi)$ . The set of elements that can appear in  $pre(\xi)$ ,  $del(\xi)$  and  $add(\xi)$  of the STRIPS action schema  $\xi$  is given by FOL interpretations of  $\Psi$  over the parameters  $pars(\xi)$ . We denote this set of FOL interpretations as  $\mathcal{I}_{\Psi, \xi}$ . For instance, in the *blocksworld* the  $\mathcal{I}_{\Psi, \xi}$  set contain eleven elements for the  $stack(v_1, v_2)$  schemata,  $\mathcal{I}_{\Psi, stack} = \{handempty, holding(v_1), holding(v_2), clear(v_1), clear(v_2), ontable(v_1), ontable(v_2), on(v_1, v_1), on(v_1, v_2), on(v_2, v_1), on(v_2, v_2)\}$ .

Despite any element of  $\mathcal{I}_{\Psi, \xi}$  can *a priori* appear in the  $pre(\xi)$ ,  $del(\xi)$  and  $add(\xi)$  of schema  $\xi$ , the space of possible STRIPS schemata is bounded by constraints of three kinds:

1. *Syntactic constraints.* STRIPS constraints require  $del(\xi) \subseteq pre(\xi)$ ,  $del(\xi) \cap add(\xi) = \emptyset$  and  $pre(\xi) \cap add(\xi) = \emptyset$ . Considering exclusively these syntactic constraints, the size of the space of possible STRIPS schemata is given by  $2^{2 \times |\mathcal{I}_{\Psi, \xi}|}$ .
2. *Domain-specific constraints.* One can introduce domain-specific knowledge to constrain further the space of possible schemata. For instance, in the *blocksworld* one can argue that  $on(v_1, v_1)$  and

```

(:action stack
  :parameters (?v1 ?v2)
  :precondition (and (holding ?v1) (clear ?v2))
  :effect (and (not (holding ?v1)) (not (clear ?v2))
              (clear ?v1) (handempty) (on ?v1 ?v2)))

(pre_holding_v1_stack) (pre_clear_v2_stack)
(eff_holding_v1_stack) (eff_clear_v2_stack)
(eff_clear_v1_stack) (eff_handempty_stack) (eff_on_v1_v2_stack)

```

Figure 1: PDDL encoding of the `stack(?v1, ?v2)` schema and our propositional representation for this same schema.

$\text{on}(v_2, v_2)$  will not appear in the  $\text{pre}(\xi)$ ,  $\text{del}(\xi)$  and  $\text{add}(\xi)$  lists of an action schema  $\xi$  because, in this specific domain, a block cannot be on top of itself. As a rule of thumb, *state invariants* constraining the possible states of a given planning domain belong to this second class of constraints [?].

3. *Observation constraints.* A sequence of state observations  $\mathcal{O}(\tau)$  depicts *semantic knowledge* that constraints further the space of possible action schemata.

### 3.3 Computing the *least-commitment* model via conformant planning

Given a task  $\langle \Phi, \mathcal{O}(\tau) \rangle$ , this section shows that we can build a *conformant planning problem*  $P_c$  whose solution induces the *least-commitment* action model for the input observation  $\mathcal{O}(\tau)$ . In more detail, we build a *conformant planning problem*  $P_c = \langle F_c, A_c, \Upsilon, G_c \rangle$  such that:

- The set of fluents  $F_c$  extends  $F$  with two new sets of fluents:
  - $\{test_j\}_{1 \leq j \leq m}$ , indicating the state observation  $s_j \in \mathcal{O}(\tau)$  where the action model is validated
  - Fluents  $\text{pre}_e\xi$  and  $\text{eff}_e\xi$  (where  $e \in \mathcal{I}_{\Psi, \xi}$ ) implementing a propositional encoding of the *preconditions*, *negative*, and *positive* effects of an action schema  $\xi$ . **Our encoding exploits the syntactic constraint of STRIPS** so, if  $\text{pre}_e\xi$  and  $\text{eff}_e\xi$  holds it means that  $e \in \mathcal{I}_{\Psi, \xi}$  is a negative effect in  $\xi$  while if  $\text{pre}_e\xi$  does not hold but  $\text{eff}_e\xi$  holds, it means that  $e \in \mathcal{I}_{\Psi, \xi}$  is a positive effect in  $\xi$ . Figure 1 shows the PDDL encoding of the `stack(?v1, ?v2)` schema and our propositional representation for this same schema.
- The set of actions  $A_c$  contains now actions of three different kinds:
  - Actions for *committing*  $\text{pre}_e\xi$  fluents to a positive/negative value (similar actions are also defined for *committing*  $\text{eff}_e\xi$  fluents to a positive/negative value).
 
$$\begin{aligned} \text{pre}(\text{commit}\top\text{-pre}_e\xi) &= \{mode_{commit}\}, \\ \text{cond}(\text{commit}\top\text{-pre}_e\xi) &= \{\text{pre}_e\xi\} \triangleright \{\text{pre}_e\xi\}, \\ &\quad \{\neg\text{pre}_e\xi\} \triangleright \{\text{pre}_e\xi\}. \end{aligned}$$
  - Actions for *committing*  $\text{pre}_e\xi$  fluents to a positive/negative value (similar actions are also defined for *committing*  $\text{eff}_e\xi$  fluents to a positive/negative value).
 
$$\begin{aligned} \text{pre}(\text{commit}\bot\text{-pre}_e\xi) &= \{mode_{commit}\}, \\ \text{cond}(\text{commit}\bot\text{-pre}_e\xi) &= \{\text{pre}_e\xi\} \triangleright \{\neg\text{pre}_e\xi\}, \\ &\quad \{\neg\text{pre}_e\xi\} \triangleright \{\neg\text{pre}_e\xi\}. \end{aligned}$$

- Actions for *validating* that committed models explain the  $s_j$  observed states,  $0 \leq j < m$ .

$$\begin{aligned} \text{pre}(\text{validate}_j) &= s_j \cup \{test_{j-1}\}, \\ \text{cond}(\text{validate}_j) &= \{\emptyset\} \triangleright \{\neg test_{j-1}, test_j, \\ &\quad \{mode_{commit}\} \triangleright \{\neg mode_{commit}, mode_{val}\}. \end{aligned}$$

- *Editable* actions whose semantics is given by the value of  $\text{pre}_e\xi$ ,  $\text{eff}_e\xi$  fluents at the current state. Figure 2 shows the PDDL encoding of an *editable* `stack(?v1, ?v2)` schema. Note that this editable schema when the set of fluents  $(\text{pre\_holding\_v1\_stack})$   $(\text{pre\_clear\_v2\_stack})$   $(\text{eff\_holding\_v1\_stack})$   $(\text{eff\_clear\_v2\_stack})$   $(\text{eff\_clear\_v1\_stack})$   $(\text{eff\_handempty\_stack})$   $(\text{eff\_on\_v1\_v2\_stack})$  hold at the current state, then it behaves exactly as the original PDDL schema defined in Figure 1. Formally, given an operator schema  $\xi \in \mathcal{M}$  its *editable* version is:

$$\begin{aligned} \text{pre}(\text{editable}_\xi) &= \{\text{pre}_e\xi \implies e\}_{e \in \mathcal{I}_{\Psi, \xi}} \\ \text{cond}(\text{editable}_\xi) &= \{\text{pre}_e\xi, \text{eff}_e\xi\} \triangleright \{\neg e\}_{e \in \mathcal{I}_{\Psi, \xi}}, \\ &\quad \{\neg\text{pre}_e\xi, \text{eff}_e\xi\} \triangleright \{e\}_{e \in \mathcal{I}_{\Psi, \xi}}. \end{aligned}$$

- The clauses in  $\Upsilon$  comprises:

1. The *unit clauses* given by the fluents that hold in the initial state  $I = s_0$  and  $mode_{commit}$  set to true.
2. The clauses representing that the actual value of fluents  $\text{pre}_e\xi$ ,  $\text{eff}_e\xi$  is unknown. In other words, that any model from the STRIPS space of models (following the previously mentioned syntactic constraints) can initially be part of the *least-commitment* action model. Formally, for every  $\xi$  and  $e \in \mathcal{I}_{\Psi, \xi}$ , then  $\Upsilon$  includes these two clauses:

$$\begin{aligned} &-\text{pre}_e\xi \vee \neg\text{pre}_e\xi. \\ &-\text{eff}_e\xi \vee \neg\text{eff}_e\xi. \end{aligned}$$

One can also add here clauses that encode *domain-specific constraints* (as mentioned in the previous section) to make the conformant planning problem easier for a specific domain.

- The new goals are  $G_c = s_m \cup \{test_m\}$ .

### 3.4 Optimization of the compilation

In fact, we do not need to compute the *least-commitment* action model (i.e. solve the  $P_c$  conformant planning problem) from scratch for all the  $\langle s_0^o, s_1^o, \dots, s_m^o \rangle$  observations. A more efficient approach is to compute the *least-commitment* action model from scratch only for the first observation, that is for  $\langle s_0^o, s_1^o \rangle$ . Then, the obtained *least-commitment* action model can be used as input for improving the computation of the two first observations  $\langle s_0^o, s_1^o, s_2^o \rangle$ . In other words, including the commits found in the previous step as new *unit clauses* in set of the *initial clauses*  $\Upsilon$ . This process is repeated until the *least-commitment* action model is computed for the full sequence of state observations  $\langle s_0^o, s_1^o, \dots, s_m^o \rangle$ .

```

(:action stack
:parameters (?o1 - object ?o2 - object)
:precondition
  (and (or (not (pre_on_v1_v1_stack)) (on ?o1 ?o1))
        (or (not (pre_on_v1_v2_stack)) (on ?o1 ?o2))
        (or (not (pre_on_v2_v1_stack)) (on ?o2 ?o1))
        (or (not (pre_on_v2_v2_stack)) (on ?o2 ?o2))
        (or (not (pre_ontable_v1_stack)) (ontable ?o1))
        (or (not (pre_ontable_v2_stack)) (ontable ?o2))
        (or (not (pre_clear_v1_stack)) (clear ?o1))
        (or (not (pre_clear_v2_stack)) (clear ?o2))
        (or (not (pre_holding_v1_stack)) (holding ?o1))
        (or (not (pre_holding_v2_stack)) (holding ?o2))
        (or (not (pre_hanempty_stack)) (handempty))))
:effect (and
  (when (and (pre_on_v1_v1_stack) (eff_on_v1_v1_stack)) (not (on ?o1 ?o1)))
  (when (and (pre_on_v1_v2_stack) (eff_on_v1_v2_stack)) (not (on ?o1 ?o2)))
  (when (and (pre_on_v2_v1_stack) (eff_on_v2_v1_stack)) (not (on ?o2 ?o1)))
  (when (and (pre_on_v2_v2_stack) (eff_on_v2_v2_stack)) (not (on ?o2 ?o2)))
  (when (and (pre_ontable_v1_stack) (eff_ontable_v1_stack)) (not (ontable ?o1)))
  (when (and (pre_ontable_v2_stack) (eff_ontable_v2_stack)) (not (ontable ?o2)))
  (when (and (pre_clear_v1_stack) (eff_clear_v1_stack)) (not (clear ?o1)))
  (when (and (pre_clear_v2_stack) (eff_clear_v2_stack)) (not (clear ?o2)))
  (when (and (pre_holding_v1_stack) (eff_holding_v1_stack)) (not (holding ?o1)))
  (when (and (pre_holding_v2_stack) (eff_holding_v2_stack)) (not (holding ?o2)))
  (when (and (pre_hanempty_stack) (eff_hanempty_stack)) (not (handempty)))
  (when (and (not (pre_on_v1_v1_stack)) (eff_on_v1_v1_stack)) (on ?o1 ?o1))
  (when (and (not (pre_on_v1_v2_stack)) (eff_on_v1_v2_stack)) (on ?o1 ?o2))
  (when (and (not (pre_on_v2_v1_stack)) (eff_on_v2_v1_stack)) (on ?o2 ?o1))
  (when (and (not (pre_on_v2_v2_stack)) (eff_on_v2_v2_stack)) (on ?o2 ?o2))
  (when (and (not (pre_ontable_v1_stack)) (eff_ontable_v1_stack)) (ontable ?o1))
  (when (and (not (pre_ontable_v2_stack)) (eff_ontable_v2_stack)) (ontable ?o2))
  (when (and (not (pre_clear_v1_stack)) (eff_clear_v1_stack)) (clear ?o1))
  (when (and (not (pre_clear_v2_stack)) (eff_clear_v2_stack)) (clear ?o2))
  (when (and (not (pre_holding_v1_stack)) (eff_holding_v1_stack)) (holding ?o1))
  (when (and (not (pre_holding_v2_stack)) (eff_holding_v2_stack)) (holding ?o2))
  (when (and (not (pre_hanempty_stack)) (eff_hanempty_stack)) (handempty))))

```

Figure 2: PDDL encoding of the editable version of the stack(?v1, ?v2) schema.

### 3.5 Compilation properties

## 4 Evaluation

## 5 Conclusions

Related work [?].

## References

- [Bonet and Geffner, 2000] Blai Bonet and Hector Geffner. Planning with incomplete information as heuristic search in belief space. In *Proceedings of the Fifth International Conference on Artificial Intelligence Planning Systems*, pages 52–61. AAAI Press, 2000.
- [Fox and Long, 1998] Maria Fox and Derek Long. The automatic inference of state invariants in tim. *Journal of Artificial Intelligence Research*, 9:367–421, 1998.
- [Goldman and Boddy, 1996] Robert P Goldman and Mark S Boddy. Expressive planning and explicit knowledge. In *AIPS*, volume 96, pages 110–117, 1996.
- [Ramírez and Geffner, 2009] Miquel Ramírez and Hector Geffner. Plan recognition as planning. In *International Joint conference on Artificial Intelligence, (IJCAI-09)*, pages 1778–1783. AAAI Press, 2009.
- [Smith and Weld, 1998] David E Smith and Daniel S Weld. Conformant graphplan. In *AAAI/IAAI*, pages 889–896, 1998.
- [Sreedharan *et al.*, 2018] Sarath Sreedharan, Tathagata Chakraborti, and Subbarao Kambhampati. Handling model uncertainty and multiplicity in explanations via

model reconciliation. In *International Conference on Automated Planning and Scheduling, (ICAPS-18)*, pages 518–526, 2018.

[Stern and Juba, 2017] Roni Stern and Brendan Juba. Efficient, safe, and probably approximately complete learning of action models. In *International Joint Conference on Artificial Intelligence, (IJCAI-17)*, pages 4405–4411, 2017.