

# Learning STRIPS action models from *state-invariants*

Diego Aineto<sup>1</sup>, Sergio Jiménez<sup>1</sup>, Eva Onaindia<sup>1</sup>

<sup>1</sup>Departamento de Sistemas Informáticos y Computación. Universitat Politècnica de València. Valencia, Spain

{dieaigar,serjice,onaindia}@dsic.upv.es

## Abstract

## 1 Introduction

The specification of planning action models is a complex process that limits, too often, the potential of model-based planning systems [Kambhampati, 2007]. The *machine learning* of action models can relieve this knowledge acquisition bottleneck of AI planning and nowadays there exists a wide range of effective approaches for learning planning action models [Arora *et al.*, 2018]. Most of current approaches for learning planning action models are *inductive*, meaning that they compute an action model that maximizes some notion of *statistical consistency* over a set of observations of plan executions so their success depends on the amount and quality of the input observations.

This paper proposes the exploitation of *deductive* knowledge to learn planning action models when input observations are minimal (just an *initial state* and the *goals*). We show that in such unfavourable scenario *state-invariant* (i.e. logic formulae that specify constraints about the possible states of a given domain) help to learn better action models. Our approach is built on top of the *classical planning* compilation for the learning of STRIPS action models [Aineto *et al.*, 2018]. This compilation approach is flexible to different kinds of input knowledge (e.g., partially/fully observations of actions of plan executions as well as partially/fully observed intermediate states) and guarantees to output an action model that is *consistent* with the input knowledge. In this paper we show how to push the flexibility of the compilation approach for learning of STRIPS action models and boost its performance with *state-invariant* when few observations are given as input.

## 2 Background

This section formalizes the *classical planning model* we follow in this work and the kind of *knowledge* that can be given as input to the task of learning STRIPS action models.

### 2.1 Classical planning with conditional effects

Let  $F$  be the set of propositional state variables (*fluents*) describing a state. A *literal*  $l$  is a valuation of a fluent  $f \in F$ ;

i.e. either  $l = f$  or  $l = \neg f$ . A set of literals  $L$  represents a partial assignment of values to fluents (without loss of generality, we will assume that  $L$  does not contain conflicting values). Given  $L$ , let  $\neg L = \{\neg l : l \in L\}$  be its complement. We use  $\mathcal{L}(F)$  to denote the set of all literal sets on  $F$ ; i.e. all partial assignments of values to fluents. A *state*  $s$  is a full assignment of values to fluents;  $|s| = |F|$ .

A *classical planning action*  $a \in A$  has: a precondition  $\text{pre}(a) \in \mathcal{L}(F)$ , a set of effects  $\text{eff}(a) \in \mathcal{L}(F)$ , and a positive action cost  $\text{cost}(a)$ . The semantics of actions  $a \in A$  is specified with two functions:  $\rho(s, a)$  denotes whether action  $a$  is *applicable* in a state  $s$  and  $\theta(s, a)$  denotes the *successor state* that results of applying action  $a$  in a state  $s$ . Then,  $\rho(s, a)$  holds iff  $\text{pre}(a) \subseteq s$ , i.e. if its precondition holds in  $s$ . The result of executing an applicable action  $a \in A$  in a state  $s$  is a new state  $\theta(s, a) = (s \setminus \neg \text{eff}(a)) \cup \text{eff}(a)$ . Subtracting the complement of  $\text{eff}(a)$  from  $s$  ensures that  $\theta(s, a)$  remains a well-defined state. The subset of action effects that assign a positive value to a state fluent is called *positive effects* and denoted by  $\text{eff}^+(a) \in \text{eff}(a)$  while  $\text{eff}^-(a) \in \text{eff}(a)$  denotes the *negative effects* of an action  $a \in A$ .

A *classical planning problem* is a tuple  $P = \langle F, A, I, G \rangle$ , where  $I$  is the initial state and  $G \in \mathcal{L}(F)$  is the set of goal conditions over the state variables. A *plan*  $\pi$  is an action sequence  $\pi = \langle a_1, \dots, a_n \rangle$ , with  $|\pi| = n$  denoting its *plan length* and  $\text{cost}(\pi) = \sum_{a \in \pi} \text{cost}(a)$  its *plan cost*. The execution of  $\pi$  on the initial state of  $P$  induces a *trajectory*  $\tau(\pi, P) = \langle s_0, a_1, s_1, \dots, a_n, s_n \rangle$  such that  $s_0 = I$  and, for each  $1 \leq i \leq n$ , it holds  $\rho(s_{i-1}, a_i)$  and  $s_i = \theta(s_{i-1}, a_i)$ . A plan  $\pi$  solves  $P$  iff the induced *trajectory*  $\tau(\pi, P)$  reaches a final state  $G \subseteq s_n$ , where all goal conditions are met. A solution plan is *optimal* iff its cost is minimal.

We also define *actions with conditional effects* because they are useful to compactly formulate our approach for *goal recognition with unknown domain models*. An action  $a_c \in A$  with conditional effects is a set of preconditions  $\text{pre}(a_c) \in \mathcal{L}(F)$  and a set of *conditional effects*  $\text{cond}(a_c)$ . Each conditional effect  $C \triangleright E \in \text{cond}(a_c)$  is composed of two sets of literals:  $C \in \mathcal{L}(F)$ , the *condition*, and  $E \in \mathcal{L}(F)$ , the *effect*. An action  $a_c$  is applicable in a state  $s$  if  $\rho(s, a_c)$  is true, and the result of applying action  $a_c$  in state  $s$  is  $\theta(s, a_c) = \{s \setminus \neg \text{eff}_c(s, a) \cup \text{eff}_c(s, a)\}$  where  $\text{eff}_c(s, a)$  are the *triggered effects* resulting from the action application

(conditional effects whose conditions hold in  $s$ ):

$$\text{eff}_c(s, a) = \bigcup_{C \supset E \in \text{cond}(a_c), C \subseteq s} E,$$

## 2.2 The observation model

Given a planning problem  $P = \langle F, A, I, G \rangle$ , a plan  $\pi$  and a trajectory  $\tau(\pi, P)$ , we define the *observation of the trajectory* as an interleaved combination of actions and states that represents the observation from the execution of  $\pi$  in  $P$ . Formally,  $\mathcal{O}(\tau) = \langle s_0^o, a_1^o, s_1^o, \dots, a_l^o, s_m^o \rangle$ ,  $s_0^o = I$ , and:

- The **observed actions** are consistent with  $\pi$ , which means that  $\langle a_1^o, \dots, a_l^o \rangle$  is a sub-sequence of  $\pi$ . The number of observed actions,  $l$ , ranges from 0 (fully unobserved action sequence) to  $|\pi|$  (fully observed action sequence).
- The **observed states**  $\langle s_0^o, s_1^o, \dots, s_m^o \rangle$  is a sequence of possibly *partially observable states*, except for the initial state  $s_0^o$ , which is fully observed. A partially observable state  $s_i^o$  is one in which  $|s_i^o| < |F|$ ; i.e., a state in which at least a fluent of  $F$  is not observable. Note that this definition also comprises the case  $|s_i^o| = 0$ , when the state is fully unobservable. Whatever the sequence of observed states of  $\mathcal{O}(\tau)$  is, it must be consistent with the sequence of states of  $\tau(\pi, P)$ , meaning that  $\forall i, s_i^o \subseteq s_i$ . The number of observed states,  $m$ , range from 1 (the initial state, at least), to  $|\pi| + 1$ , and each *observed* states comprises  $[1, |F|]$  fluents (the observation can still miss intermediate states that are *unobserved*).

We assume a bijective monotone mapping between actions/states of trajectories and observations [Ramírez and Geffner, 2009], thus also granting the inverse consistency relationship (the trajectory is a superset of the observation). Therefore, transiting between two consecutive observed states in  $\mathcal{O}(\tau)$  may require the execution of more than a single action ( $\theta(s_i^o, \langle a_1, \dots, a_k \rangle) = s_{i+1}^o$ , where  $k \geq 1$  is unknown but finite. In other words, having an input observation  $\mathcal{O}(\tau)$  does not imply knowing the actual length of  $\pi$ ).

## 2.3 State-invariants

The notion of *state-constraint* is very general and has been used in different areas of AI and for different purposes. If we restrict ourselves to planning, *state-constraints* are abstractions for compactly specifying sets of states. For instance, *state-constraints* in planning allow to specify the set of states where a given action is applicable, the set of states where a given *derived predicate* holds or the set of states that are considered goal states.

*State invariants* is a kind of state-constraints useful for computing more compact state representations [Helmert, 2009] or making *satisfiability planning* and *backward search* more efficient [Rintanen, 2014; Alcázar and Torralba, 2015]. Given a classical planning problem  $P = \langle F, A, I, G \rangle$ , a *state invariant* is a formula  $\phi$  that holds at the initial state of a given classical planning problem,  $I \models \phi$ , and at every state  $s$ , built from  $F$ , that is reachable from  $I$  by applying actions in  $A$ .

The formula  $\phi_{I,A}^*$  represents the *strongest invariant* and exactly characterizes the set of all states reachable from  $I$  with

the actions in  $A$ . For instance Figure 1 shows five clauses that define the *strongest invariant* for the *blocksworld* planning domain [Slaney and Thiébaux, 2001]. There are infinitely many strongest invariants, but they are all logically equivalent, and computing the strongest invariant is PSPACE-hard (as hard as testing plan existence [Bylander, 1994]).

$$\begin{aligned} \forall x_1, x_2 \text{ ontable}(x_1) &\leftrightarrow \neg \text{on}(x_1, x_2). \\ \forall x_1, x_2 \text{ clear}(x_1) &\leftrightarrow \neg \text{on}(x_2, x_1). \\ \forall x_1, x_2, x_3 \neg \text{on}(x_1, x_2) \vee \neg \text{on}(x_1, x_3) &\text{ such that } x_2 \neq x_3. \\ \forall x_1, x_2, x_3 \neg \text{on}(x_2, x_1) \vee \neg \text{on}(x_3, x_1) &\text{ such that } x_2 \neq x_3. \\ \forall x_1, \dots, x_n \neg(\text{on}(x_1, x_2) \wedge \text{on}(x_2, x_3) \wedge \dots \wedge \text{on}(x_{n-1}, x_n) \wedge & \\ \text{on}(x_n, x_1)) &. \end{aligned}$$

Figure 1: *Strongest invariant* for the *blocksworld* domain.

A *mutex* (mutually exclusive) is a state invariant that takes the form of a binary clause and indicates a pair of different properties that cannot be simultaneously true [Kautz and Selman, 1999]. For instance in a three-block *blocksworld*,  $\phi_1 = \neg \text{on}(\text{block}_A, \text{block}_B) \vee \neg \text{on}(\text{block}_A, \text{block}_C)$  is a mutex because *block<sub>A</sub>* can only be on top of a single block.

A *domain invariant* is an instance-independent invariant, i.e. holds for any possible initial state and set of objects. Therefore, if a given state  $s$  holds  $s \not\models \phi$  such that  $\phi$  is a *domain invariant*, it means that  $s$  is not a valid state. Domain invariants are often compactly defined as *lifted invariants* (also called schematic invariants) [Rintanen and others, 2017]. For instance,  $\phi_2 = \forall x : (\neg \text{handempty} \vee \neg \text{holding}(x))$ , is a *domain mutex* for the *blocksworld* because the robot hand is never empty and holding a block at the same time.

## 3 Learning action models from state-invariants

We define the task of learning action models as a tuple  $\Lambda = \langle M, \Phi, \mathcal{O}(\tau) \rangle$ , where:

- $M$  is the *space of possible action model*. This set is the *full* space of action models, when learning from scratch, or *partially specified*, when some fragments of the action models are known a priori. A set of action models can be defined *explicitly*, enumerating all the models that belong to the set or *implicitly*, enumerating all the constraints that must satisfy any model that belongs to the set. A *partially specified* STRIPS action model is then a formalism for the *implicit* representation of a set of STRIPS schemes [Sreedharan *et al.*, 2018].
- $\Phi$ , a set of *state-invariants* that constrain the set of possible states in the given domain.
- $\mathcal{O}(\tau)$  is an observation of a trajectory  $\tau(\pi, P)$  produced by the execution of an unknown plan  $\pi$  that solves  $P_{\mathcal{O}}$  (i.e. the classical planning problem  $P_{\mathcal{O}} = \langle F, A[\cdot], s_0^o, s_m^o \rangle$  where  $A[\cdot]$  is a set of actions s.t., the semantics of each action  $a \in A[\cdot]$  is unknown since functions  $\rho$  and/or  $\theta$  of  $a$  are undefined).

A *solution* to a  $\Lambda = \langle M, \Phi, \mathcal{O}(\tau) \rangle$  learning task is a model  $M' \in M$  that is *consistent* with the given input knowledge  $\Phi$  and  $\mathcal{O}(\tau)$ . This means that let there exists a plan  $\pi$  such that

the observation  $\mathcal{O}(\tau)$  is consistent with  $\pi$  when the dynamics of  $A[\cdot]$  is given by  $\mathcal{M}' \in M$  and such that any state produced in the trajectory  $\tau(\pi, P_{\mathcal{O}})$  satisfies the *state-invariants*  $\Phi$ .

Next we show that the set  $M$  of possible action models can be encoded as a set of propositional variables and a set of constraints over those variables. Then, we show how to exploit this encoding to solve the  $\Lambda = \langle M, \Phi, \mathcal{O}(\tau) \rangle$  learning task with an off-the-shelf classical planner.

### 3.1 A propositional encoding for the space of STRIPS action models

A STRIPS *action schema*  $\xi$  is defined by four lists: A list of *parameters*  $\text{pars}(\xi)$ , and three list of predicates (namely  $\text{pre}(\xi)$ ,  $\text{del}(\xi)$  and  $\text{add}(\xi)$ ) that shape the kind of fluents that can appear in the *preconditions*, *negative effects* and *positive effects* of the actions induced from that schema. Let be  $\Psi$  the set of *predicates* that shape the propositional state variables  $F$ , and a list of *parameters*  $\text{pars}(\xi)$ . The set of elements that can appear in  $\text{pre}(\xi)$ ,  $\text{del}(\xi)$  and  $\text{add}(\xi)$  of the STRIPS action schema  $\xi$  is given by FOL interpretations of  $\Psi$  over the parameters  $\text{pars}(\xi)$  and is denoted as  $\mathcal{I}_{\Psi, \xi}$ .

For instance in a four-operator *blocksworld* [Slaney and Thiébaux, 2001], the  $\mathcal{I}_{\Psi, \xi}$  set contains only five elements for the `pickup( $v_1$ )` schemata,  $\mathcal{I}_{\Psi, \text{pickup}} = \{\text{handempty}, \text{holding}(v_1), \text{clear}(v_1), \text{ontable}(v_1), \text{on}(v_1, v_1)\}$  while it contains eleven elements for the `stack( $v_1, v_2$ )` schemata,  $\mathcal{I}_{\Psi, \text{stack}} = \{\text{handempty}, \text{holding}(v_1), \text{holding}(v_2), \text{clear}(v_1), \text{clear}(v_2), \text{ontable}(v_1), \text{ontable}(v_2), \text{on}(v_1, v_1), \text{on}(v_1, v_2), \text{on}(v_2, v_1), \text{on}(v_2, v_2)\}$ .

Despite any element of  $\mathcal{I}_{\Psi, \xi}$  can *a priori* appear in the  $\text{pre}(\xi)$ ,  $\text{del}(\xi)$  and  $\text{add}(\xi)$  of schema  $\xi$ , the actual space of possible STRIPS schemata is bounded by constraints of three kinds:

1. **Syntactic constraints.** STRIPS constraints require  $\text{del}(\xi) \subseteq \text{pre}(\xi)$ ,  $\text{del}(\xi) \cap \text{add}(\xi) = \emptyset$  and  $\text{pre}(\xi) \cap \text{add}(\xi) = \emptyset$ . Considering exclusively these syntactic constraints, the size of the space of possible STRIPS schemata is given by  $2^{2 \times |\mathcal{I}_{\Psi, \xi}|}$ . *Typing constraints* are also of this kind [McDermott *et al.*, 1998].
2. **Domain-specific constraints.** One can introduce domain-specific knowledge to constrain further the space of possible schemata. For instance, in the *blocksworld* one can argue that  $\text{on}(v_1, v_1)$  and  $\text{on}(v_2, v_2)$  will not appear in the  $\text{pre}(\xi)$ ,  $\text{del}(\xi)$  and  $\text{add}(\xi)$  lists of an action schema  $\xi$  because, in this specific domain, a block cannot be on top of itself. *State invariants* are also constraints of this kind.
3. **Observation constraints.** An observation  $\mathcal{O}(\tau)$  depicts *semantic knowledge* that constraints further the space of possible action schemata.

In this work we introduce a propositional encoding of the *preconditions*, *negative*, and *positive* effects of a STRIPS action schema  $\xi$  using only fluents of two kinds  $\text{pre}_e \xi$  and  $\text{eff}_e \xi$  (where  $e \in \mathcal{I}_{\Psi, \xi}$ ). This encoding exploits the syntactic constraints of STRIPS so it is more compact than the one previously proposed by Aineto *et al.* 2018 for learning

```
(:action stack
:parameters (?v1 ?v2)
:precondition (and (holding ?v1) (clear ?v2))
:effect (and (not (holding ?v1)) (not (clear ?v2))
             (clear ?v1) (handempty) (on ?v1 ?v2)))

(pre_holding_v1_stack) (pre_clear_v2_stack)
(eff_holding_v1_stack) (eff_clear_v2_stack)
(eff_clear_v1_stack) (eff_handempty_stack) (eff_on_v1_v2_stack)
```

Figure 2: PDDL encoding of the `stack(?v1, ?v2)` schema and our propositional representation for this same schema.

classical planning action models. In more detail, if  $\text{pre}_e \xi$  holds it means that  $e \in \mathcal{I}_{\Psi, \xi}$  is a *precondition* in  $\xi$ . If  $\text{pre}_e \xi$  and  $\text{eff}_e \xi$  holds it means that  $e \in \mathcal{I}_{\Psi, \xi}$  is a *negative effect* in  $\xi$  while if  $\text{pre}_e \xi$  does not hold but  $\text{eff}_e \xi$  holds, it means that  $e \in \mathcal{I}_{\Psi, \xi}$  is a *positive effect* in  $\xi$ . Figure 2 shows the PDDL encoding of the `stack(?v1, ?v2)` schema and our propositional representation for this same schema with  $\text{pre}_e \text{stack}$  and  $\text{eff}_e \text{stack}$  fluents ( $e \in \mathcal{I}_{\Psi, \text{stack}}$ ).

### 3.2 Learning STRIPS action models with classical planning

Given a  $\Lambda = \langle M, \Phi, \mathcal{O}(\tau) \rangle$  where  $\Phi$  is a set of *domain mutex*  $\phi \in \Phi$ , we create a classical planning problem  $P' = \langle F', A', I, G \rangle$  such that:

- $F'$  extends  $F$  with a fluent *mode<sub>insert</sub>*, to indicate whether action models are being programmed, and the fluents for the propositional encoding of the corresponding space of STRIPS action models. This is a set of fluents of the type  $\{\text{pre}_e \xi, \text{eff}_e \xi\}_{\forall e \in \mathcal{I}_{\Psi, \xi}}$  such that  $e \in \mathcal{I}_{\Psi, \xi}$  is a single element from the set of FOL interpretations of predicates  $\Psi$  over the corresponding action parameters  $\text{pars}(\xi)$ .
- $A'$  replaces the actions in  $A$  with two types of actions.
  1. Actions for *inserting* a *precondition*, *positive* effect or *negative* effect in  $\xi$  following the syntactic constraints of STRIPS models. In the particular case that  $M$  is a *partially specified model* then only the actions for inserting a possible *precondition* or *effect* are necessary.
    - Actions which support the addition of a *precondition*  $p \in \Psi_{\xi}$  to the action model  $\xi$ . A precondition  $p$  is inserted in  $\xi$  when neither  $\text{pre}_p$ ,  $\text{eff}_p$  exist in  $\xi$ .

$$\begin{aligned} \text{pre}(\text{insertPre}_{p, \xi}) &= \{\neg \text{pre}_p(\xi), \neg \text{eff}_p(\xi), \text{mode}_{\text{insert}}\}, \\ \text{cond}(\text{insertPre}_{p, \xi}) &= \{\emptyset\} \triangleright \{\text{pre}_p(\xi)\}. \end{aligned}$$

- Actions which support the addition of a *negative* or *positive* effect  $p \in \Psi_{\xi}$  to the action model  $\xi$ .

$$\begin{aligned} \text{pre}(\text{insertEff}_{p, \xi}) &= \{\neg \text{eff}_p(\xi), \text{mode}_{\text{insert}}\}, \\ \text{cond}(\text{insertEff}_{p, \xi}) &= \{\emptyset\} \triangleright \{\text{eff}_p(\xi)\}. \end{aligned}$$

2. Actions for *applying* an action model  $\xi$  built by the *insert* actions and bounded to objects  $\omega \subseteq$

$\Omega^{|pars(\xi)|}$  (where  $\Omega$  is the set of *objects* used to induce the fluents  $F$  by assigning objects in  $\Omega$  to the  $\Psi$  predicates and  $\Omega^k$  is the  $k$ -th Cartesian power of  $\Omega$ ). The action parameters,  $pars(\xi)$ , are bound to the objects in  $\omega$  that appear in the same position.

$$\begin{aligned} \text{pre}(\text{apply}_{\xi, \omega}) &= \{pre_p(\xi) \implies p(\omega)\}_{\forall p \in \Psi_\xi}, \\ \text{cond}(\text{apply}_{\xi, \omega}) &= \{pre_p(\xi) \wedge eff_p(\xi)\} \triangleright \{\neg p(\omega)\}_{\forall p \in \Psi_\xi}, \\ &\quad \{\neg pre_p(\xi) \wedge eff_p(\xi)\} \triangleright \{p(\omega)\}_{\forall p \in \Psi_\xi}, \\ &\quad \{\emptyset\} \triangleright \{\neg mode_{insert}\}. \end{aligned}$$

*domain mutex* are useful to reduce the amount of applicable actions for programming a precondition or an effect for a given action schema. For example given the *domain mutex*  $\phi = (\neg f_1 \vee \neg f_2)$  such that  $f_1 \in F_v(\xi)$  and  $f_2 \in F_v(\xi)$ , we can redefine the corresponding programming actions for **removing** the *precondition*  $f_1 \in F_v(\xi)$  from the action schema  $\xi \in \mathcal{M}$  as:

## 4 Evaluation

## 5 Conclusions

## References

- [Aineto *et al.*, 2018] Diego Aineto, Sergio Jiménez, and Eva Onaindia. Learning STRIPS action models with classical planning. In *International Conference on Automated Planning and Scheduling, (ICAPS-18)*, pages 399–407. AAAI Press, 2018.
- [Alcázar and Torralba, 2015] Vidal Alcázar and Alvaro Torralba. A reminder about the importance of computing and exploiting invariants in planning. In *ICAPS*, pages 2–6. AAAI Press, 2015.
- [Arora *et al.*, 2018] Ankuj Arora, Humbert Fiorino, Damien Pellier, Marc Métivier, and Sylvie Pesty. A review of learning planning action models. *The Knowledge Engineering Review*, 2018.
- [Bylander, 1994] Tom Bylander. The computational complexity of propositional strips planning. *Artificial Intelligence*, 69(1-2):165–204, 1994.
- [Helmert, 2009] Malte Helmert. Concise finite-domain representations for pddl planning tasks. *Artificial Intelligence*, 173(5-6):503–535, 2009.
- [Kambhampati, 2007] Subbarao Kambhampati. Model-lite planning for the web age masses: The challenges of planning with incomplete and evolving domain models. In *National Conference on Artificial Intelligence, (AAAI-07)*, 2007.
- [Kautz and Selman, 1999] Henry Kautz and Bart Selman. Unifying SAT-based and graph-based planning. In *IJCAI*, volume 99, pages 318–325, 1999.
- [McDermott *et al.*, 1998] Drew McDermott, Malik Ghallab, Adele Howe, Craig Knoblock, Ashwin Ram, Manuela Veloso, Daniel Weld, and David Wilkins. PDDL – The Planning Domain Definition Language, 1998.
- [Ramírez and Geffner, 2009] Miquel Ramírez and Hector Geffner. Plan recognition as planning. In *International Joint conference on Artificial Intelligence, (IJCAI-09)*, pages 1778–1783. AAAI Press, 2009.
- [Rintanen and others, 2017] Jussi Rintanen et al. Schematic invariants by reduction to ground invariants. In *AAAI*, pages 3644–3650, 2017.
- [Rintanen, 2014] Jussi Rintanen. Madagascar: Scalable planning with SAT. In *International Planning Competition, (IPC-2014)*, 2014.
- [Slaney and Thiébaux, 2001] John Slaney and Sylvie Thiébaux. Blocks world revisited. *Artificial Intelligence*, 125(1-2):119–153, 2001.
- [Sreedharan *et al.*, 2018] Sarath Sreedharan, Tathagata Chakraborti, and Subbarao Kambhampati. Handling model uncertainty and multiplicity in explanations via model reconciliation. In *International Conference on Automated Planning and Scheduling, (ICAPS-18)*, pages 518–526, 2018.