

# Learning STRIPS action models from *state-invariants*

Diego Aineto<sup>1</sup>, Sergio Jiménez<sup>1</sup>, Eva Onaindia<sup>1</sup>

<sup>1</sup>Departamento de Sistemas Informáticos y Computación. Universitat Politècnica de València. Valencia, Spain

{dieaigar,serjice,onaindia}@dsic.upv.es, miquel.ramirez@unimelb.edu.au

## Abstract

## 1 Introduction

*Classical planing* is an interesting approach for learning STRIPS action models since it is flexible to different kinds of input knowledge (e.g., partially/fully observations of actions of plan executions as well as partially/fully observed intermediate states) [Aineto *et al.*, 2018]. This paper shows that this flexibility of the compilation approach goes one step further since it also allows learning from *state-invariants*, logic formulae that constrain the set of possible states of a given domain. The experimental results show that *state-invariants* boost the performance of the *classical planing* compilation for learning STRIPS action models.

## 2 Background

This section formalizes the *classical planning model* we follow in this work and the kind of *knowledge* that can be given as input to the task of learning STRIPS action models.

### 2.1 Classical planning with conditional effects

Let  $F$  be the set of propositional state variables (*fluents*) describing a state. A *literal*  $l$  is a valuation of a fluent  $f \in F$ ; i.e. either  $l = f$  or  $l = \neg f$ . A set of literals  $L$  represents a partial assignment of values to fluents (without loss of generality, we will assume that  $L$  does not contain conflicting values). Given  $L$ , let  $\neg L = \{\neg l : l \in L\}$  be its complement. We use  $\mathcal{L}(F)$  to denote the set of all literal sets on  $F$ ; i.e. all partial assignments of values to fluents. A *state*  $s$  is a full assignment of values to fluents;  $|s| = |F|$ .

A *classical planning action*  $a \in A$  has: a precondition  $\text{pre}(a) \in \mathcal{L}(F)$ , a set of effects  $\text{eff}(a) \in \mathcal{L}(F)$ , and a positive action cost  $\text{cost}(a)$ . The semantics of actions  $a \in A$  is specified with two functions:  $\rho(s, a)$  denotes whether action  $a$  is *applicable* in a state  $s$  and  $\theta(s, a)$  denotes the *successor state* that results of applying action  $a$  in a state  $s$ . Then,  $\rho(s, a)$  holds iff  $\text{pre}(a) \subseteq s$ , i.e. if its precondition holds in  $s$ . The result of executing an applicable action  $a \in A$  in a state  $s$  is a new state  $\theta(s, a) = (s \setminus \neg \text{eff}(a)) \cup \text{eff}(a)$ . Subtracting the complement of  $\text{eff}(a)$  from  $s$  ensures that  $\theta(s, a)$  remains a well-defined state. The subset of action effects that assign

a positive value to a state fluent is called *positive effects* and denoted by  $\text{eff}^+(a) \in \text{eff}(a)$  while  $\text{eff}^-(a) \in \text{eff}(a)$  denotes the *negative effects* of an action  $a \in A$ .

A *classical planning problem* is a tuple  $P = \langle F, A, I, G \rangle$ , where  $I$  is the initial state and  $G \in \mathcal{L}(F)$  is the set of goal conditions over the state variables. A *plan*  $\pi$  is an action sequence  $\pi = \langle a_1, \dots, a_n \rangle$ , with  $|\pi| = n$  denoting its *plan length* and  $\text{cost}(\pi) = \sum_{a \in \pi} \text{cost}(a)$  its *plan cost*. The execution of  $\pi$  on the initial state of  $P$  induces a *trajectory*  $\tau(\pi, P) = \langle s_0, a_1, s_1, \dots, a_n, s_n \rangle$  such that  $s_0 = I$  and, for each  $1 \leq i \leq n$ , it holds  $\rho(s_{i-1}, a_i)$  and  $s_i = \theta(s_{i-1}, a_i)$ . A plan  $\pi$  solves  $P$  iff the induced *trajectory*  $\tau(\pi, P)$  reaches a final state  $G \subseteq s_n$ , where all goal conditions are met. A solution plan is *optimal* iff its cost is minimal.

We also define *actions with conditional effects* because they are useful to compactly formulate our approach for *goal recognition with unknown domain models*. An action  $a_c \in A$  with conditional effects is a set of preconditions  $\text{pre}(a_c) \in \mathcal{L}(F)$  and a set of *conditional effects*  $\text{cond}(a_c)$ . Each conditional effect  $C \triangleright E \in \text{cond}(a_c)$  is composed of two sets of literals:  $C \in \mathcal{L}(F)$ , the *condition*, and  $E \in \mathcal{L}(F)$ , the *effect*. An action  $a_c$  is applicable in a state  $s$  if  $\rho(s, a_c)$  is true, and the result of applying action  $a_c$  in state  $s$  is  $\theta(s, a_c) = \{s \setminus \neg \text{eff}_c(s, a) \cup \text{eff}_c(s, a)\}$  where  $\text{eff}_c(s, a)$  are the *triggered effects* resulting from the action application (conditional effects whose conditions hold in  $s$ ):

$$\text{eff}_c(s, a) = \bigcup_{C \triangleright E \in \text{cond}(a_c), C \subseteq s} E,$$

### 2.2 The observation model

Given a planning problem  $P = \langle F, A, I, G \rangle$ , a plan  $\pi$  and a trajectory  $\tau(\pi, P)$ , we define the *observation of the trajectory* as an interleaved combination of actions and states that represents the observation from the execution of  $\pi$  in  $P$ . Formally,  $O(\tau) = \langle s_0^o, a_1^o, s_1^o, \dots, a_l^o, s_l^o, s_0^o = I \rangle$ , and:

- The **observed actions** are consistent with  $\pi$ , which means that  $\langle a_1^o, \dots, a_l^o \rangle$  is a sub-sequence of  $\pi$ . The number of observed actions,  $l$ , ranges from 0 (fully unobserved action sequence) to  $|\pi|$  (fully observed action sequence).
- The **observed states**  $\langle s_0^o, s_1^o, \dots, s_m^o \rangle$  is a sequence of possibly *partially observable states*, except for the initial state  $s_0^o$ , which is fully observed. A partially observable

state  $s_i^o$  is one in which  $|s_i^o| < |F|$ ; i.e., a state in which at least a fluent of  $F$  is not observable. Note that this definition also comprises the case  $|s_i^o| = 0$ , when the state is fully unobservable. Whatever the sequence of observed states of  $\mathcal{O}(\tau)$  is, it must be consistent with the sequence of states of  $\tau(\pi, P)$ , meaning that  $\forall i, s_i^o \subseteq s_i$ . The number of observed states,  $m$ , range from 1 (the initial state, at least), to  $|\pi| + 1$ , and each *observed* states comprises  $[1, |F|]$  fluents (the observation can still miss intermediate states that are *unobserved*).

We assume a bijective monotone mapping between actions/states of trajectories and observations [Ramírez and Geffner, 2009], thus also granting the inverse consistency relationship (the trajectory is a superset of the observation). Therefore, transiting between two consecutive observed states in  $\mathcal{O}(\tau)$  may require the execution of more than a single action ( $\theta(s_i^o, \langle a_1, \dots, a_k \rangle) = s_{i+1}^o$ , where  $k \geq 1$  is unknown but finite. In other words, having an input observation  $\mathcal{O}(\tau)$  does not imply knowing the actual length of  $\pi$ .

### 2.3 State-invariants

The notion of *state-constraint* is very general and has been used in different areas of AI and for different purposes. If we restrict ourselves to planning, *state-constraints* are abstractions for compactly specifying sets of states. For instance, *state-constraints* in planning allow to specify the set of states where a given action is applicable, the set of states where a given *derived predicate* holds or the set of states that are considered goal states.

*State invariants* is a kind of state-constraints useful for computing more compact state representations [Helmert, 2009] or making *satisfiability planning* and *backward search* more efficient [Rintanen, 2014; Alcázar and Torralba, 2015]. Given a classical planning problem  $P = \langle F, A, I, G \rangle$ , a *state invariant* is a formula  $\phi$  that holds at the initial state of a given classical planning problem,  $I \models \phi$ , and at every state  $s$ , built from  $F$ , that is reachable from  $I$  by applying actions in  $A$ .

The formula  $\phi_{I,A}^*$  represents the *strongest invariant* and exactly characterizes the set of all states reachable from  $I$  with the actions in  $A$ . For instance Figure 1 shows five clauses that define the *strongest invariant* for *blocksworld*. There are infinitely many strongest invariants, but they are all logically equivalent, and computing the strongest invariant is PSPACE-hard as hard as testing plan existence.

$$\begin{aligned} \forall x_1, x_2 \text{ ontable}(x_1) &\leftrightarrow \neg \text{on}(x_1, x_2). \\ \forall x_1, x_2 \text{ clear}(x_1) &\leftrightarrow \neg \text{on}(x_2, x_1). \\ \forall x_1, x_2, x_3 \neg \text{on}(x_1, x_2) \vee \neg \text{on}(x_1, x_3) &\text{ such that } x_2 \neq x_3. \\ \forall x_1, x_2, x_3 \neg \text{on}(x_2, x_1) \vee \neg \text{on}(x_3, x_1) &\text{ such that } x_2 \neq x_3. \\ \forall x_1, \dots, x_n \neg (\text{on}(x_1, x_2) \wedge \text{on}(x_2, x_3) \wedge \dots \wedge \text{on}(x_{n-1}, x_n) \wedge \text{on}(x_n, x_1)). \end{aligned}$$

Figure 1: Example of the *strongest invariant* for the *blocksworld* domain.

A *mutex* (mutually exclusive) is a state invariant that takes the form of a binary clause and indicates a pair of different properties that cannot be simultaneously true [Kautz and Selman, 1999]. For instance in a three-block *blocksworld*,

$\phi_1 = \neg \text{on}(\text{block}_A, \text{block}_B) \vee \neg \text{on}(\text{block}_A, \text{block}_C)$  is a mutex because  $\text{block}_A$  can only be on top of a single block.

A *domain invariant* is an instance-independent invariant, i.e. holds for any possible initial state and set of objects. Therefore, if a given state  $s$  holds  $s \models \phi$  such that  $\phi$  is a *domain invariant*, it means that  $s$  is not a valid state. Domain invariants are often compactly defined as *lifted invariants* (also called schematic invariants) [Rintanen and others, 2017]. For instance,  $\phi_2 = \forall x : (\neg \text{handempty} \vee \neg \text{holding}(x))$ , is a *domain mutex* for the *blocksworld* because the robot hand is never empty and holding a block at the same time.

## 3 Deductive learning of planning action models

The task of learning action models is defined by  $\Lambda = \langle \mathcal{M}, \Phi, \mathcal{O}(\tau) \rangle$ :

- $\mathcal{M}$ , is the **initial domain model** (set of action models). This set is *empty*, when learning from scratch, or *partially specified*, when some fragments of the action models are known a priori.
- $\Phi$ , a set of *state-invariants* that constraint the set of possible states in the given domain.
- $\mathcal{O}(\tau)$  is an observation of a trajectory  $\tau(\pi, P)$  produced by the execution of an unknown plan  $\pi$  that reaches the goals  $G \in G[\cdot]$  starting from the initial state  $I$  in  $P$ .

A *solution* to a learning task  $\Lambda = \langle \mathcal{M}, \mathcal{O}(\tau) \rangle$  is a domain model  $\mathcal{M}'$  that is consistent with the information of  $\mathcal{M}$  and with the observed plan trace  $\tau$ . This means that the action sequence (plan) that solves the planning problem  $\langle s_0, s_n \rangle$  with  $\mathcal{M}'$  along with the state trajectory induced by this plan is consistent with the input observation  $\mathcal{O}(\tau)$  and only traverses states that satisfy the states invariants.

### 3.1 Learning of planning action models with classical planning

Because of the combinatorial nature of the search for a solution plan, the sooner unpromising nodes are pruned from the search the more efficient the computation of a solution plan. Constraints can be used to confine earlier the set of possible STRIPS action models and reduce then the learning hypothesis space. With regard to our compilation, *domain mutex* are useful to reduce the amount of applicable actions for programming a precondition or an effect for a given action schema. For example given the *domain mutex*  $\phi = (\neg f_1 \vee \neg f_2)$  such that  $f_1 \in F_v(\xi)$  and  $f_2 \in F_v(\xi)$ , we can redefine the corresponding programming actions for **removing** the *precondition*  $f_1 \in F_v(\xi)$  from the action schema  $\xi \in \mathcal{M}$  as:

$$\begin{aligned} \text{pre}(\text{programPre}_{f_1, \xi}) &= \{\neg \text{del}_{f_1}(\xi), \neg \text{add}_{f_1}(\xi), \text{mode}_{\text{prog}}, \text{pre}_{f_1}(\xi), \text{pre}_{f_2}(\xi)\} \\ \text{cond}(\text{programPre}_{f_1, \xi}) &= \{\emptyset\} \triangleright \{\neg \text{pre}_{f_1}(\xi)\}. \end{aligned}$$

## 4 Evaluation

## 5 Conclusions

## References

- [Aineto *et al.*, 2018] Diego Aineto, Sergio Jiménez, and Eva Onaindia. Learning STRIPS action models with classical planning. In *International Conference on Automated Planning and Scheduling, (ICAPS-18)*, pages 399–407. AAAI Press, 2018.
- [Alcázar and Torralba, 2015] Vidal Alcázar and Alvaro Torralba. A reminder about the importance of computing and exploiting invariants in planning. In *ICAPS*, pages 2–6. AAAI Press, 2015.
- [Helmert, 2009] Malte Helmert. Concise finite-domain representations for pddl planning tasks. *Artificial Intelligence*, 173(5-6):503–535, 2009.
- [Kautz and Selman, 1999] Henry Kautz and Bart Selman. Unifying SAT-based and graph-based planning. In *IJCAI*, volume 99, pages 318–325, 1999.
- [Ramírez and Geffner, 2009] Miquel Ramírez and Hector Geffner. Plan recognition as planning. In *International Joint conference on Artificial Intelligence, (IJCAI-09)*, pages 1778–1783. AAAI Press, 2009.
- [Rintanen and others, 2017] Jussi Rintanen et al. Schematic invariants by reduction to ground invariants. In *AAAI*, pages 3644–3650, 2017.
- [Rintanen, 2014] Jussi Rintanen. Madagascar: Scalable planning with SAT. In *International Planning Competition, (IPC-2014)*, 2014.