

Learning STRIPS action models from *state-invariants*

Diego Aineto¹, Sergio Jiménez¹, Eva Onaindia¹

¹Departamento de Sistemas Informáticos y Computación. Universitat Politècnica de València. Valencia, Spain

{dieaigar,serjice,onaindia}@dsic.upv.es, miquel.ramirez@unimelb.edu.au

Abstract

1 Introduction

The specification of planning action models is a complex task, even for planning experts, which frequently causes a knowledge acquisition bottleneck that limits the potential of AI planning [Kambhampati, 2007]. The *machine learning* of action models aims to relieve this knowledge acquisition bottleneck. Nowadays there is a wide range of different approaches for learning planning action models [Arora *et al.*, 2018]. However, most of these approaches aim maximizing some measure of *statistical consistency* over large sets of input examples that typically are sequential plans.

Classical planning is an interesting approach for learning STRIPS action models since it is flexible to different kinds of input knowledge (e.g., partially/fully observations of actions of plan executions as well as partially/fully observed intermediate states) and guarantees to output an action model that is *consistent* with the input knowledge [Aineto *et al.*, 2018]. In this paper we show that this flexibility of the compilation approach goes one step further since it also allows learning from *state-invariants*, logic formulae that constrain the set of possible states of a given domain. The experimental results show that *state-invariants* boost the performance of the *classical planing* compilation for learning STRIPS action models.

2 Background

This section formalizes the *classical planning model* we follow in this work and the kind of *knowledge* that can be given as input to the task of learning STRIPS action models.

2.1 Classical planning with conditional effects

Let F be the set of propositional state variables (*fluents*) describing a state. A *literal* l is a valuation of a fluent $f \in F$; i.e. either $l = f$ or $l = \neg f$. A set of literals L represents a partial assignment of values to fluents (without loss of generality, we will assume that L does not contain conflicting values). Given L , let $\neg L = \{\neg l : l \in L\}$ be its complement. We use $\mathcal{L}(F)$ to denote the set of all literal sets on F ; i.e. all partial assignments of values to fluents. A *state* s is a full assignment of values to fluents; $|s| = |F|$.

A *classical planning action* $a \in A$ has: a precondition $\text{pre}(a) \in \mathcal{L}(F)$, a set of effects $\text{eff}(a) \in \mathcal{L}(F)$, and a positive action cost $\text{cost}(a)$. The semantics of actions $a \in A$ is specified with two functions: $\rho(s, a)$ denotes whether action a is *applicable* in a state s and $\theta(s, a)$ denotes the *successor state* that results of applying action a in a state s . Then, $\rho(s, a)$ holds iff $\text{pre}(a) \subseteq s$, i.e. if its precondition holds in s . The result of executing an applicable action $a \in A$ in a state s is a new state $\theta(s, a) = (s \setminus \neg \text{eff}(a)) \cup \text{eff}(a)$. Subtracting the complement of $\text{eff}(a)$ from s ensures that $\theta(s, a)$ remains a well-defined state. The subset of action effects that assign a positive value to a state fluent is called *positive effects* and denoted by $\text{eff}^+(a) \in \text{eff}(a)$ while $\text{eff}^-(a) \in \text{eff}(a)$ denotes the *negative effects* of an action $a \in A$.

A *classical planning problem* is a tuple $P = \langle F, A, I, G \rangle$, where I is the initial state and $G \in \mathcal{L}(F)$ is the set of goal conditions over the state variables. A *plan* π is an action sequence $\pi = \langle a_1, \dots, a_n \rangle$, with $|\pi| = n$ denoting its *plan length* and $\text{cost}(\pi) = \sum_{a \in \pi} \text{cost}(a)$ its *plan cost*. The execution of π on the initial state of P induces a *trajectory* $\tau(\pi, P) = \langle s_0, a_1, s_1, \dots, a_n, s_n \rangle$ such that $s_0 = I$ and, for each $1 \leq i \leq n$, it holds $\rho(s_{i-1}, a_i)$ and $s_i = \theta(s_{i-1}, a_i)$. A plan π solves P iff the induced *trajectory* $\tau(\pi, P)$ reaches a final state $G \subseteq s_n$, where all goal conditions are met. A solution plan is *optimal* iff its cost is minimal.

We also define *actions with conditional effects* because they are useful to compactly formulate our approach for *goal recognition with unknown domain models*. An action $a_c \in A$ with conditional effects is a set of preconditions $\text{pre}(a_c) \in \mathcal{L}(F)$ and a set of *conditional effects* $\text{cond}(a_c)$. Each conditional effect $C \triangleright E \in \text{cond}(a_c)$ is composed of two sets of literals: $C \in \mathcal{L}(F)$, the *condition*, and $E \in \mathcal{L}(F)$, the *effect*. An action a_c is applicable in a state s if $\rho(s, a_c)$ is true, and the result of applying action a_c in state s is $\theta(s, a_c) = \{s \setminus \neg \text{eff}_c(s, a) \cup \text{eff}_c(s, a)\}$ where $\text{eff}_c(s, a)$ are the *triggered effects* resulting from the action application (conditional effects whose conditions hold in s):

$$\text{eff}_c(s, a) = \bigcup_{C \triangleright E \in \text{cond}(a_c), C \subseteq s} E,$$

2.2 The observation model

Given a planning problem $P = \langle F, A, I, G \rangle$, a plan π and a trajectory $\tau(\pi, P)$, we define the *observation of the trajectory*

as an interleaved combination of actions and states that represents the observation from the execution of π in P . Formally, $\mathcal{O}(\tau) = \langle s_0^o, a_1^o, s_1^o, \dots, a_l^o, s_m^o \rangle$, $s_0^o = I$, and:

- The **observed actions** are consistent with π , which means that $\langle a_1^o, \dots, a_l^o \rangle$ is a sub-sequence of π . The number of observed actions, l , ranges from 0 (fully unobserved action sequence) to $|\pi|$ (fully observed action sequence).
- The **observed states** $\langle s_0^o, s_1^o, \dots, s_m^o \rangle$ is a sequence of possibly *partially observable states*, except for the initial state s_0^o , which is fully observed. A partially observable state s_i^o is one in which $|s_i^o| < |F|$; i.e., a state in which at least a fluent of F is not observable. Note that this definition also comprises the case $|s_i^o| = 0$, when the state is fully unobservable. Whatever the sequence of observed states of $\mathcal{O}(\tau)$ is, it must be consistent with the sequence of states of $\tau(\pi, P)$, meaning that $\forall i, s_i^o \subseteq s_i$. The number of observed states, m , range from 1 (the initial state, at least), to $|\pi| + 1$, and each *observed states* comprises $[1, |F|]$ fluents (the observation can still miss intermediate states that are *unobserved*).

We assume a bijective monotone mapping between actions/states of trajectories and observations [Ramírez and Geffner, 2009], thus also granting the inverse consistency relationship (the trajectory is a superset of the observation). Therefore, transiting between two consecutive observed states in $\mathcal{O}(\tau)$ may require the execution of more than a single action ($\theta(s_i^o, \langle a_1, \dots, a_k \rangle) = s_{i+1}^o$, where $k \geq 1$ is unknown but finite. In other words, having an input observation $\mathcal{O}(\tau)$ does not imply knowing the actual length of π).

2.3 State-invariants

The notion of *state-constraint* is very general and has been used in different areas of AI and for different purposes. If we restrict ourselves to planning, *state-constraints* are abstractions for compactly specifying sets of states. For instance, *state-constraints* in planning allow to specify the set of states where a given action is applicable, the set of states where a given *derived predicate* holds or the set of states that are considered goal states.

State invariants is a kind of state-constraints useful for computing more compact state representations [Helmert, 2009] or making *satisfiability planning* and *backward search* more efficient [Rintanen, 2014; Alcázar and Torralba, 2015]. Given a classical planning problem $P = \langle F, A, I, G \rangle$, a *state invariant* is a formula ϕ that holds at the initial state of a given classical planning problem, $I \models \phi$, and at every state s , built from F , that is reachable from I by applying actions in A .

The formula $\phi_{I,A}^*$ represents the *strongest invariant* and exactly characterizes the set of all states reachable from I with the actions in A . For instance Figure 1 shows five clauses that define the *strongest invariant* for the *blocksworld* planning domain [Slaney and Thiébaux, 2001]. There are infinitely many strongest invariants, but they are all logically equivalent, and computing the strongest invariant is PSPACE-hard (as hard as testing plan existence [Bylander, 1994]).

A *mutex* (mutually exclusive) is a state invariant that takes the form of a binary clause and indicates a pair of differ-

$$\begin{aligned} \forall x_1, x_2 \text{ ontable}(x_1) &\leftrightarrow \neg \text{on}(x_1, x_2). \\ \forall x_1, x_2 \text{ clear}(x_1) &\leftrightarrow \neg \text{on}(x_2, x_1). \\ \forall x_1, x_2, x_3 \neg \text{on}(x_1, x_2) \vee \neg \text{on}(x_1, x_3) &\text{ such that } x_2 \neq x_3. \\ \forall x_1, x_2, x_3 \neg \text{on}(x_2, x_1) \vee \neg \text{on}(x_3, x_1) &\text{ such that } x_2 \neq x_3. \\ \forall x_1, \dots, x_n \neg (\text{on}(x_1, x_2) \wedge \text{on}(x_2, x_3) \wedge \dots \wedge \text{on}(x_{n-1}, x_n) \wedge & \\ \text{on}(x_n, x_1)) &. \end{aligned}$$

Figure 1: *Strongest invariant* for the *blocksworld* domain.

ent properties that cannot be simultaneously true [Kautz and Selman, 1999]. For instance in a three-block *blocksworld*, $\phi_1 = \neg \text{on}(\text{block}_A, \text{block}_B) \vee \neg \text{on}(\text{block}_A, \text{block}_C)$ is a mutex because *block_A* can only be on top of a single block.

A *domain invariant* is an instance-independent invariant, i.e. holds for any possible initial state and set of objects. Therefore, if a given state s holds $s \not\models \phi$ such that ϕ is a *domain invariant*, it means that s is not a valid state. Domain invariants are often compactly defined as *lifted invariants* (also called schematic invariants) [Rintanen and others, 2017]. For instance, $\phi_2 = \forall x : (\neg \text{handempty} \vee \neg \text{holding}(x))$, is a *domain mutex* for the *blocksworld* because the robot hand is never empty and holding a block at the same time.

3 Learning action models from state-invariants

We define the task of learning action models as a tuple $\Lambda = \langle M, \Phi, \mathcal{O}(\tau) \rangle$, where:

- M is the *space of possible action model*. This set is the *full* space of action models, when learning from scratch, or *partially specified*, when some fragments of the action models are known a priori. A set of action models can be defined *explicitly*, enumerating all the models that belong to the set or *implicitly*, enumerating all the constraints that must satisfy any model that belongs to the set. A *partially specified* STRIPS action model is then a formalism for the *implicit* representation of a set of STRIPS schemes [Sreedharan et al., 2018].
- Φ , a set of *state-invariants* that constrain the set of possible states in the given domain.
- $\mathcal{O}(\tau)$ is an observation of a trajectory $\tau(\pi, P)$ produced by the execution of an unknown plan π that solves $P_{\mathcal{O}}$ (i.e. the classical planning problem $P_{\mathcal{O}} = \langle F, A[\cdot], s_0^o, s_m^o \rangle$ where $A[\cdot]$ is a set of actions s.t., the semantics of each action $a \in A[\cdot]$ is unknown since functions ρ and/or θ of a are undefined).

A *solution* to a $\Lambda = \langle M, \Phi, \mathcal{O}(\tau) \rangle$ learning task is a model $\mathcal{M}' \in M$ that is *consistent* with the given input knowledge Φ and $\mathcal{O}(\tau)$. This means that let there exists a plan π such that the observation $\mathcal{O}(\tau)$ is consistent with π when the dynamics of $A[\cdot]$ is given by $\mathcal{M}' \in M$ and such that any state produced in the trajectory $\tau(\pi, P_{\mathcal{O}})$ satisfies the *state-invariants* Φ .

Next we show that the set M of possible action models can be encoded as a set of propositional variables and a set of constraints over those variables. Then, we show how to exploit this encoding to solve the $\Lambda = \langle M, \Phi, \mathcal{O}(\tau) \rangle$ learning task with an off-the-shelf classical planner.

```

(:action stack
  :parameters (?v1 ?v2)
  :precondition (and (holding ?v1) (clear ?v2))
  :effect (and (not (holding ?v1)) (not (clear ?v2))
              (clear ?v1) (handempty) (on ?v1 ?v2)))

(pre_holding_v1_stack) (pre_clear_v2_stack)
(ef_holding_v1_stack) (ef_clear_v2_stack)
(ef_clear_v1_stack) (ef_handempty_stack) (ef_on_v1_v2_stack)

```

Figure 2: PDDL encoding of the `stack(?v1, ?v2)` schema and our propositional representation for this same schema.

3.1 A propositional encoding for the space of STRIPS action models

A STRIPS *action schema* ξ is defined by four lists: A list of *parameters* $pars(\xi)$, and three list of predicates (namely $pre(\xi)$, $del(\xi)$ and $add(\xi)$) that shape the kind of fluents that can appear in the *preconditions*, *negative effects* and *positive effects* of the actions induced from that schema. Let be Ψ the set of *predicates* that shape the propositional state variables F , and a list of *parameters* $pars(\xi)$. The set of elements that can appear in $pre(\xi)$, $del(\xi)$ and $add(\xi)$ of the STRIPS action schema ξ is given by FOL interpretations of Ψ over the parameters $pars(\xi)$ and is denoted as $\mathcal{I}_{\Psi, \xi}$.

For instance in a four-operator *blocksworld* [Slaney and Thiébaux, 2001], the $\mathcal{I}_{\Psi, \xi}$ set contains only five elements for the `pickup(v_1)` schemata, $\mathcal{I}_{\Psi, pickup} = \{\text{handempty}, \text{holding}(v_1), \text{clear}(v_1), \text{ontable}(v_1), \text{on}(v_1, v_1)\}$ while it contains eleven elements for the `stack(v_1, v_2)` schemata, $\mathcal{I}_{\Psi, stack} = \{\text{handempty}, \text{holding}(v_1), \text{holding}(v_2), \text{clear}(v_1), \text{clear}(v_2), \text{ontable}(v_1), \text{ontable}(v_2), \text{on}(v_1, v_1), \text{on}(v_1, v_2), \text{on}(v_2, v_1), \text{on}(v_2, v_2)\}$.

Despite any element of $\mathcal{I}_{\Psi, \xi}$ can *a priori* appear in the $pre(\xi)$, $del(\xi)$ and $add(\xi)$ of schema ξ , the actual space of possible STRIPS schemata is bounded by constraints of three kinds:

1. **Syntactic constraints.** STRIPS constraints require $del(\xi) \subseteq pre(\xi)$, $del(\xi) \cap add(\xi) = \emptyset$ and $pre(\xi) \cap add(\xi) = \emptyset$. Considering exclusively these syntactic constraints, the size of the space of possible STRIPS schemata is given by $2^{2 \times |\mathcal{I}_{\Psi, \xi}|}$. *Typing constraints* are also of this kind [McDermott *et al.*, 1998].
2. **Domain-specific constraints.** One can introduce domain-specific knowledge to constrain further the space of possible schemata. For instance, in the *blocksworld* one can argue that $\text{on}(v_1, v_1)$ and $\text{on}(v_2, v_2)$ will not appear in the $pre(\xi)$, $del(\xi)$ and $add(\xi)$ lists of an action schema ξ because, in this specific domain, a block cannot be on top of itself. *State invariants* are also constraints of this kind.
3. **Observation constraints.** An observation $\mathcal{O}(\tau)$ depicts *semantic knowledge* that constraints further the space of possible action schemata.

In this work we introduce a propositional encoding of the *preconditions*, *negative*, and *positive effects* of a STRIPS action schema ξ using only fluents of two kinds $pre_e_ \xi$ and

$eff_e_ \xi$ (where $e \in \mathcal{I}_{\Psi, \xi}$). This encoding exploits the syntactic constraints of STRIPS so it is more compact than the one previously proposed by Aineto *et al.* 2018 for learning classical planning action models. In more detail, if $pre_e_ \xi$ holds it means that $e \in \mathcal{I}_{\Psi, \xi}$ is a *precondition* in ξ . If $pre_e_ \xi$ and $eff_e_ \xi$ holds it means that $e \in \mathcal{I}_{\Psi, \xi}$ is a *negative effect* in ξ while if $pre_e_ \xi$ does not hold but $eff_e_ \xi$ holds, it means that $e \in \mathcal{I}_{\Psi, \xi}$ is a *positive effect* in ξ . Figure 2 shows the PDDL encoding of the `stack(?v1, ?v2)` schema and our propositional representation for this same schema with pre_e_stack and eff_e_stack fluents ($e \in \mathcal{I}_{\Psi, stack}$).

3.2 Learning STRIPS action models with classical planning

Given a $\Lambda = \langle M, \Phi, \mathcal{O}(\tau) \rangle$ where Φ is a set of *domain mutex* $\phi \in \Phi$, we create a classical planning problem $P' = \langle F', A', I, G \rangle$ such that:

- F' extends F with a fluent $mode_{insert}$, to indicate whether action models are being programmed, and the fluents for the propositional encoding of the corresponding space of STRIPS action models. This is a set of fluents of the type $\{pre_e_ \xi, eff_e_ \xi\}_{e \in \mathcal{I}_{\Psi, \xi}}$ such that $e \in \mathcal{I}_{\Psi, \xi}$ is a single element from the set of FOL interpretations of predicates Ψ over the corresponding action parameters $pars(\xi)$.
- A' replaces the actions in A with two types of actions.

1. Actions for *inserting* a *precondition*, *positive effect* or *negative effect* in ξ following the syntactic constraints of STRIPS models. In the particular case that M is a *partially specified model* then only the actions for inserting a possible *precondition* or *effect* are necessary.

- Actions which support the addition of a *precondition* $p \in \Psi_\xi$ to the action model ξ . A precondition p is inserted in ξ when neither pre_p , eff_p exist in ξ .

$$\begin{aligned} pre(insertPre_p, \xi) &= \{\neg pre_p(\xi), \neg eff_p(\xi), mode_{insert}\}, \\ cond(insertPre_p, \xi) &= \{\emptyset\} \triangleright \{pre_p(\xi)\}. \end{aligned}$$

- Actions which support the addition of a *negative* or *positive effect* $p \in \Psi_\xi$ to the action model ξ .

$$\begin{aligned} pre(insertEff_p, \xi) &= \{\neg eff_p(\xi), mode_{insert}\}, \\ cond(insertEff_p, \xi) &= \{\emptyset\} \triangleright \{eff_p(\xi)\}. \end{aligned}$$

2. Actions for *applying* an action model ξ built by the *insert* actions and bounded to objects $\omega \subseteq \Omega^{|pars(\xi)|}$ (where Ω is the set of *objects* used to induce the fluents F by assigning objects in Ω to the Ψ predicates and Ω^k is the k -th Cartesian power of Ω). The action parameters, $pars(\xi)$, are bound to the objects in ω that appear in the same position.

$$\begin{aligned}
\text{pre}(\text{apply}_{\xi, \omega}) &= \{pre_p(\xi) \implies p(\omega)\}_{\forall p \in \Psi_\xi}, \\
\text{cond}(\text{apply}_{\xi, \omega}) &= \{pre_p(\xi) \wedge eff_p(\xi)\} \triangleright \{\neg p(\omega)\}_{\forall p \in \Psi_\xi}, \\
&\quad \{\neg pre_p(\xi) \wedge eff_p(\xi)\} \triangleright \{p(\omega)\}_{\forall p \in \Psi_\xi}, \\
&\quad \{\emptyset\} \triangleright \{\neg mode_{insert}\}.
\end{aligned}$$

domain mutex are useful to reduce the amount of applicable actions for programming a precondition or an effect for a given action schema. For example given the *domain mutex* $\phi = (\neg f_1 \vee \neg f_2)$ such that $f_1 \in F_v(\xi)$ and $f_2 \in F_v(\xi)$, we can redefine the corresponding programming actions for **removing** the *precondition* $f_1 \in F_v(\xi)$ from the action schema $\xi \in \mathcal{M}$ as:

4 Evaluation

5 Conclusions

References

- [Aineto *et al.*, 2018] Diego Aineto, Sergio Jiménez, and Eva Onaindia. Learning STRIPS action models with classical planning. In *International Conference on Automated Planning and Scheduling, (ICAPS-18)*, pages 399–407. AAAI Press, 2018.
- [Alcázar and Torralba, 2015] Vidal Alcázar and Alvaro Torralba. A reminder about the importance of computing and exploiting invariants in planning. In *ICAPS*, pages 2–6. AAAI Press, 2015.
- [Arora *et al.*, 2018] Ankuj Arora, Humbert Fiorino, Damien Pellier, Marc Métivier, and Sylvie Pesty. A review of learning planning action models. *The Knowledge Engineering Review*, 2018.
- [Bylander, 1994] Tom Bylander. The computational complexity of propositional strips planning. *Artificial Intelligence*, 69(1-2):165–204, 1994.
- [Helmert, 2009] Malte Helmert. Concise finite-domain representations for pddl planning tasks. *Artificial Intelligence*, 173(5-6):503–535, 2009.
- [Kambhampati, 2007] Subbarao Kambhampati. Model-lite planning for the web age masses: The challenges of planning with incomplete and evolving domain models. In *National Conference on Artificial Intelligence, (AAAI-07)*, 2007.
- [Kautz and Selman, 1999] Henry Kautz and Bart Selman. Unifying SAT-based and graph-based planning. In *IJCAI*, volume 99, pages 318–325, 1999.
- [McDermott *et al.*, 1998] Drew McDermott, Malik Ghallab, Adele Howe, Craig Knoblock, Ashwin Ram, Manuela Veloso, Daniel Weld, and David Wilkins. PDDL – The Planning Domain Definition Language, 1998.
- [Ramírez and Geffner, 2009] Miquel Ramírez and Hector Geffner. Plan recognition as planning. In *International Joint conference on Artificial Intelligence, (IJCAI-09)*, pages 1778–1783. AAAI Press, 2009.
- [Rintanen and others, 2017] Jussi Rintanen et al. Schematic invariants by reduction to ground invariants. In *AAAI*, pages 3644–3650, 2017.
- [Rintanen, 2014] Jussi Rintanen. Madagascar: Scalable planning with SAT. In *International Planning Competition, (IPC-2014)*, 2014.
- [Slaney and Thiébaux, 2001] John Slaney and Sylvie Thiébaux. Blocks world revisited. *Artificial Intelligence*, 125(1-2):119–153, 2001.
- [Sreedharan *et al.*, 2018] Sarath Sreedharan, Tathagata Chakraborti, and Subbarao Kambhampati. Handling model uncertainty and multiplicity in explanations via model reconciliation. In *International Conference on Automated Planning and Scheduling, (ICAPS-18)*, pages 518–526, 2018.